

# Introduction to HPC2N, Kebnekaise and HPC

Birgitte Brydsö,  
Pedro Ojeda May, and others at HPC2N

HPC2N  
Umeå University

8. September 2021



# HPC2N (HPC2N at a glance)

- ▶ **High Performance Computing Center North (HPC2N)** is a national center for Scientific and Parallel Computing



# HPC2N (HPC2N at a glance)

- ▶ **High Performance Computing Center North (HPC2N)** is a national center for Scientific and Parallel Computing



- ▶ A part of **Swedish National Infrastructure for Computing (SNIC)**



# HPC2N (HPC2N at a glance)

Provides state-of-the-art resources and expertise:

- ▶ Scalable and parallel **HPC**

# HPC2N (HPC2N at a glance)

Provides state-of-the-art resources and expertise:

- ▶ Scalable and parallel **HPC**
- ▶ Large-scale **storage facilities** (Project storage (Lustre), SweStore, Tape)

# HPC2N (HPC2N at a glance)

Provides state-of-the-art resources and expertise:

- ▶ Scalable and parallel **HPC**
- ▶ Large-scale **storage facilities** (Project storage (Lustre), SweStore, Tape)
- ▶ **Grid and cloud** computing (WLCG NT1, SNIC Cloud)

# HPC2N (HPC2N at a glance)

Provides state-of-the-art resources and expertise:

- ▶ Scalable and parallel **HPC**
- ▶ Large-scale **storage facilities** (Project storage (Lustre), SweStore, Tape)
- ▶ **Grid and cloud** computing (WLCG NT1, SNIC Cloud)
- ▶ Support
  - ▶ Primary, advanced, dedicated
  - ▶ Application Experts (AEs)

# HPC2N (HPC2N at a glance)

Provides state-of-the-art resources and expertise:

- ▶ Scalable and parallel **HPC**
- ▶ Large-scale **storage facilities** (Project storage (Lustre), SweStore, Tape)
- ▶ **Grid and cloud** computing (WLCG NT1, SNIC Cloud)
- ▶ Support
  - ▶ Primary, advanced, dedicated
  - ▶ Application Experts (AEs)
- ▶ International network for **research and development**



# HPC2N (partners)

HPC2N has five **partners**:

- ▶ Luleå University of Technology
- ▶ Mid Sweden University
- ▶ Swedish Institute of Space Physics
- ▶ Swedish University of Agricultural Sciences (SLU)
- ▶ Umeå University

# HPC2N (funding)

- ▶ Funded by **Swedish Research Council (VR)**, **SNIC** and various **partners**



Swedish  
Research  
Council



# HPC2N (funding)

- ▶ Funded by **Swedish Research Council (VR)**, **SNIC** and various **partners**



Swedish  
Research  
Council



- ▶ Involved in several **projects and collaborations**
  - ▶ EGI, PRACE, EISCAT, eSENCE, NOSEG, SNIC Science Cloud, ...

# HPC2N (training and other services)

- ▶ **User support** (primary, advanced, dedicated)
  - ▶ Research group meetings @ UmU
  - ▶ Also at the partner sites

# HPC2N (training and other services)

- ▶ **User support** (primary, advanced, dedicated)
  - ▶ Research group meetings @ UmU
  - ▶ Also at the partner sites
- ▶ **User training and education program**
  - ▶ 0.5 – 3 days; ready-to-run exercises
  - ▶ Introduction to HPC2N and Kebnekaise
  - ▶ Parallel programming and tools (e.g., OpenMP, MPI, debugging, performance analyzers, Matlab, R, MD simulation, Deep Learning, GPU, ...)

# HPC2N (training and other services)

- ▶ **User support** (primary, advanced, dedicated)
  - ▶ Research group meetings @ UmU
  - ▶ Also at the partner sites
- ▶ **User training and education program**
  - ▶ 0.5 – 3 days; ready-to-run exercises
  - ▶ Introduction to HPC2N and Kebnekaise
  - ▶ Parallel programming and tools (e.g., OpenMP, MPI, debugging, performance analyzers, Matlab, R, MD simulation, Deep Learning, GPU, ...)
- ▶ NGSSC / SeSE & university courses

# HPC2N (training and other services)

- ▶ **User support** (primary, advanced, dedicated)
  - ▶ Research group meetings @ UmU
  - ▶ Also at the partner sites
- ▶ **User training and education program**
  - ▶ 0.5 – 3 days; ready-to-run exercises
  - ▶ Introduction to HPC2N and Kebnekaise
  - ▶ Parallel programming and tools (e.g., OpenMP, MPI, debugging, performance analyzers, Matlab, R, MD simulation, Deep Learning, GPU, ...)
- ▶ NGSSC / SeSE & university courses
- ▶ Workshops and seminars

# HPC2N (personnel)

## Management

- ▶ Paolo Bientinesi, **new** director
- ▶ Björn Torkelsson, deputy director
- ▶ Lena Hellman, administrator



# HPC2N (personnel)

## Management

- ▶ Paolo Bientinesi, **new** director
- ▶ Björn Torkelsson, deputy director
- ▶ Lena Hellman, administrator

## Application experts

- ▶ Jerry Eriksson
- ▶ Pedro Ojeda May

# HPC2N (personnel)

## Management

- ▶ Paolo Bientinesi, **new** director
- ▶ Björn Torkelsson, deputy director
- ▶ Lena Hellman, administrator

## Application experts

- ▶ Jerry Eriksson
- ▶ Pedro Ojeda May

## Others

- ▶ Bo Kågström
- ▶ Mikael Rännar (WLCG coord)
- ▶ Anders Backman
- ▶ Kenneth Bodin
- ▶ Claude Lacoursière (Algoryx)

# HPC2N (personnel)

## Management

- ▶ Paolo Bientinesi, **new** director
- ▶ Björn Torkelsson, deputy director
- ▶ Lena Hellman, administrator

## Application experts

- ▶ Jerry Eriksson
- ▶ Pedro Ojeda May

## Others

- ▶ Bo Kågström
- ▶ Mikael Rännar (WLCG coord)
- ▶ Anders Backman
- ▶ Kenneth Bodin
- ▶ Claude Lacoursière (Algoryx)

## System and support

- ▶ Erik Andersson
- ▶ Birgitte Brydsö
- ▶ Niklas Edmundsson (Tape coord)
- ▶ Ingemar Fällman
- ▶ Magnus Jonsson
- ▶ Roger Oscarsson
- ▶ Åke Sandgren
- ▶ Mattias Wadenstein (NeIC, Tier1)
- ▶ Lars Viklund

# HPC2N (application experts)

- ▶ HPC2N provides advanced and dedicated support in the form of **Application Experts (AEs)**:

---

<sup>1</sup><https://www.snic.se/support/dedicated-user-support/>

# HPC2N (application experts)

- ▶ HPC2N provides advanced and dedicated support in the form of **Application Experts (AEs)**:

**Jerry Eriksson**   Profiling, Machine learning (DNN), MPI,  
OpenMP, OpenACC

---

<sup>1</sup><https://www.snic.se/support/dedicated-user-support/>

# HPC2N (application experts)

- ▶ HPC2N provides advanced and dedicated support in the form of **Application Experts (AEs)**:

**Jerry Eriksson** Profiling, Machine learning (DNN), MPI, OpenMP, OpenACC

**Pedro Ojeda May** Molecular dynamics, Profiling, QM/MM, NAMD, Amber, Gromacs, GAUSSIAN, R

---

<sup>1</sup><https://www.snic.se/support/dedicated-user-support/>

# HPC2N (application experts)

- ▶ HPC2N provides advanced and dedicated support in the form of **Application Experts (AEs)**:

**Jerry Eriksson**   Profiling, Machine learning (DNN), MPI, OpenMP, OpenACC

**Pedro Ojeda May**   Molecular dynamics, Profiling, QM/MM, NAMD, Amber, Gromacs, GAUSSIAN, R

**Åke Sandgren**   General high level programming assistance, VASP, Gromacs, Amber

---

<sup>1</sup><https://www.snlic.se/support/dedicated-user-support/>

# HPC2N (application experts)

- ▶ HPC2N provides advanced and dedicated support in the form of **Application Experts (AEs)**:

**Jerry Eriksson** Profiling, Machine learning (DNN), MPI, OpenMP, OpenACC

**Pedro Ojeda May** Molecular dynamics, Profiling, QM/MM, NAMD, Amber, Gromacs, GAUSSIAN, R

**Åke Sandgren** General high level programming assistance, VASP, Gromacs, Amber

- ▶ Contact through regular support or dedicated support form<sup>1</sup>
  - ▶ If you have a specific problem/question and/or need consultation (up to 100 h)

---

<sup>1</sup><https://www.snlic.se/support/dedicated-user-support/>



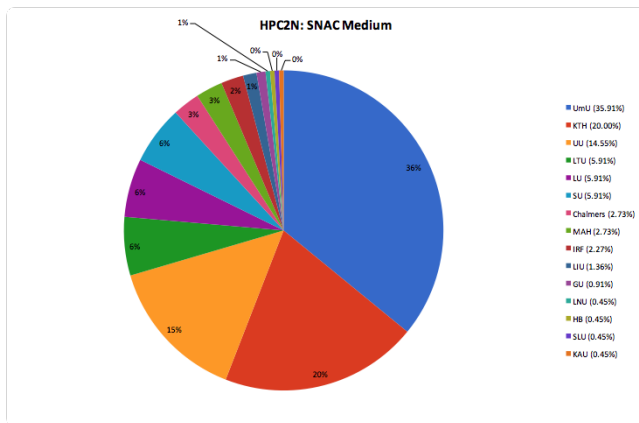
# HPC2N (users by discipline)

- ▶ Users from several scientific disciplines:
  - ▶ Biosciences and medicine
  - ▶ Chemistry
  - ▶ Computing science
  - ▶ Engineering
  - ▶ Materials science
  - ▶ Mathematics and statistics
  - ▶ Physics including space physics
  - ▶ Deep learning and artificial intelligence

# HPC2N (users by discipline, largest users)

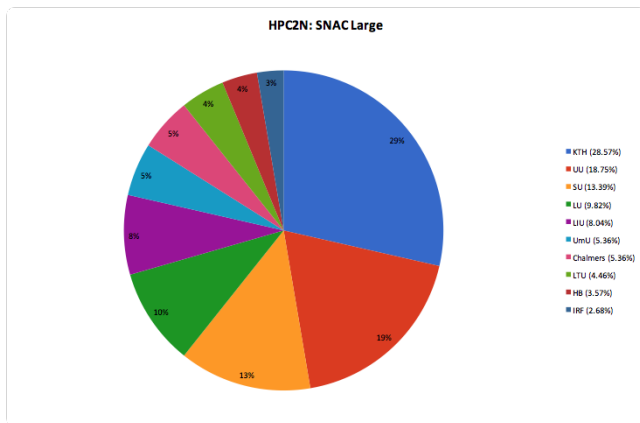
- ▶ Users from several scientific disciplines:
  - ▶ Biosciences and medicine
  - ▶ **Chemistry**
  - ▶ Computing science
  - ▶ Engineering
  - ▶ **Materials science**
  - ▶ Mathematics and statistics
  - ▶ **Physics including space physics**
  - ▶ **Deep learning and artificial intelligence** (several new projects)

# HPC2N (medium users by university)



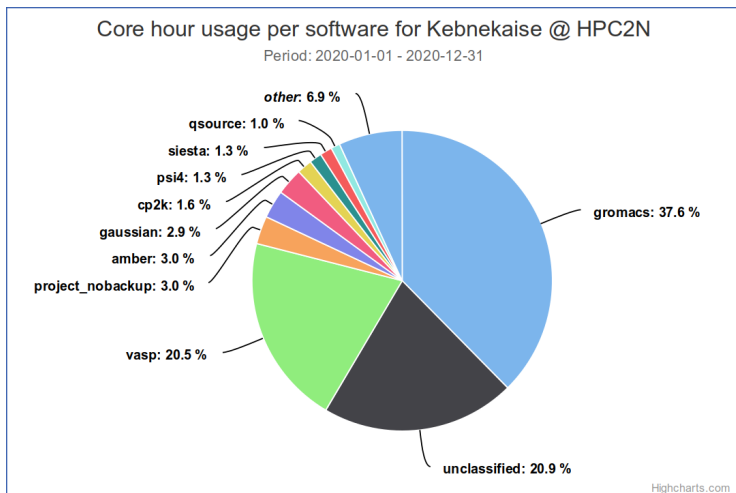
Projects with allocations at HPC2N: 2014-01-01 to 2016-05-30

# HPC2N (large users by university)



Projects with allocations at HPC2N: 2014-01-01 to 2016-05-30

# HPC2N (users by software)



- ▶ **Latest supercomputer at HPC2N**

- ▶ **Latest supercomputer at HPC2N**
- ▶ Named after a massif (contains some of Sweden's highest mountain peaks)

# Kebnekaise

- ▶ **Latest supercomputer at HPC2N**
- ▶ Named after a massif (contains some of Sweden's highest mountain peaks)
- ▶ Kebnekaise was
  - ▶ delivered by Lenovo and
  - ▶ **installed during the summer 2016**



# Kebnekaise

- ▶ **Latest supercomputer at HPC2N**
- ▶ Named after a massif (contains some of Sweden's highest mountain peaks)
- ▶ Kebnekaise was
  - ▶ delivered by Lenovo and
  - ▶ **installed during the summer 2016**
- ▶ Opened up for general availability on November 7, 2016

# Kebnekaise

- ▶ **Latest supercomputer at HPC2N**
- ▶ Named after a massif (contains some of Sweden's highest mountain peaks)
- ▶ Kebnekaise was
  - ▶ delivered by Lenovo and
  - ▶ **installed during the summer 2016**
- ▶ Opened up for general availability on November 7, 2016
- ▶ In 2018, Kebnekaise was **extended** with
  - ▶ 52 Intel Xeon Gold 6132 (Skylake) nodes, as well as
  - ▶ 10 NVidia V100 (Volta) GPU nodes

## Kebnekaise (compute nodes)

Name	#	Description
Compute	432	Intel Xeon E5-2690v4, <b>2 x 14 cores</b> , <b>128 GB</b> , FDR Infiniband

## Kebnekaise (compute nodes)

Name	#	Description
Compute	432	Intel Xeon E5-2690v4, <b>2 x 14 cores</b> , <b>128 GB</b> , FDR Infiniband
Compute-skylake	52	Intel Xeon Gold 6132, 2 x 14 cores, <b>192 GB</b> , EDR Infiniband, <b>AVX-512</b>

## Kebnekaise (compute nodes)

Name	#	Description
Compute	432	Intel Xeon E5-2690v4, <b>2 x 14 cores</b> , <b>128 GB</b> , FDR Infiniband
Compute-skylake	52	Intel Xeon Gold 6132, 2 x 14 cores, <b>192 GB</b> , EDR Infiniband, <b>AVX-512</b>
Large Memory	20	Intel Xeon E7-8860v4, <b>4 x 18 cores</b> , <b>3072 GB</b> , EDR Infiniband

## Kebnekaise (compute nodes)

Name	#	Description
Compute	432	Intel Xeon E5-2690v4, <b>2 x 14 cores</b> , <b>128 GB</b> , FDR Infiniband
Compute-skylake	52	Intel Xeon Gold 6132, 2 x 14 cores, <b>192 GB</b> , EDR Infiniband, <b>AVX-512</b>
Large Memory	20	Intel Xeon E7-8860v4, <b>4 x 18 cores</b> , <b>3072 GB</b> , EDR Infiniband
KNL	36	Intel <b>Xeon Phi</b> 7250 (Knight's Landing), 68 cores, 192 GB, 16 GB MCDRAM, FDR Infiniband

## Kebnekaise (GPU nodes)

Name	#	Description
2xGPU	32	Intel Xeon E5-2690v4, 2 x 14 cores, 128 GB, FDR Infiniband, <b>2 x NVidia K80</b> 4 x 2496 CUDA cores, 4 x 12 GB VRAM
-----		

## Kebnekaise (GPU nodes)

Name	#	Description
2xGPU	32	Intel Xeon E5-2690v4, 2 x 14 cores, 128 GB, FDR Infiniband, <b>2 x NVidia K80</b> 4 x 2496 CUDA cores, 4 x 12 GB VRAM
4xGPU	4	Intel Xeon E5-2690v4, 2 x 14 cores, 128 GB, FDR Infiniband, <b>4 x NVidia K80</b> 8 x 2496 CUDA cores, 8 x 12 GB VRAM



## Kebnekaise (GPU nodes)

Name	#	Description
2xGPU	32	Intel Xeon E5-2690v4, 2 x 14 cores, 128 GB, FDR Infiniband, <b>2 x NVidia K80</b> 4 x 2496 CUDA cores, 4 x 12 GB VRAM
4xGPU	4	Intel Xeon E5-2690v4, 2 x 14 cores, 128 GB, FDR Infiniband, <b>4 x NVidia K80</b> 8 x 2496 CUDA cores, 8 x 12 GB VRAM
GPU-volta	10	Intel Xeon Gold 6132, 2 x 14 cores, 192 GB, EDR Infiniband, <b>2 x NVidia V100,</b> 2 x 5120 CUDA cores, 2 x 16 GB VRAM, <b>2 x 640 Tensor cores</b>

# Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks

## Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks
- ▶ **19288 cores** (of which 2448 cores are KNL-cores)
  - ▶ 18840 available for users (the rest are for managing the cluster)

## Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks
- ▶ **19288 cores** (of which 2448 cores are KNL-cores)
  - ▶ 18840 available for users (the rest are for managing the cluster)
- ▶ More than **136 TB memory**

## Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks
- ▶ **19288 cores** (of which 2448 cores are KNL-cores)
  - ▶ 18840 available for users (the rest are for managing the cluster)
- ▶ More than **136 TB memory**
- ▶ 71 switches (Infiniband, Access and Management networks)

## Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks
- ▶ **19288 cores** (of which 2448 cores are KNL-cores)
  - ▶ 18840 available for users (the rest are for managing the cluster)
- ▶ More than **136 TB memory**
- ▶ 71 switches (Infiniband, Access and Management networks)
- ▶ 728 TFlops/s Peak performance (expansion not included)

## Kebnekaise (in numbers)

- ▶ 602 nodes in 15 racks
- ▶ **19288 cores** (of which 2448 cores are KNL-cores)
  - ▶ 18840 available for users (the rest are for managing the cluster)
- ▶ More than **136 TB memory**
- ▶ 71 switches (Infiniband, Access and Management networks)
- ▶ 728 TFlops/s Peak performance (expansion not included)
- ▶ **629 TFlops/s** Linpack (all parts, except expansion)
  - ▶ 86% of Peak performance

## Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:



# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ Home directory
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned

# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ Home directory
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned
  - ▶ Project storage
    - ▶ /proj/nobackup/abc
    - ▶ Shared among project members

# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ **Home directory**
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned
  - ▶ **Project storage**
    - ▶ /proj/nobackup/abc
    - ▶ Shared among project members
  - ▶ **Parallel file system **Deprecated!****
    - ▶ /pfs/nobackup/home/X/Xyz, user owned

# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ **Home directory**
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned
  - ▶ **Project storage**
    - ▶ /proj/nobackup/abc
    - ▶ Shared among project members
  - ▶ **Parallel file system **Deprecated!****
    - ▶ /pfs/nobackup/home/X/Xyz, user owned
  - ▶ **Local scratch space**
    - ▶ \$SNIC\_TMP
    - ▶ SSD (170GB), per job, per node, "volatile"

# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ **Home directory**
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned
  - ▶ **Project storage**
    - ▶ /proj/nobackup/abc
    - ▶ Shared among project members
  - ▶ **Parallel file system Deprecated!**
    - ▶ /pfs/nobackup/home/X/Xyz, user owned
  - ▶ **Local scratch space**
    - ▶ \$SNIC\_TMP
    - ▶ SSD (170GB), per job, per node, "volatile"
  - ▶ **SweStore** — disk based (dCache)
    - ▶ part of SNIC Storage, **nationally accessible storage**

# Kebnekaise (HPC2N storage)

- ▶ Basically five types of storage are available at HPC2N:
  - ▶ **Home directory**
    - ▶ /home/X/Xyz, \$HOME, ~
    - ▶ 25 GB, user owned
  - ▶ **Project storage**
    - ▶ /proj/nobackup/abc
    - ▶ Shared among project members
  - ▶ **Parallel file system Deprecated!**
    - ▶ /pfs/nobackup/home/X/Xyz, user owned
  - ▶ **Local scratch space**
    - ▶ \$SNIC\_TMP
    - ▶ SSD (170GB), per job, per node, "volatile"
  - ▶ **SweStore** — disk based (dCache)
    - ▶ part of SNIC Storage, **nationally accessible storage**
  - ▶ **Tape Storage**
    - ▶ Backup
    - ▶ **Long term storage**

# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**

# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**
  - ▶ A compute project has a certain number of **core hours** allocated for it per month



# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**
  - ▶ A compute project has a certain number of **core hours** allocated for it per month
  - ▶ A regular CPU core cost 1 core hour per hour, other resources (e.g., GPUs) cost more

# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**
  - ▶ A compute project has a certain number of **core hours** allocated for it per month
  - ▶ A regular CPU core cost 1 core hour per hour, other resources (e.g., GPUs) cost more
  - ▶ Not a hard limit but projects that go over the allocation get lower priority

# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**
  - ▶ A compute project has a certain number of **core hours** allocated for it per month
  - ▶ A regular CPU core cost 1 core hour per hour, other resources (e.g., GPUs) cost more
  - ▶ Not a hard limit but projects that go over the allocation get lower priority
- ▶ A compute project contains a certain amount of storage
  - ▶ If more storage is required, you must be a member of a **storage project**

# Kebnekaise (projects)

- ▶ In order to use Kebnekaise, you must be a member of a **compute project**
  - ▶ A compute project has a certain number of **core hours** allocated for it per month
  - ▶ A regular CPU core cost 1 core hour per hour, other resources (e.g., GPUs) cost more
  - ▶ Not a hard limit but projects that go over the allocation get lower priority
- ▶ A compute project contains a certain amount of storage
  - ▶ If more storage is required, you must be a member of a **storage project**
- ▶ I will cover more details in the next section, where we go more in to detail about HPC2N and Kebnekaise

# High Performance Computing (definition)

“High Performance Computing most generally refers to the practice of **aggregating computing power** in a way that delivers much **higher performance** than one could get out of a typical desktop computer or workstation in order to **solve large problems** in science, engineering, or business.”<sup>2</sup>

---

<sup>2</sup><https://insidehpc.com/hpc-basic-training/what-is-hpc/>

# High Performance Computing (opening the definition)

- ▶ **Aggregating computing power**

- ▶ 602 nodes in 15 racks totalling 19288 cores
- ▶ Compared to 4 cores in a modern laptop

---

<sup>3</sup>728 trillion (billion)

<sup>4</sup>200 billion (milliard)



# High Performance Computing (opening the definition)

- ▶ **Aggregating computing power**

- ▶ 602 nodes in 15 racks totalling 19288 cores
- ▶ Compared to 4 cores in a modern laptop

- ▶ **Higher performance**

- ▶ 728 000 000 000 000 arithmetical operations per second<sup>3</sup>
- ▶ Compared to 200 000 000 000 Flops in a modern laptop<sup>4</sup>

---

<sup>3</sup>728 trillion (billion)

<sup>4</sup>200 billion (milliard)

# High Performance Computing (opening the definition)

- ▶ **Aggregating computing power**

- ▶ 602 nodes in 15 racks totalling 19288 cores
- ▶ Compared to 4 cores in a modern laptop

- ▶ **Higher performance**

- ▶ 728 000 000 000 000 arithmetical operations per second<sup>3</sup>
- ▶ Compared to 200 000 000 000 Flops in a modern laptop<sup>4</sup>

- ▶ **Solve large problems**

- ▶ When does a problem become large enough for HPC?
- ▶ Are there other reasons for using HPC resources? (Memory, software, support, etc.)

---

<sup>3</sup>728 trillion (billion)

<sup>4</sup>200 billion (milliard)



# High Performance Computing (large problems)

- ▶ A problem can be large for two main reasons:
  1. **Execution time**: The time required to form a solution to the problem is very long
  2. **Memory / storage use**: The solution of the problem requires a lot of memory and/or storage


# High Performance Computing (large problems)

- ▶ A problem can be large for two main reasons:
  1. **Execution time**: The time required to form a solution to the problem is very long
  2. **Memory / storage use**: The solution of the problem requires a lot of memory and/or storage
- ▶ The former can be remedied by **increasing the performance**
  - ▶ More cores, more nodes, GPUs, ...


# High Performance Computing (large problems)

- ▶ A problem can be large for two main reasons:
  1. **Execution time**: The time required to form a solution to the problem is very long
  2. **Memory / storage use**: The solution of the problem requires a lot of memory and/or storage
- ▶ The former can be remedied by **increasing the performance**
  - ▶ More cores, more nodes, GPUs, ...
- ▶ The latter by **adding more memory / storage**
  - ▶ More memory per node (including large memory nodes), more nodes, ...
  - ▶ Large storage solutions, ...


# High Performance Computing (what counts as HPC)



# High Performance Computing (what counts as HPC)



# High Performance Computing (what counts as HPC)



# High Performance Computing (other reasons)

- ▶ Specialized (expensive) hardware

# High Performance Computing (other reasons)

- ▶ Specialized (expensive) hardware
  - ▶ GPUs, **Nvidia Tesla V100 GPUs** are optimized for AI



# High Performance Computing (other reasons)

- ▶ Specialized (expensive) hardware
  - ▶ GPUs, **Nvidia Tesla V100 GPUs** are optimized for AI
  - ▶ Intel Xeon Phi

# High Performance Computing (other reasons)

- ▶ Specialized (expensive) hardware
  - ▶ GPUs, **Nvidia Tesla V100 GPUs** are optimized for AI
  - ▶ Intel Xeon Phi
  - ▶ High-end CPUs (AVX-512 etc) and ECC memory

# High Performance Computing (other reasons)

- ▶ Specialized (expensive) hardware
  - ▶ GPUs, **Nvidia Tesla V100 GPUs** are optimized for AI
  - ▶ Intel Xeon Phi
  - ▶ High-end CPUs (AVX-512 etc) and ECC memory
- ▶ Software
  - ▶ HPC2N holds **licenses** for several softwares
  - ▶ Software is **pre-configured and ready-to-use**

# High Performance Computing (other reasons)


- ▶ Specialized (expensive) hardware
  - ▶ GPUs, **Nvidia Tesla V100 GPUs** are optimized for AI
  - ▶ Intel Xeon Phi
  - ▶ High-end CPUs (AVX-512 etc) and ECC memory
- ▶ Software
  - ▶ HPC2N holds **licenses** for several softwares
  - ▶ Software is **pre-configured and ready-to-use**
- ▶ **Support and documentation**

# High Performance Computing (memory models)

- ▶ Two memory models are relevant for HPC:

# High Performance Computing (memory models)


- ▶ Two memory models are relevant for HPC:
  - ▶ **Shared memory**: Single memory space for all data.



- ▶ **Everyone can access the same data**
- ▶ Straightforward to use


# High Performance Computing (memory models)

- ▶ Two memory models are relevant for HPC:
  - ▶ **Shared memory**: Single memory space for all data.




- ▶ **Everyone can access the same data**
- ▶ Straightforward to use

- ▶ **Distributed memory**: Multiple **distinct** memory spaces.



- ▶ Everyone has direct access **only to the local data**
- ▶ Requires **communication**

# High Performance Computing (memory models)





# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...
    - ▶ **Data distribution and communication**

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...
    - ▶ **Data distribution and communication**
- ▶ **GPUs**: **CUDA**, OpenCL, OpenACC, OpenMP, ...

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...
    - ▶ **Data distribution and communication**
- ▶ **GPUs**: **CUDA**, OpenCL, OpenACC, OpenMP, ...
  - ▶ **Many lightweight** streams of operations



# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...
    - ▶ **Data distribution and communication**
- ▶ **GPUs**: **CUDA**, OpenCL, OpenACC, OpenMP, ...
  - ▶ **Many lightweight** streams of operations
  - ▶ Work distribution, coordination (synchronization, etc), ...

# High Performance Computing (programming models)

- ▶ The programming model changes when we aim for extra performance and/or memory:
  1. **Single-core**: Matlab, Python, C, Fortran, ...
    - ▶ Single stream of operations
  2. **Multi-core**: Vectorized Matlab, pthreads, **OpenMP**
    - ▶ **Multiple streams** of operations
    - ▶ **Work distribution, coordination** (synchronization, etc), ...
  3. **Distributed memory**: **MPI**, ...
    - ▶ Multiple streams of operations
    - ▶ Work distribution, coordination (synchronization, etc), ...
    - ▶ **Data distribution and communication**
- ▶ **GPUs**: **CUDA**, OpenCL, OpenACC, OpenMP, ...
  - ▶ **Many lightweight** streams of operations
  - ▶ Work distribution, coordination (synchronization, etc), ...
  - ▶ **Data distribution across memory spaces and movement**

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc
    - ▶ **Allocate** correct number of cores, **configure** software to use correct number of cores, ...

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc
    - ▶ **Allocate** correct number of cores, **configure** software to use correct number of cores, ...
  3. **Distributed memory**: ScaLAPACK, ...
    - ▶ Load correct toolchain etc

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc
    - ▶ **Allocate** correct number of cores, **configure** software to use correct number of cores, ...
  3. **Distributed memory**: ScaLAPACK, ...
    - ▶ Load correct toolchain etc
    - ▶ Allocate correct number of **nodes and cores**, configure software to use correct number of **nodes and cores**, ...



# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc
    - ▶ **Allocate** correct number of cores, **configure** software to use correct number of cores, ...
  3. **Distributed memory**: ScaLAPACK, ...
    - ▶ Load correct toolchain etc
    - ▶ Allocate correct number of **nodes and cores**, configure software to use correct number of **nodes and cores**, ...
    - ▶ Data distribution, storage, ...

# High Performance Computing (software)

- ▶ Complexity grows when we aim for extra performance and/or memory/storage:
  1. **Single-core**: LAPACK, ...
    - ▶ Load correct toolchain etc
  2. **Multi-core**: LAPACK + parallel BLAS, ...
    - ▶ Load correct toolchain etc
    - ▶ **Allocate** correct number of cores, **configure** software to use correct number of cores, ...
  3. **Distributed memory**: ScaLAPACK, ...
    - ▶ Load correct toolchain etc
    - ▶ Allocate correct number of **nodes and cores**, configure software to use correct number of **nodes and cores**, ...
    - ▶ Data distribution, storage, ...
- ▶ **GPUs**: MAGMA, TensorFlow, ...
  - ▶ Load correct toolchain etc
  - ▶ Allocate correct number of **cores and GPUs**, configure software to use correct number of **cores and GPUs**, ...

End (questions?)

Questions?