

Introduction to HPC2N

Birgitte Brydsø, Pedro Ojeda-May

HPC2N, Umeå University

21 September 2023



- Projects - compute and storage
- Using our systems
- The File System
- The Module System
 - Overview
 - Compiler Tool Chains
 - Examples
- Compiling/linking with libraries
- The Batch System (SLURM)
 - Overview
 - Examples

Projects - compute and storage

Join a project or apply for one. **From 2023-01-01**

- Apply for a **compute project** in SUPR

<https://supr.snic.se/round/compute/>

- Small (≤ 5000 core-h/month, at least PhD student to apply)
- Medium (monthly rounds, at least assistant professor to apply)
- Large (bi-annual rounds)
- When applying for compute projects, you can apply for default storage (small is 500GB)
- More information:

<https://snic.se/allocations/compute/>

- If the above mentioned default storage is not enough, you will need to apply for a **storage project**

<https://supr.snic.se/round/storage/>

- Small (up to 2 TB, at least PhD student to apply)
- Medium (monthly rounds)
- Large (bi-annual rounds)
- More information:

<https://snic.se/allocations/storage/>

Projects - compute and storage

- As default, you have 25GB in your home directory.
- If you need more, you can accept the “default storage” you will be offered after applying for compute resources.
- The default storage is 500GB.
- If you need more than that, you will have to apply for a storage project. When you have both, link them together. It is done from the **storage** project.

Projects - compute and storage

After indicating the amount of compute resources you want, you will be asked if you want the default amount of storage:

[Admin](#)[User](#)[Start](#)[Reviews](#)[Reviewer Competence](#)[Proposals](#)[SNIC 2021/22-43](#) (HPC2N)[SNIC 2021/23-36](#) (HPC2N)[SNIC 2021/23-43](#)[hpc2n2021-001](#) (HPC2N)[Rounds](#)[Projects](#)[hpc2n2021-001](#) (HPC2N)[hpc2nstor2021-001](#) (HPC2N)[SNIC 2020/5-235](#) (multicentre)[SNIC 2020/9-215](#) (HPC2N)[SNIC 2021/23-36](#) (HPC2N)[Groups](#)[application-experts](#)[hpc2n-staff](#)[SONC](#)[Accounts](#)[Personal Information](#)[Support](#)[Logout](#)

Logged in as:

Birgitte Brydøb
(brydsoe@ics.umu.se)

Turn on warning colour.

[Start](#) / [Proposals](#) / [SNIC Small Compute 2021](#) / [SNIC 2021/22-43](#) / [Kebnekaise](#) / [Default Storage](#)

Default Storage for Kebnekaise @ HPC2N in SNIC 2021/22-43

Resource Kebnekaise added to proposal.

If you get compute time on Kebnekaise you can also get default storage for your project.

You may want to **decline** the default storage if

- you already have a separate storage project accessible from Kebnekaise that you intend to use for the storage related to this compute project, or
- the default storage offered is not enough for your needs, so you are going to create a proposal for a separate storage project.

In all other cases, we advise you to **accept** the default storage.


Default Storage Offered

Resource	Centre	Default	Unit	Default	Unit
Nobackup	HPC2N	500	GiB	1 000 000	files

[Decline Default Storage](#)[Accept Default Storage](#)

Projects - compute and storage

You will also be asked what the default storage directory should be called:



AdminUser

Start
Reviews
Reviewer Competence
Proposals
SNIC 2021/22-43 (HPC2N)
SNIC 2021/23-36 (HPC2N)
SNIC 2021/23-43
hpc2n2021-001 (HPC2N)
Rounds
Projects
hpc2n2021-001 (HPC2N)
hpc2nstor2021-001 (HPC2N)
SNIC 2020/5-235 (multicentre)
SNIC 2020/9-215 (HPC2N)
SNIC 2021/23-36 (HPC2N)
Groups
application-experts
hpc2n-staff
SONC
Accounts
Personal Information
Support
Logout

Logged in as:
Birgitte Brydsø
(bbryds@cs.umu.se)
Turn on warning colour.

Start / Proposals / SNIC Small Compute 2021 / SNIC 2021/22-43 / Select Directory Name

Select Directory Name for SNIC 2021/22-43

Default storage for Kebnekaise accepted. You now need to select a directory name for it.

A project has an associated directory name. Here, you will select the last part of the directory pathname used to access your storage.

Directory Name Type *

(no directory name selected yet)

Directory Name

Select

Explanation of directory name types

Inherit directory name "*DIRNAME*" as continuation of *PROJECT*

You will keep the directory and its contents from the earlier project, but now associated with your new project. This option will be available if you use the [continuation page](#) to select a project this is a continuation of.

Choose a new custom directory name

The new directory be empty from start.

You will have to provide a suitable name of 3-24 characters. It can contain lowercase a-z and the digits 0-9. Also, it can contain "-" and "_" but not at the start or the end. You cannot choose names that are already used by other projects, names that look like default directory names and names that are blocked by other use at centres.

Accept a new default directory name "snic2021-22-43"

The new directory will be empty from start.

Projects - compute and storage

- After applying on SUPR, the project(s) will be reviewed.
- When (if) the projects are approved, the PI needs to link the compute and storage projects together if they applied for both.
- Linking them together is done from the **storage** project.
- This way all members of the compute project also becomes members of the storage project.

Projects - compute and storage

Linking a compute project to a storage project:

Membership in this project can also be requested using the [Request Membership in Project](#) function. The Principal Investigator and the proxy will then get an email with a link in it, that is used to approve or deny the membership request.

Compute projects linked to this storage project

Members of linked compute projects that are active become extended members of this storage project and can access its storage.

Search for Compute Project to Link

Resources

Allocation shows the current allocation.

Storage

Percentage field to the right, compares **Usage** with the allocation. **Last Updated** shows the time at which the usage was last updated.

Resource	Allocation	Usage	Unit	Percentage	Allocation	Usage	Unit	Percentage	Last Updated
Nobackup @ HPC2N	2 000		GiB		999 999		files		


Resource Allocation and Usage

Storage

Show Storage Allocation and Usage on Nobackup @ HPC2N

Projects - compute and storage

Pick a compute project to link:



[Admin](#) [User](#)

Start
[Reviews](#)
[Reviewer Competence](#)
Proposals
[SNIC 2021/23-36](#) (HPC2N)
[hpc2n2021-001](#) (HPC2N)
Rounds
Projects
[hpc2n2021-001](#) (HPC2N)
[hpc2nstor2021-001](#) (HPC2N)
[SNIC 2020/5-235](#) (multicentre)
[SNIC 2020/9-215](#) (HPC2N)
[SNIC 2021/23-36](#) (HPC2N)
Groups
[application-experts](#)
[hpc2n-staff](#)
[SONC](#)
Accounts
Personal Information
Support
Logout

Logged in as:
Birgitte Brydsö
(bbrydsöe@cs.umu.se)
Turn on warning colour.

[Start](#) / [Projects](#) / [SNIC 2021/23-36](#) / [Link Compute Project](#)

Link Compute Project to SNIC 2021/23-36

Members of linked compute projects that are active become extended members of this storage project and can access its storage. Please note that if you link to a project belonging to another PI, you are allowing that PI to give access to your storage by adding members to his/her project's member list.

Use the search box below to search for matching projects (an empty search box will show projects with the same PI as this project). Projects that are already linked will not be shown.

Linkable projects with the same PI

Link	Compute Project	Title	PI
<input type="button" value="Link"/>	SNIC 2020/9-215	Introduction to HPC2N	Birgitte Brydsö

Projects - compute and storage

Showing linked projects:

Compute projects linked to this storage project

Members of linked compute projects that are active become extended members of this storage project and can access its storage.

Compute Project	Title	PI	Remove
SNIC 2020/9-215	Introduction to HPC2N	Birgitte Brydsö	<button>Remove</button>

Resources

Allocation shows the current allocation.

Storage

Percentage field to the right, compares **Usage** with the allocation. **Last Updated** shows the time at which the usage was last updated.

Resource	Allocation	Usage	Unit	Percentage	Allocation	Usage	Unit	Percentage	Last Updated
Nobackup @ HPC2N	2 000		GiB		999 999		files		

Projects - compute and storage

Members of the storage project after linking:

Project Members

These persons are direct members of this project and do not depend on [linked compute projects](#) for their access to the storage.

Name	Email	Organization / Department	Remove
Birgitte Brydsö	bbrydsoe@cs.umu.se	Umeå universitet / HPC2N	<i>PI cannot be removed</i>
Mirko Myllykoski	mirkom@cs.umu.se	Umeå universitet / Department of Computing Science and HPC2N	<button>Remove</button>
Pedro Ojeda	pedro.ojeda-may@umu.se	Umeå universitet / High Performance Computing Center North (HPC2N)	<button>Remove</button>

Add Member

Add Members from Other Project

Membership in this project can also be requested using the [Request Membership in Project](#) function. The Principal Investigator and the proxy will then get an email with a link in it, that is used to approve or deny the membership request.

Extended Project Members

These persons are members of this project or of active compute projects linked to this storage project. All extended members can access the storage belonging to this storage project.

Name	Email	Organization / Department
------	-------	---------------------------

Projects - after 2022-12-31?

- Nothing known for sure yet, but hopefully it will be clearer in a few weeks.
- Non-UmU PI: You can continue working at Kebnekaise until 2022-12-31, but you should be prepared for the situation where you will have to migrate to another centre.
- Hopefully Kebnekaise will at least exist as a local resource after the new year, until it can be replaced with a newer suitable resource.
- Local (UmU) PIs who have a need of using Kebnekaise (or another HPC2N resource) are encouraged to voice this to the people at UmU with the power to influence this.

Using our systems - example process

- 1 Connect to (SSH):

`kebnekaise.hpc2n.umu.se`

or use ThinLinc and connect to:

`kebnekaise-tl.hpc2n.umu.se`

- 2 Transfer your files and data (optionally, if needed)
- 3 Load modules (if needed)
- 4 Compile own code, install software, or run pre-installed software
- 5 Create batch script, submit batch job
- 6 Download data/results to local/other machine (optionally)

Using our systems

Connecting to HPC2N's systems - ThinLinc

ThinLinc: a cross-platform remote desktop server from Cendio AB. Especially useful when you need software with a graphical interface.

- We **recommend** ThinLinc if you don't have a preferred SSH client.
- Download the client from <https://www.cendio.com/thinlinc/download>. Install it.
- Start the client. Enter the name of the server: **kebnekaise-tl.hpc2n.umu.se**. Enter your username.
- Go to "Options" – > "Security". Check that authentication method is set to password.
- Go to "Options" – > "Screen". Uncheck "Full screen mode".
- Enter your HPC2N password. Click "Connect"
- Click "Continue" when you are being told that the server's host key is not in the registry. Wait for the ThinLinc desktop to open.

Using our systems

Transfer your files and data

- **Linux, OS X:**

- Use scp for file transfer:

```
local> scp username@kebnekaise.hpc2n.umu.se:file .  
local> scp file username@kebnekaise.hpc2n.umu.se:file
```

- **Windows:**

- Download client: WinSCP, FileZilla (sftp), PSCP/PSFTP, ...
- Transfer with sftp or scp

- <https://www.hpc2n.umu.se/documentation/filesystems/filetransfer>

- **Mac/OSX:**

- Transfer with sftp or scp (as for Linux) using Terminal
- Or download client: Cyberduck, Fetch, ...

- More info in guides (see previous slide) and here:

<https://www.hpc2n.umu.se/documentation/filesystems/filetransfer>

Editing your files

- Various editors: vi, vim, nano, emacs ...
- Example, nano:
 - `nano <filename>`
 - Save and exit nano: `Ctrl-x`
- Example, Emacs:
 - Start with: `emacs`
 - Open (or create) file: `Ctrl-x Ctrl-f`
 - Save: `Ctrl-x Ctrl-s`
 - Exit Emacs: `Ctrl-x Ctrl-c`
 - (If you want to run in an a separate emacs window, and with full functionality, you need to login with ThinLinc, `ssh -Y` or similar, for X11 forwarding)

The File System

More info here: <http://www.hpc2n.umu.se/filesystems/overview>

	Project storage	\$HOME	/scratch
Recommended for batch jobs	Yes	No (size)	Yes
Backed up	No	Yes	No
Accessible by batch system	Yes	Yes	Yes (node only)
Performance	High	High	Medium
Default readability	Group only	Owner	Owner
Permissions management	chmod, chgrp, ACL	chmod, chgrp, ACL	N/A for batch jobs
Notes	Storage your group get allocated through the storage projects	Your home-directory	Per node

Using project storage

- If you have a storage project, you should use that to run your jobs.
- You (your PI) will either choose a directory name when you/they apply for the storage project or get the SNIC id as default name.
- The location of the storage project in the file system is `/proj/nobackup/<name-you-picked>`.
- Since the storage project is shared between all users of the project, you should go to that directory and create a subdirectory for your things, which you will then be using.
- For this course the storage is in `/proj/nobackup/snic2022-22-642`

The Module System (Lmod)

Most programs are accessed by first loading them as a 'module'

Modules are

- used to set up your environment (paths to executables, libraries, etc.) for using a particular (set of) software package(s)
- a tool to help users manage their Unix/Linux shell environment, allowing groups of related environment-variable settings to be made or removed dynamically
- allows having multiple versions of a program or package available by just loading the proper module
- are installed in a hierarchical layout. This means that some modules are only available after loading a specific compiler and/or MPI version.

The Module System (Lmod)

Useful commands (Lmod)

- See which modules exists:
`module spider` or `ml spider`
- Modules depending only on what is currently loaded:
`module avail` or `ml av`
- See which modules are currently loaded:
`module list` or `ml`
- Example: loading a compiler toolchain, here for GCC:
`module load foss/version` or `ml foss/version`
- Example: Unload the above module:
`module unload foss` or `ml -foss`
- More information about a module:
`ml show <module>` or `module show <module>`
- Unload all modules except the 'sticky' modules:
`ml purge`

The Module System

Compiler Toolchains

Compiler toolchains load bundles of software making up a complete environment for compiling/using a specific prebuilt software. Includes some/all of: compiler suite, MPI, BLAS, LAPACK, ScaLapack, FFTW, CUDA.

- Some currently available toolchains (check `m1 av` for versions and full, updated list):
 - **GCC**: GCC only
 - **gcccuda**: GCC and CUDA
 - **foss**: GCC, OpenMPI, OpenBLAS/LAPACK, FFTW, ScaLAPACK
 - **gomp**: GCC, OpenMPI
 - **gompic**: GCC, OpenMPI, CUDA
 - **gomkl**: GCC, OpenMPI, MKL
 - **iccifort**: icc, ifort
 - **iccifortcuda**: icc, ifort, CUDA
 - **iimpi**: icc, ifort, IntelMPI
 - **intel**: icc, ifort, IntelMPI, IntelMKL
 - **intelcuda**: intel and CUDA
 - **iompi**: iccifort and OpenMPI

Compiling and Linking with Libraries

Linking

Figuring out how to link

- Intel and Intel MKL linking:

<https://software.intel.com/en-us/articles/intel-mkl-link-line-advisor>

- GCC, etc. **Use buildenv**

- After loading a compiler toolchain, load 'buildenv' and use 'ml show buildenv' to get useful linking info
- Example, foss (add relevant version):

```
ml foss/version
ml buildenv
ml show buildenv
```


- Using the environment variable (prefaced with \$) for linking is highly recommended!
- You have to load the buildenv module in order to use the environment variable for linking!

The Batch System (SLURM)

- Large/long/parallel jobs **must** be run through the batch system
- SLURM is an Open Source job scheduler, which provides three key functions
 - Keeps track of available system resources
 - Enforces local system resource usage and job scheduling policies
 - Manages a job queue, distributing work across resources according to policies
- In order to run a batch job, you need to create and submit a SLURM submit file (also called a batch submit file, a batch script, or a job script).
- Guides and documentation at:
<http://www.hpc2n.umu.se/support>


The Batch System

Accounting, Compute nodes, Kebnekaise




The Batch System

Accounting, largemem nodes, Kebnekaise



The Batch System

Accounting, GPU nodes, Kebnekaise. Same for the V100 as for the K80.



Note: V100s accounts like K80s and have **one** engine per card.

The Batch System (SLURM)

Useful Commands

- Submit job: `sbatch <jobscrip>`
- Get list of your jobs: `squeue -u <username>`
- `srun <commands for your job/program>`
- Check on a specific job: `scontrol show job <job id>`
- Delete a specific job: `scancel <job id>`
- Delete all your own jobs: `scancel -u <user>`
- More detailed info about jobs:
`sacct -l -j <jobid> -o jobname,NTasks,nodelist,MaxRSS,MaxVMSize...`
 - More flags can be found with `man sacct`
 - The output will be **very** wide. To view, use
`sacct -l -j | less -S`
(makes it sideways scrollable, using the left/right arrow key)
- Web url with graphical info about a job: `job-usage <job-id>`

Use `man sbatch`, `man srun`, `man` for more information

The Batch System (SLURM)

Job Output

- Output and errors in:
`slurm-<job id>.out`
- Look at it with `vi`, `nano`, `emacs`, `cat`, `less`...
- To get output and error files split up, you can give these flags in the submit script:
`#SBATCH --error=job.%J.err`
`#SBATCH --output=job.%J.out`

The Batch System (SLURM)

Using different parts of Kebnekaise

- To run on the 'fat' nodes, add this flag to your script:
`#SBATCH -p largemem` (Kebnekaise - largemem does not have general access)
- Specifying Intel Broadwell or Skylake CPUs only on Kebnekaise:
`#SBATCH --constraint=broadwell`
or
`#SBATCH --constraint=skylake`
- Using the GPU nodes on Kebnekaise
`#SBATCH --gres=gpu:<type-of-card>:x` where
<type-of-card> is either k80 or v100 and $x = 1, 2, \text{ or } 4$ (4 only for the K80 type).

More on

https://www.hpc2n.umu.se/documentation/guides/using_kebnekaise

The Batch System (SLURM)

Simple example, serial

Example: Serial job on Kebnekaise, compiler toolchain 'foss'

```
#!/bin/bash
# Project id - change to your own after the course!
#SBATCH -A SNIC2022-22-642
# Asking for 1 core
#SBATCH -n 1
# Asking for a walltime of 5 min
#SBATCH --time=00:05:00

# Purge modules before loading new ones in a script.
ml purge
ml foss/2021b

./my_serial_program
```

Submit with:

```
sbatch <jobscript>
```

The Batch System (SLURM)

Example, MPI C program

```
#include <stdio.h>
#include <mpi.h>

int main (int argc, char *argv[])

int myrank, size;

MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
MPI_Comm_size(MPI_COMM_WORLD, &size);

printf("Processor %d of %d:  Hello World!\n", myrank,
size);

MPI_Finalize();
```

The Batch System (SLURM)

Simple example, parallel

Example: MPI job on Kebnekaise, compiler toolchain 'foss'

```
#!/bin/bash
#SBATCH -A SNIC2021-22-642
#SBATCH -n 14
#SBATCH --time=00:05:00
##SBATCH --exclusive
#SBATCH --reservation=intro-cpu

module purge
ml foss/2021b

srun ./my_parallel_program
```


The Batch System (SLURM)

Simple example, output

Example: Output from a MPI job on Kebnekaise, run on 14 cores (one NUMA island)

```
b-an01 [~/slurm]$ cat slurm-15952.out
```

The following modules were not unloaded:

(Use "module --force purge" to unload all):

```
1) systemdefault 2) snicenvironment  
Processor 12 of 14: Hello World!  
Processor 5 of 14: Hello World!  
Processor 9 of 14: Hello World!  
Processor 4 of 14: Hello World!  
Processor 11 of 14: Hello World!  
Processor 13 of 14: Hello World!  
Processor 0 of 14: Hello World!  
Processor 1 of 14: Hello World!  
Processor 2 of 14: Hello World!  
Processor 3 of 14: Hello World!  
Processor 6 of 14: Hello World!  
Processor 7 of 14: Hello World!  
Processor 8 of 14: Hello World!  
Processor 10 of 14: Hello World!
```

The Batch System (SLURM)

Starting more than one serial job in the same submit file

```
#!/bin/bash
#SBATCH -A SNIC2021-22-642
#SBATCH -n 5
#SBATCH --time=00:15:00

module purge
ml foss/2021b

srun -n 1 ./job1.batch &
srun -n 1 ./job2.batch &
srun -n 1 ./job3.batch &
srun -n 1 ./job4.batch &
srun -n 1 ./job5.batch
wait
```

The Batch System (SLURM)

Multiple Parallel Jobs Sequentially

```
#!/bin/bash
#SBATCH -A SNIC2021-22-642
#SBATCH -c 28
# Remember to ask for enough time for all jobs to complete
#SBATCH --time=02:00:00

module purge
ml foss/2021b

# Here 14 tasks with 2 cores per task. Output to file.
# Not needed if your job creates output in a file
# I also copy the output somewhere else and then run
# another executable...

srun -n 14 -c 2 ./a.out > myoutput1 2>&1
cp myoutput1 /pfs/nobackup/home/u/username/mydatadir
srun -n 14 -c 2 ./b.out > myoutput2 2>&1
cp myoutput2 /pfs/nobackup/home/u/username/mydatadir
srun -n 14 -c 2 ./c.out > myoutput3 2>&1
cp myoutput3 /pfs/nobackup/home/u/username/mydatadir
```

The Batch System (SLURM)

Multiple Parallel Jobs Simultaneously

Make sure you ask for enough cores that all jobs can run at the same time, and have enough memory. Of course, this will also work for serial jobs - just remove the `srun` from the command line.

```
#!/bin/bash
#SBATCH -A SNIC2021-22-642
# Total number of cores the jobs need
#SBATCH -n 56
# Remember to ask for enough time for all of the jobs to
# complete, even the longest
#SBATCH --time=02:00:00

module purge
ml foss/2021b

srun -n 14 --cpu_bind=cores ./a.out &
srun -n 28 --cpu_bind=cores ./b.out &
srun -n 14 --cpu_bind=cores ./c.out &
...
wait
```

The Batch System (SLURM)

GPU Job

```
#!/bin/bash
#SBATCH -A SNIC2021-22-642
# Expected time for job to complete
#SBATCH --time=00:10:00
# Number of GPU cards needed. Here asking for 2 K80 cards
#SBATCH --gres=k80:2

module purge
ml fosscuda/2021b

./my-program
```

Important information

- The course project has the following SNIC ID:
SNIC2021-22-642
- In order to use it in a batch job, add this to the batch script:
 - `#SBATCH -A SNIC2021-22-642`
- There are some nodes reserved for the course, in order to let us run small examples without having to wait for too long. Two nodes are regular Broadwell CPU nodes, two nodes are K80 GPU nodes.
- These reservations are **ONLY** valid during the course:
 - `intro-cpu`
(add with `#SBATCH -reservation=intro-cpu`)
 - `intro-gpu`
(add with `#SBATCH -reservation=intro-gpu`)
- We have a storage project linked to the compute project. It is `snic2021-22-642`. You find it in
`/proj/nobackup/snic2021-22-642`

Questions? Now: Ask me or one of the other support or application experts present.

OR

- Documentation: <https://www.hpc2n.umu.se/support>
- Support questions to: <https://supr.snic.se/support/> or support@hpc2n.umu.se