



Applications' usage in HPC2N

P. Ojeda, B. Brydsö, M. Myllykoski, J. Eriksson



Why do we need HPC?

Scientific codes use loops extensively:

```
do j=1,N
    do i=1,N
        do k=1,N
            p(i,j) = p(i,j) + v1(i,k)* v2(k,j)
        enddo
    enddo
enddo
```

Heavy loops
MD, QM, CFD, ...

```
!$omp parallel do private(i,j,k) shared(p,v1,v2) collapse(2)
    do j=1,N
        do i=1,N
            s=0.0d0
            do k=1,N
                s = s + v1(i,k)* v2(k,j)
            enddo
            p(i,j) = s
        enddo
    enddo
!$omp end parallel do
```

Parallelized loops

Processes communication

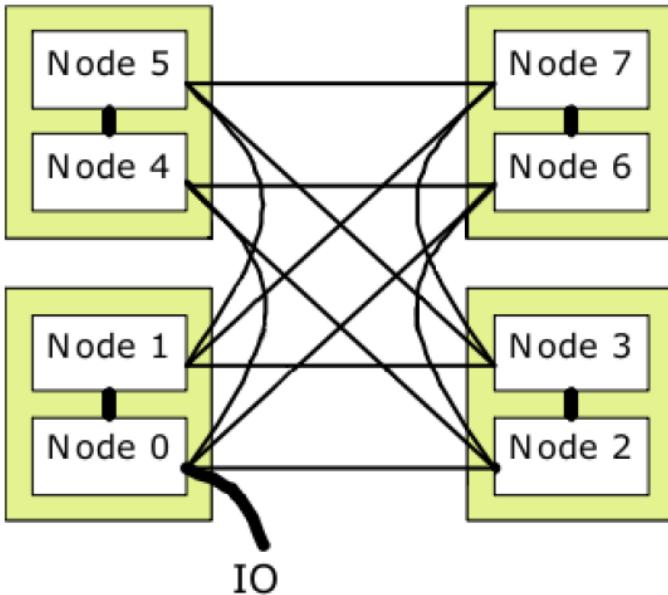


Figure : Nodes.

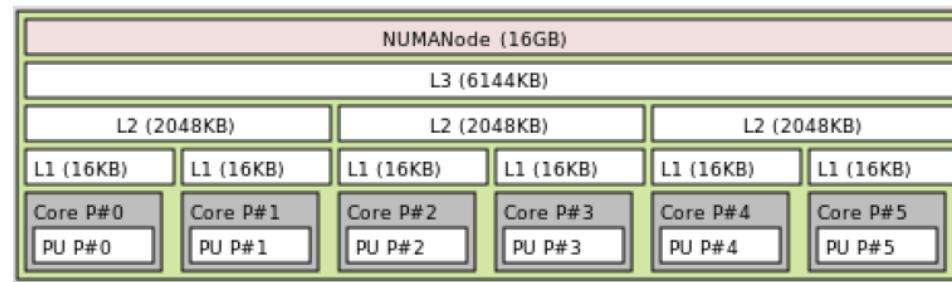
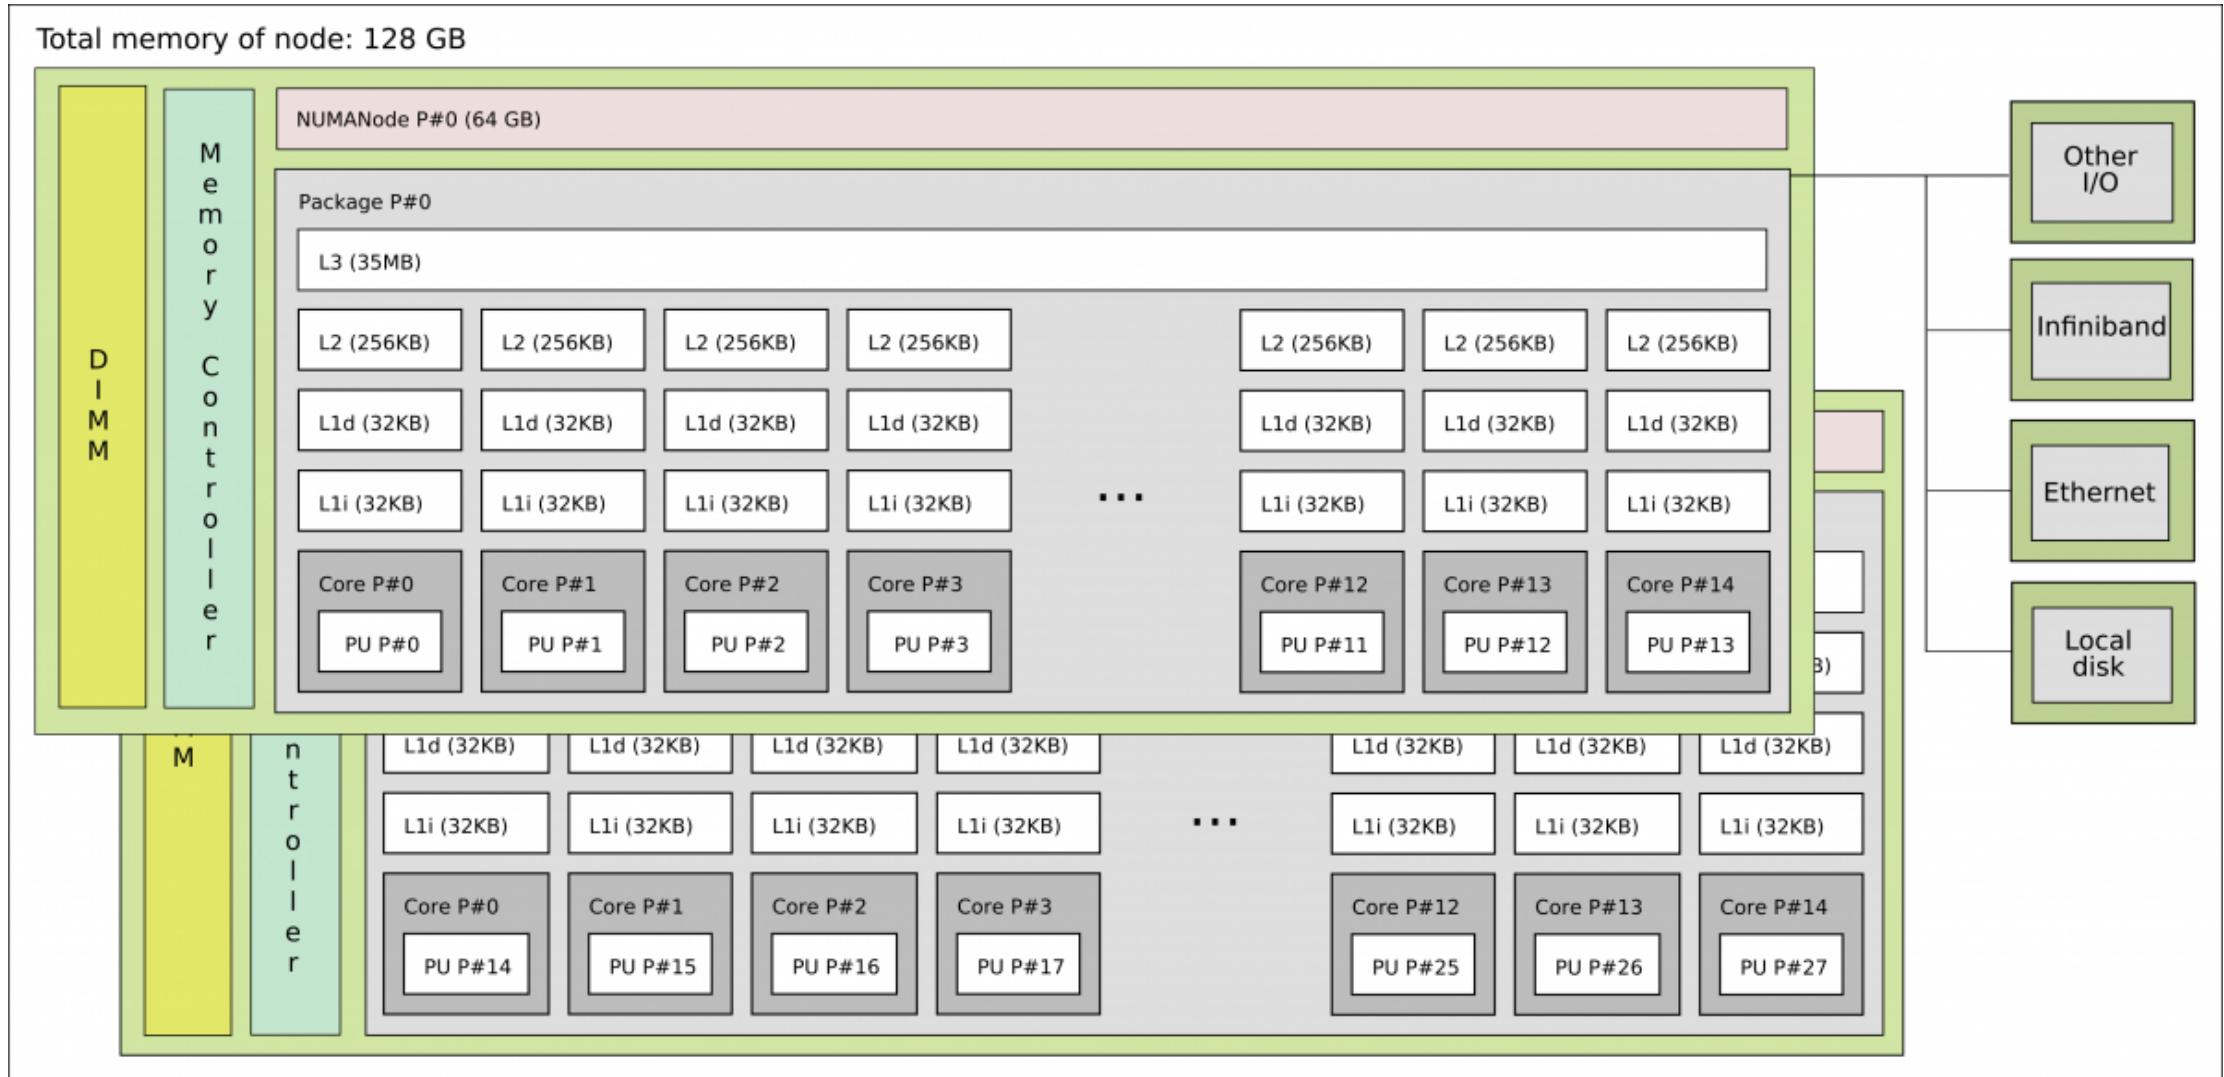


Figure : NUMA machine.

$$T_p = \frac{T_s}{P} + T_{comm}$$

Broadwell node on Kebnekaise



Accelerators



GPU showing the independent units
Streaming Multiprocessors (SM).



Single tile



KNL, composed of several Tiles

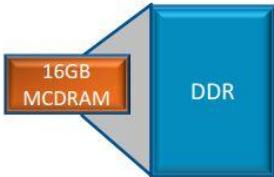
$$Z(i) = A * X(i) + Y(i) \quad (\text{Vector Op.})$$

KNL

Memory Modes

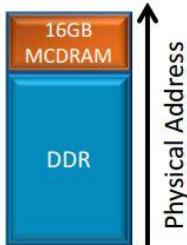
Three Modes. Selected at boot

Cache Mode



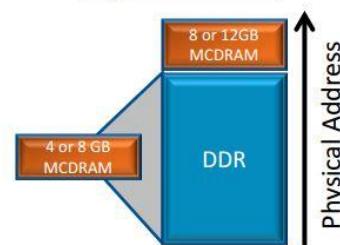
- SW-Transparent, Mem-side cache
- Direct mapped. 64B lines.
- Tags part of line
- Covers whole DDR range

Flat Mode



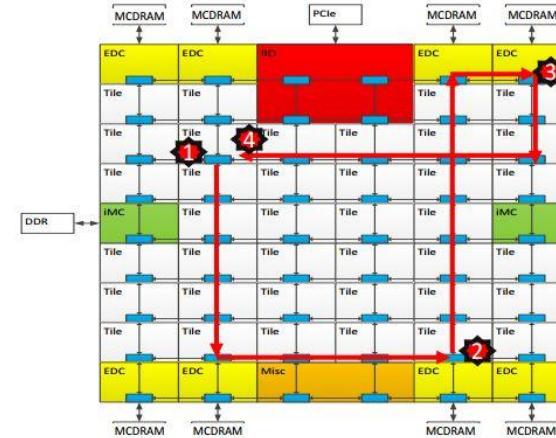
- MCDRAM as regular memory
- SW-Managed
- Same address space

Hybrid Mode



- Part cache, Part memory
- 25% or 50% cache
- Benefits of both

Cluster Mode: All-to-All



Address uniformly hashed across all distributed directories

No affinity between Tile, Directory and Memory

Most general mode. Lower performance than other modes.

Typical Read L2 miss

1. L2 miss encountered
2. Send request to the distributed directory
3. Miss in the directory. Forward to memory
4. Memory sends the data to the requestor

#SBATCH --constraint=a2a,cache

#SBATCH --gres=hbm:4G

Credits: PRACE Best practice KNL (2017)

KNL Thread affinity

There are physical 68 cores with 4 hyperthreads on each.



Bind the threads by using OpenMP env. var.
export OMP_NUM_THREADS=4
export OMP_PROC_BIND=spread
export OMP_PLACES=cores

`srun -n 68 -c 4 --cpu_bind=cores a_knl.out`

Alternatively, use Intel var.

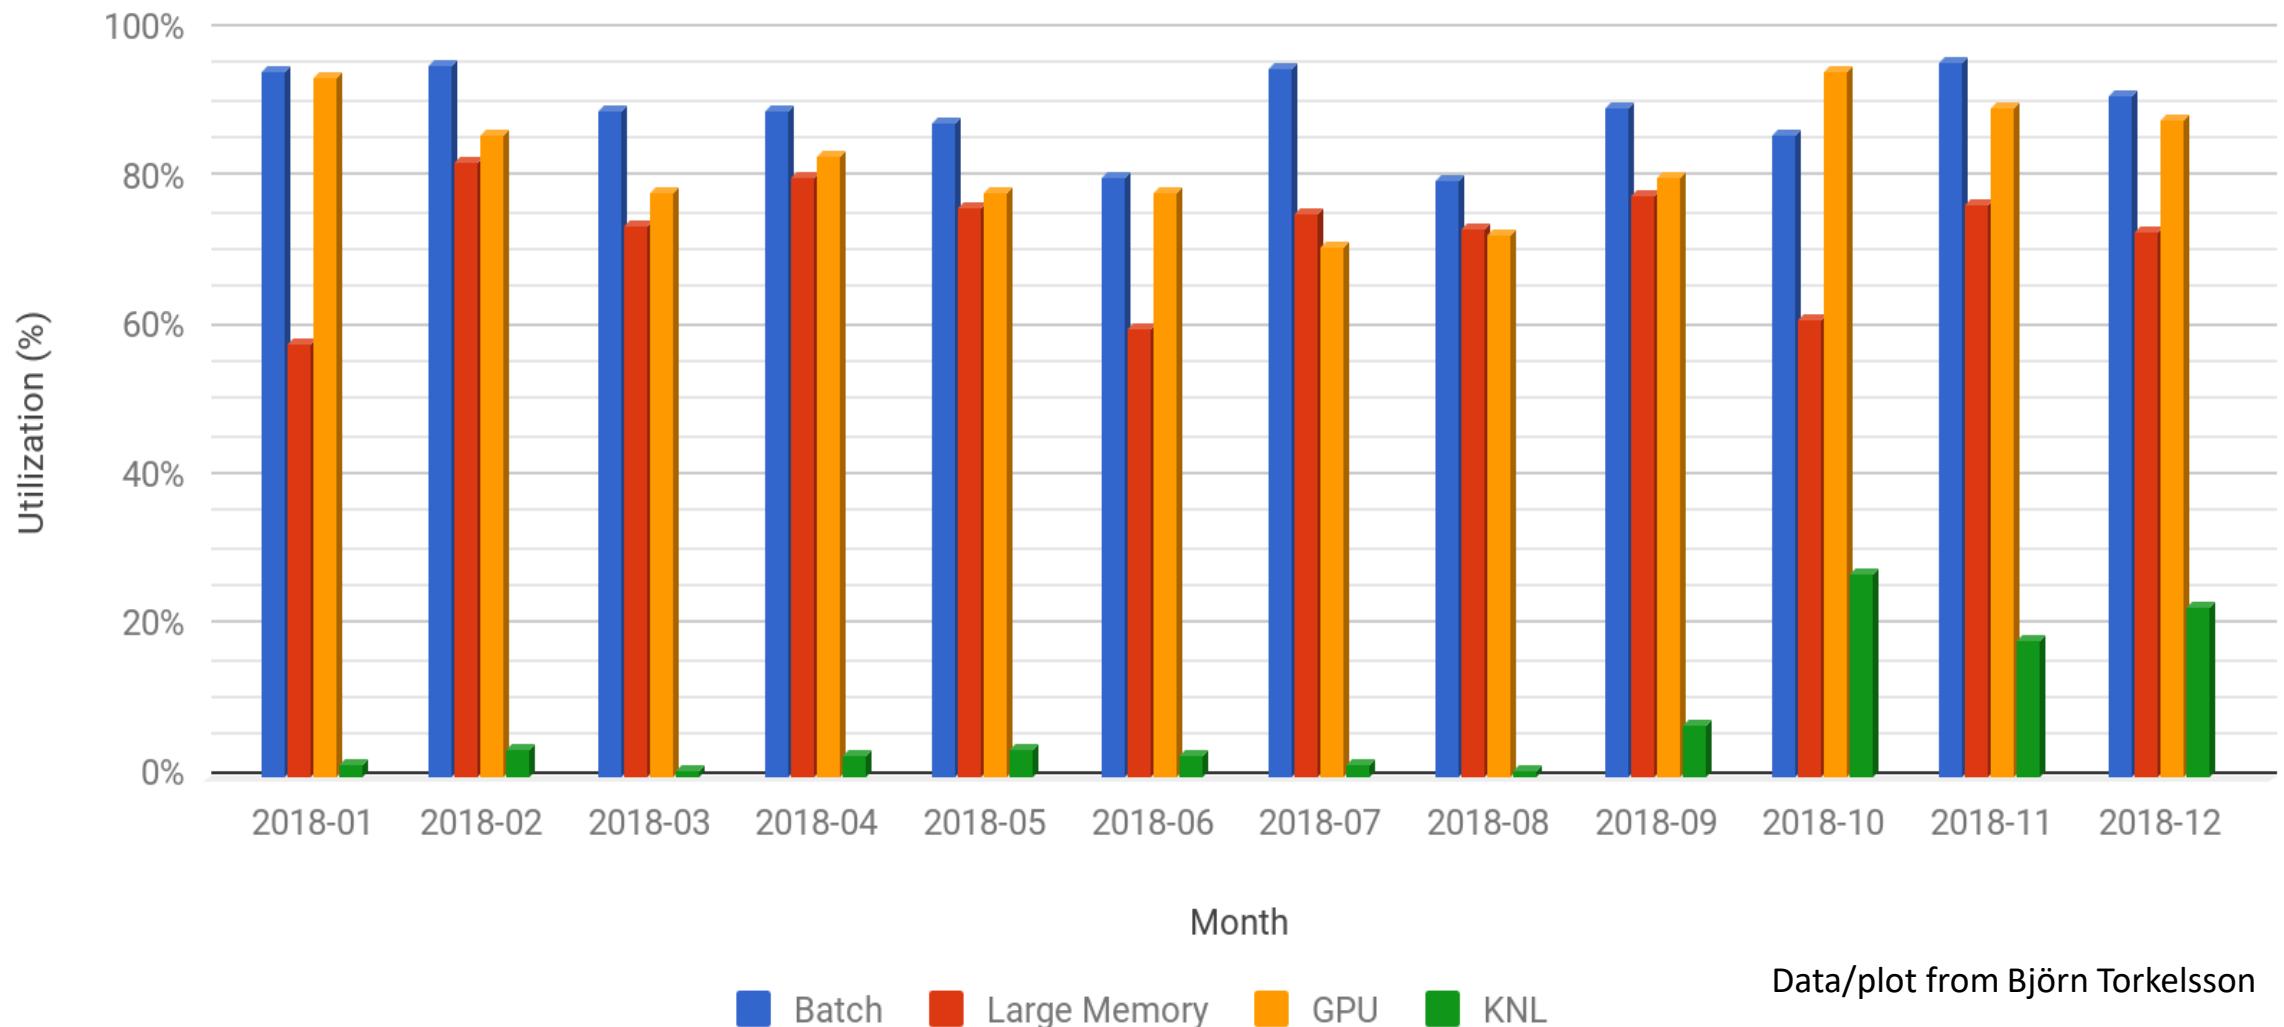
Credits: Intel

KNL Thread affinity

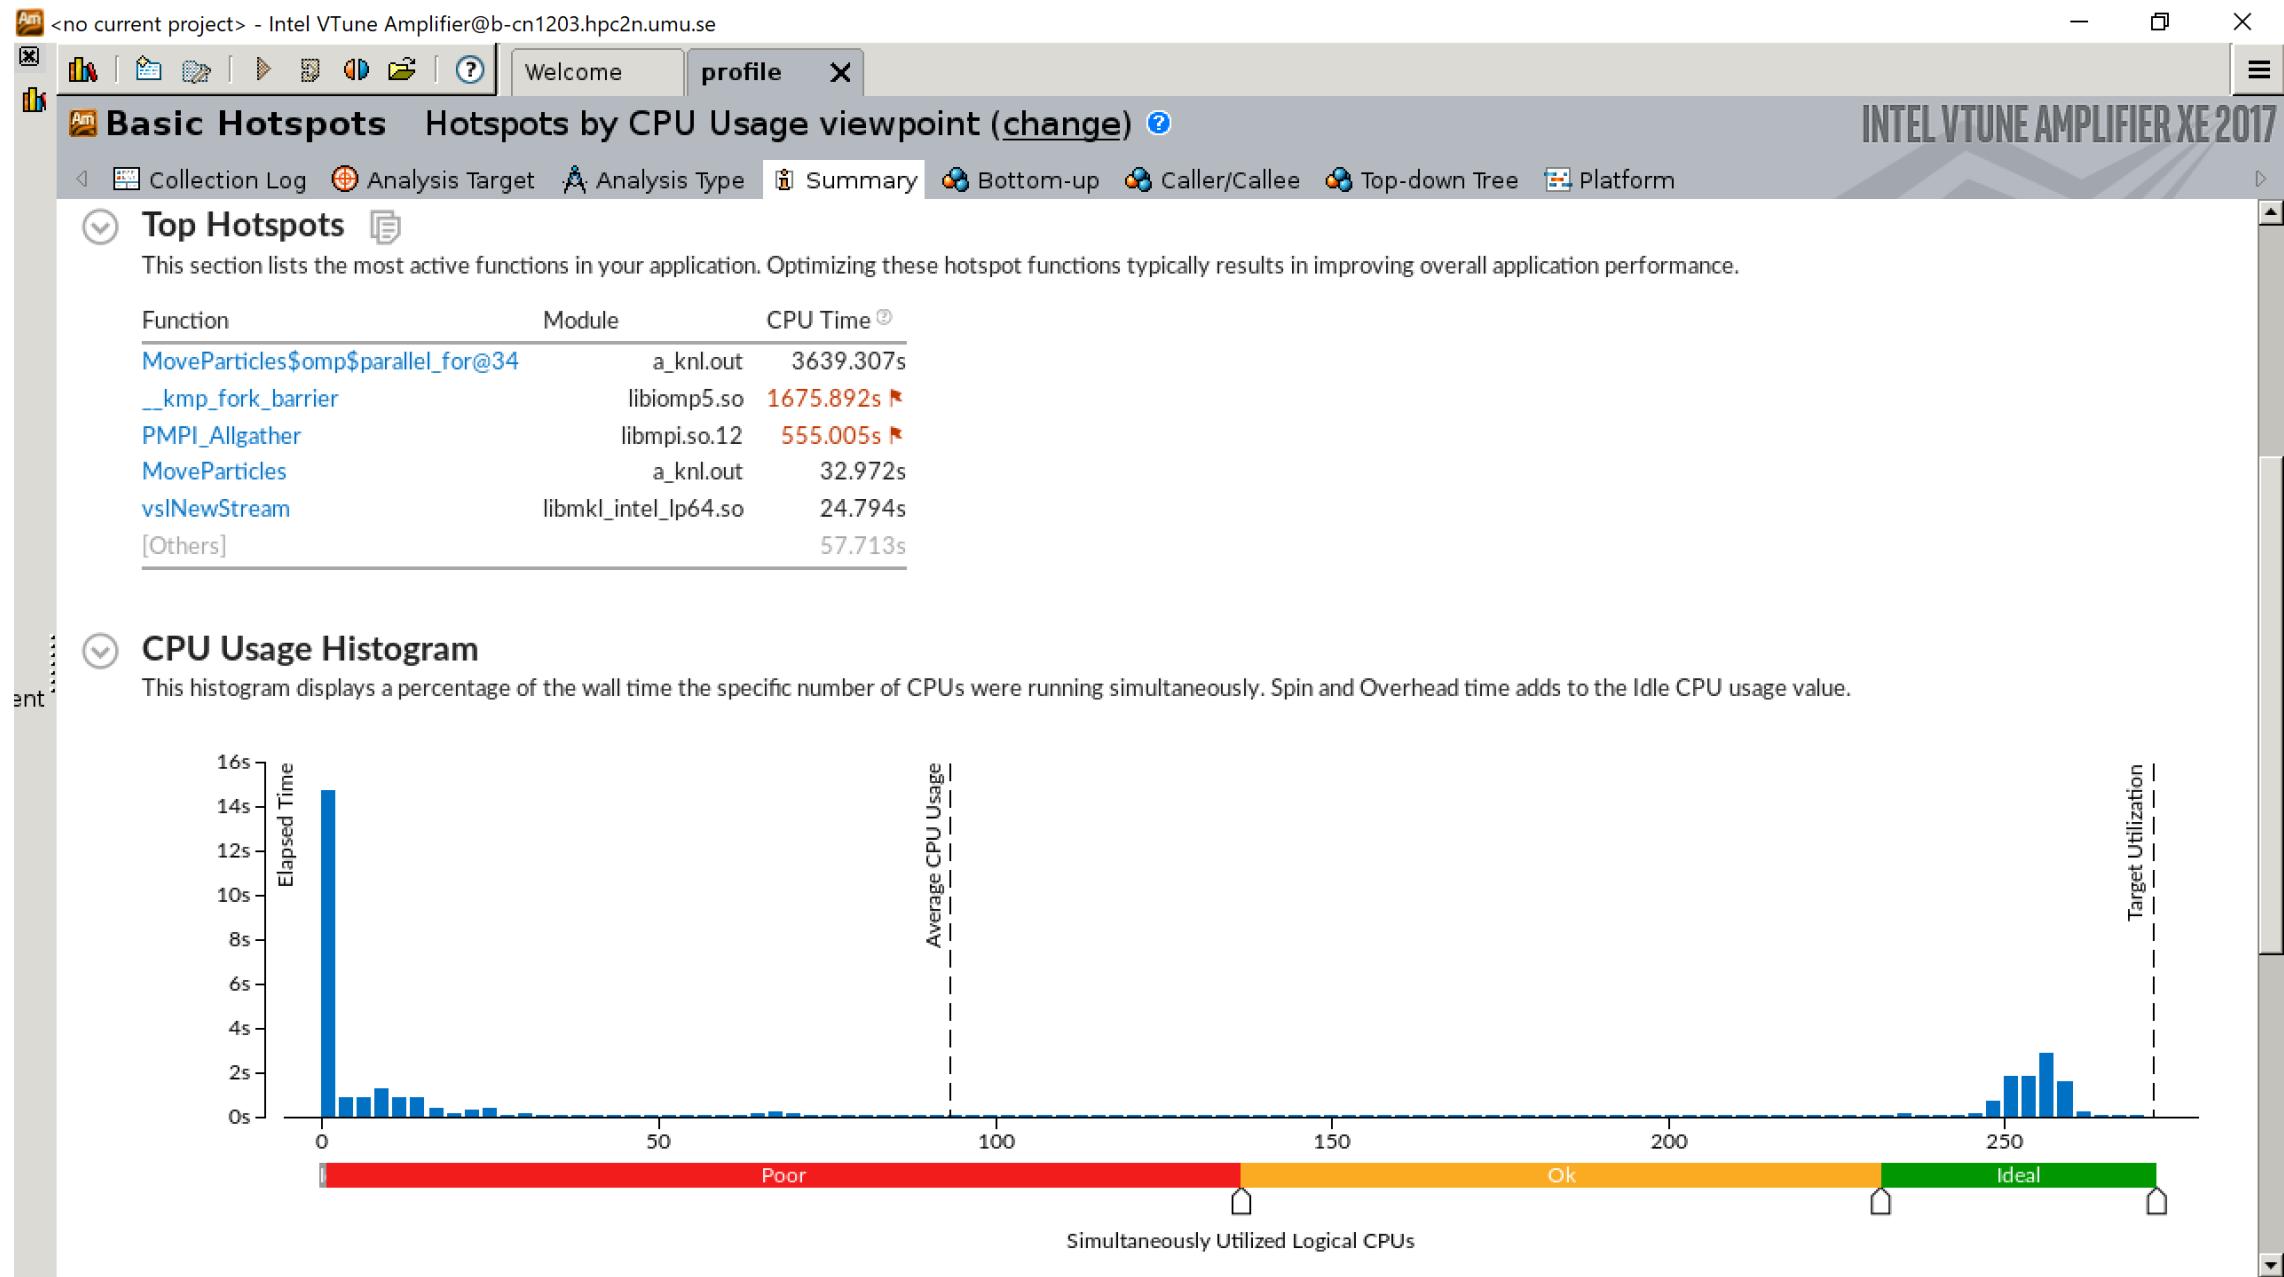
```
#export OMP_PROC_BIND=spread, close, etc.  
#export OMP_PLACES=threads, cores, etc.  
export OMP_NUM_THREADS=4  
srun -n 16 -c 4 --cpu_bind=cores ./xthi
```

Hello from rank 0, thread 0, on b-cn1209.hpc2n.umu.se. (core affinity = 0)
Hello from rank 0, thread 1, on b-cn1209.hpc2n.umu.se. (core affinity = 68)
Hello from rank 0, thread 2, on b-cn1209.hpc2n.umu.se. (core affinity = 136)
Hello from rank 0, thread 3, on b-cn1209.hpc2n.umu.se. (core affinity = 204)
Hello from rank 1, thread 0, on b-cn1209.hpc2n.umu.se. (core affinity = 1)
Hello from rank 1, thread 1, on b-cn1209.hpc2n.umu.se. (core affinity = 69)
Hello from rank 1, thread 2, on b-cn1209.hpc2n.umu.se. (core affinity = 137)
Hello from rank 1, thread 3, on b-cn1209.hpc2n.umu.se. (core affinity = 205)

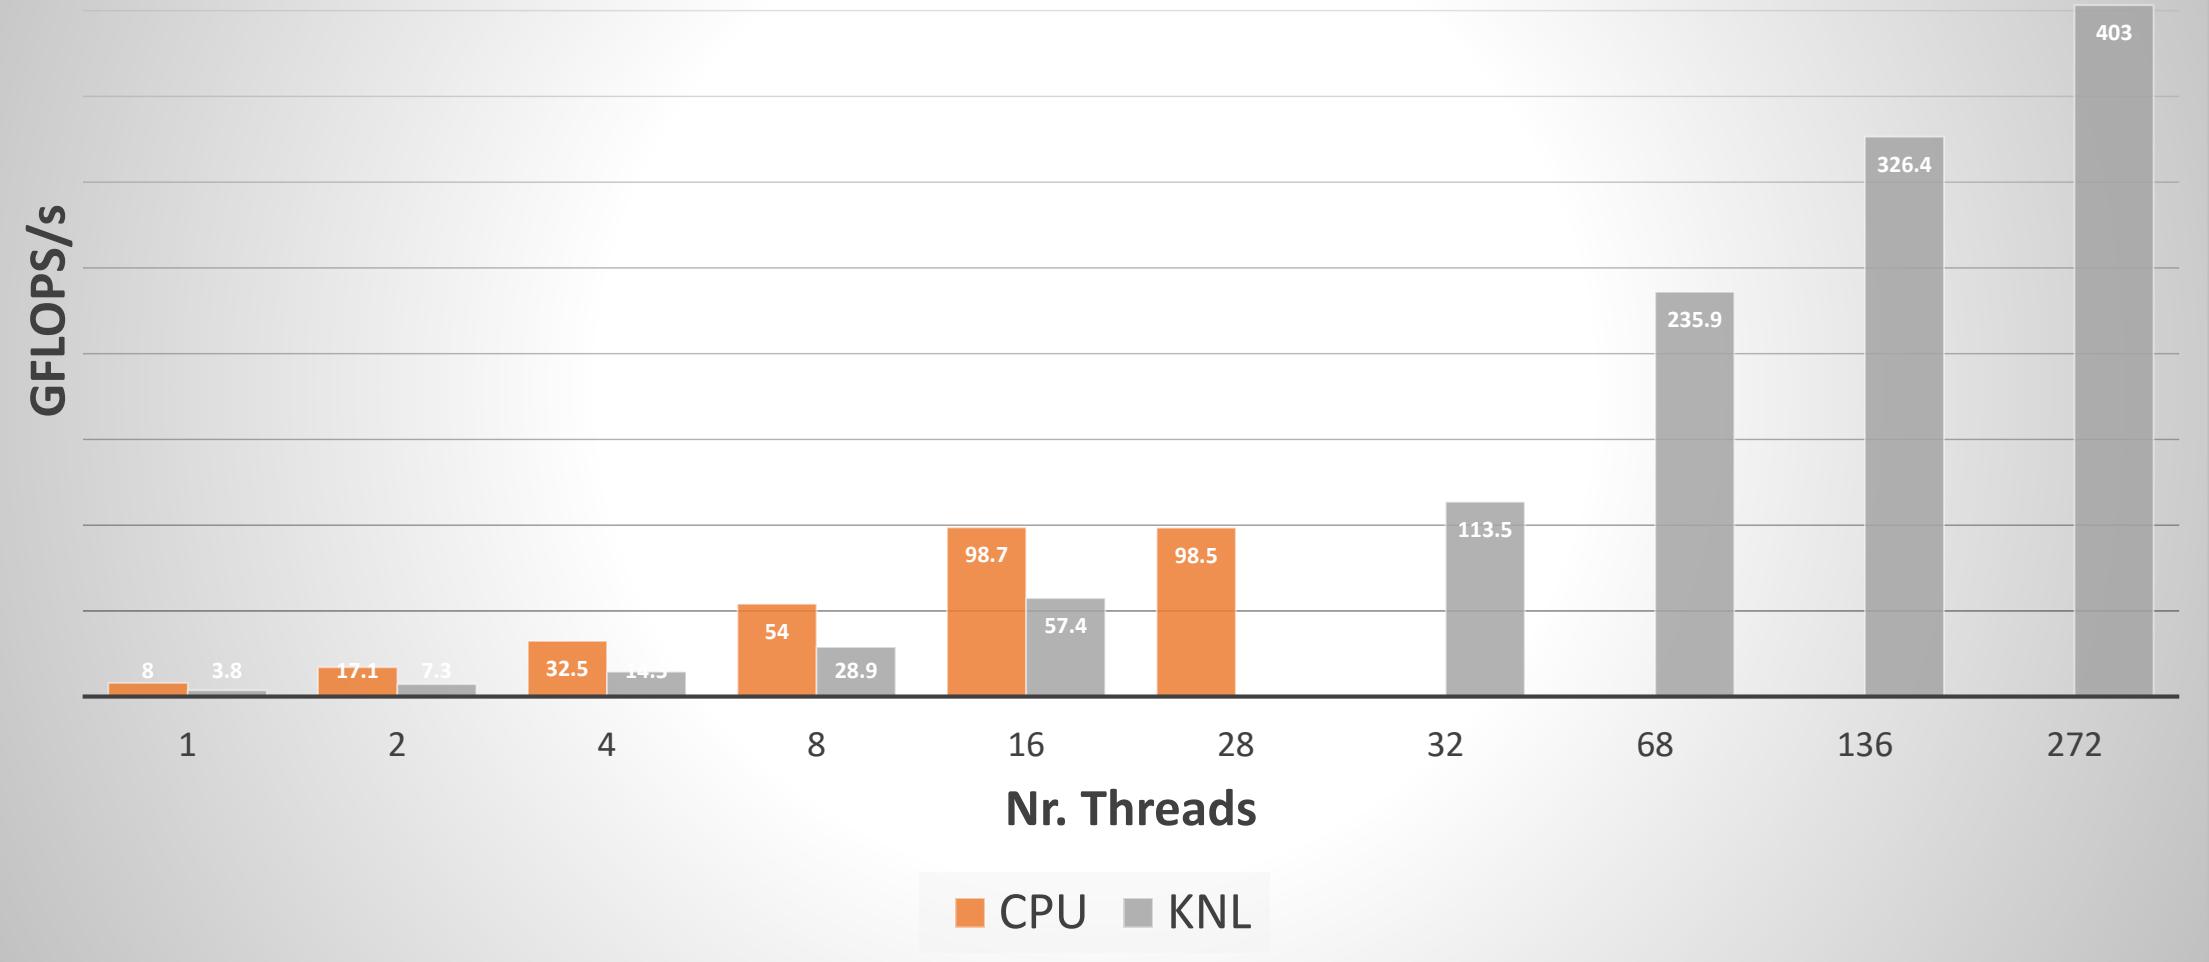
Kebnekaise: Utilization 2018 (per partition)

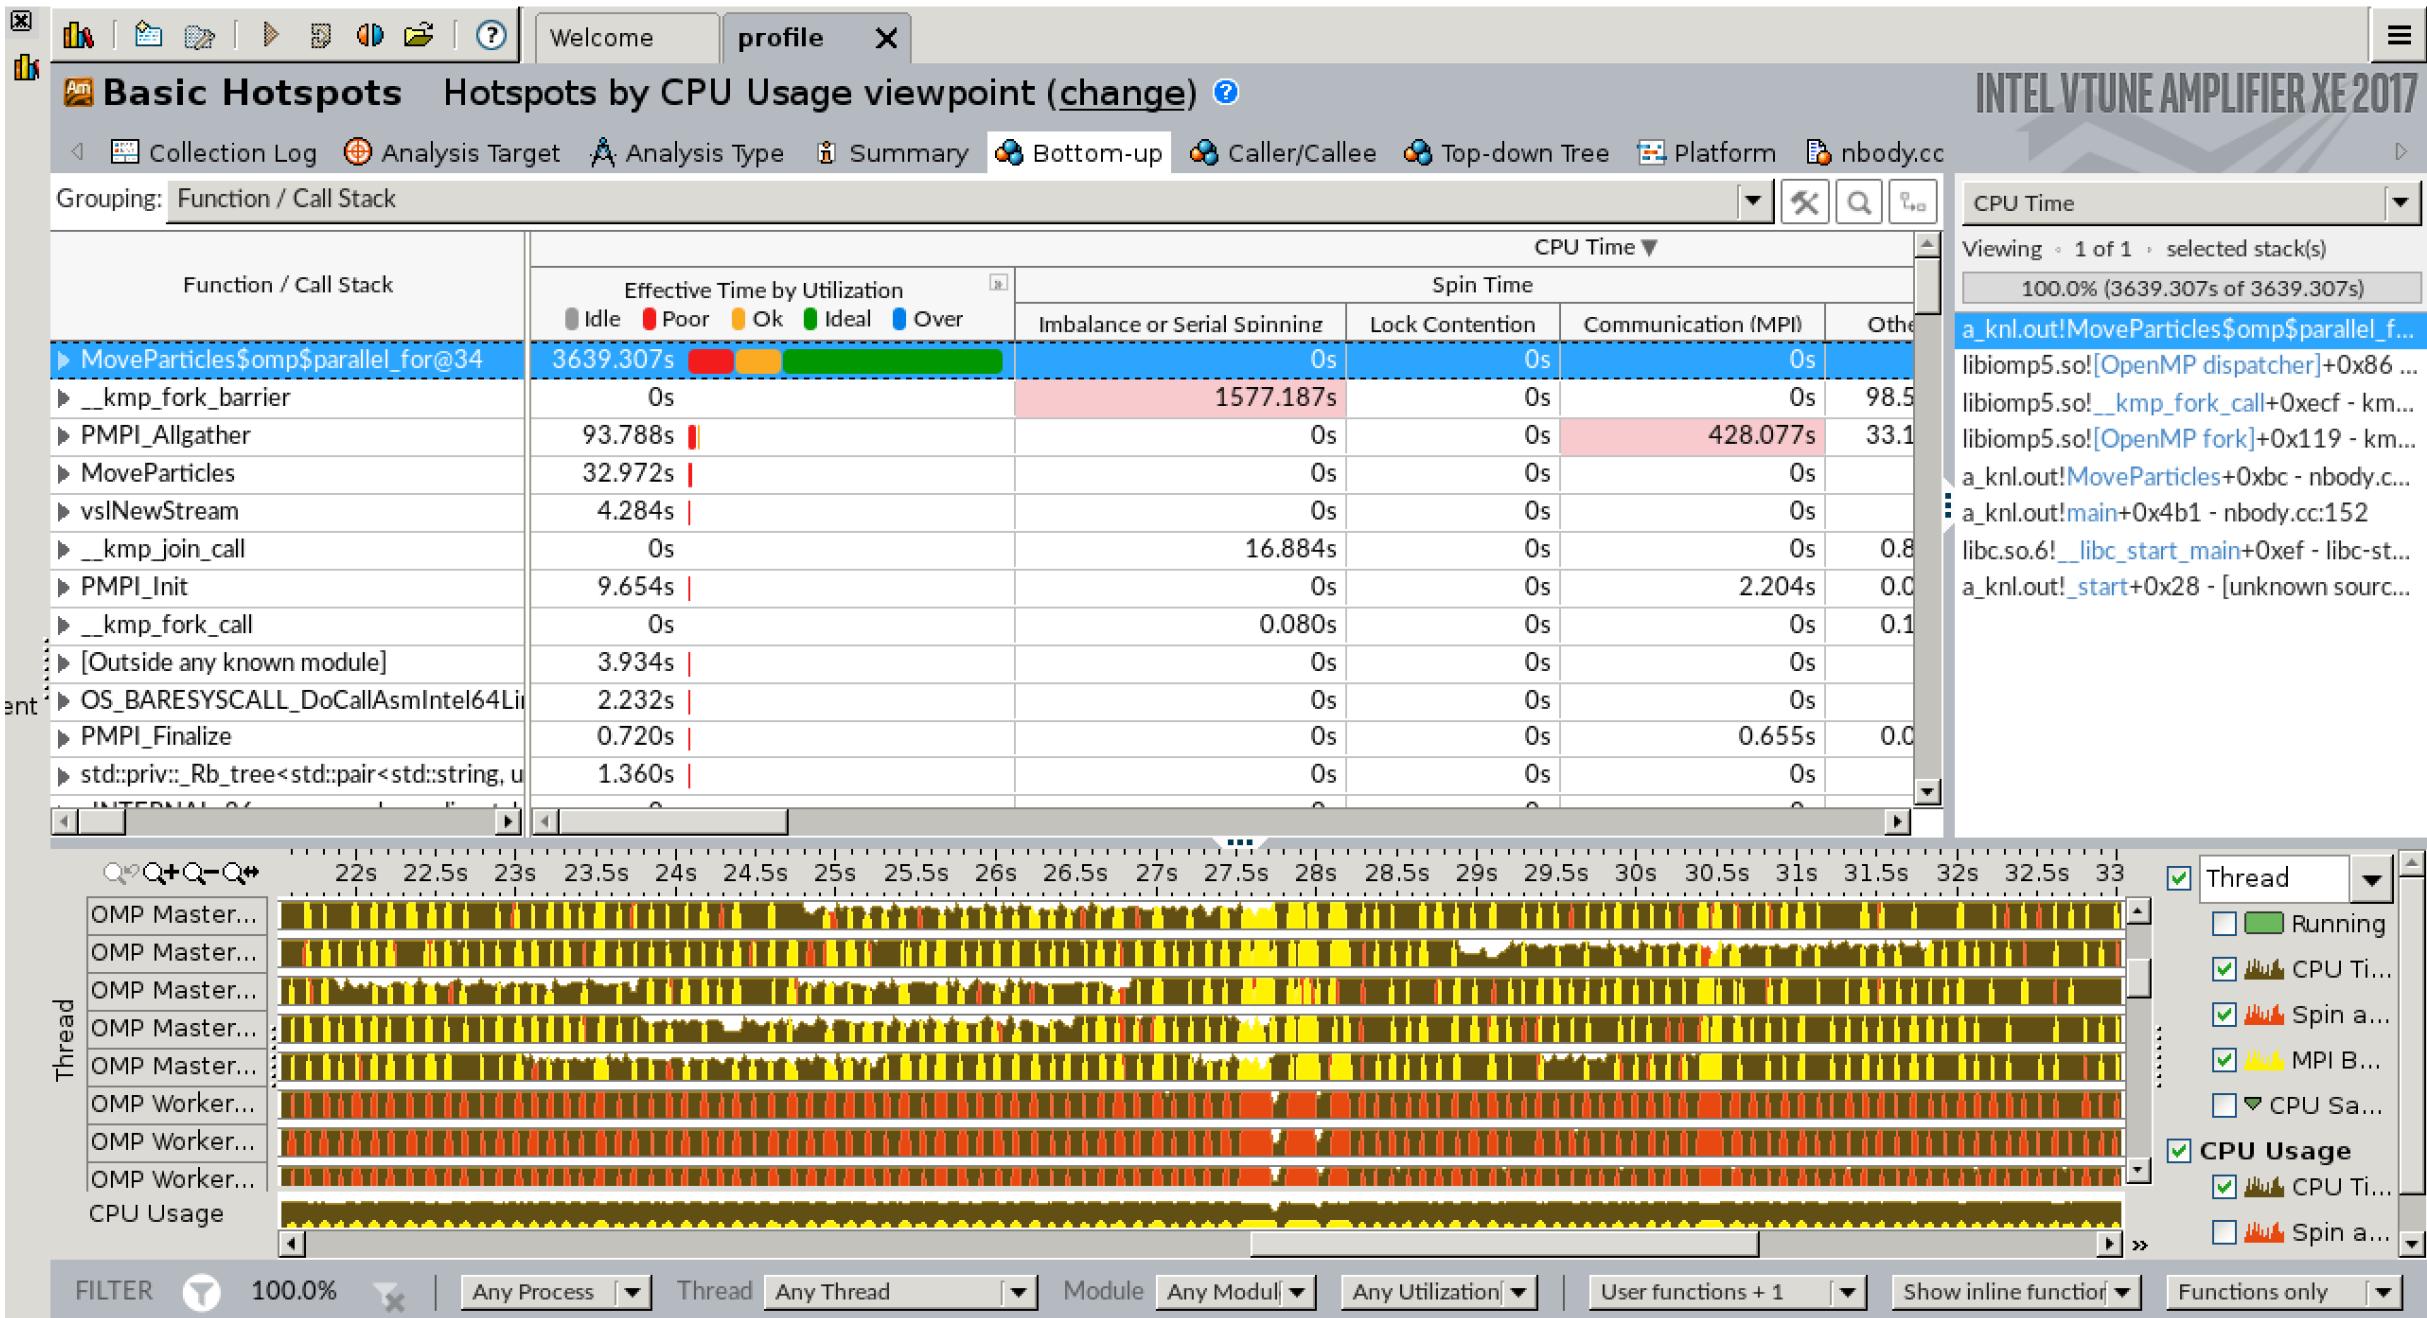


Tuning your own Applications



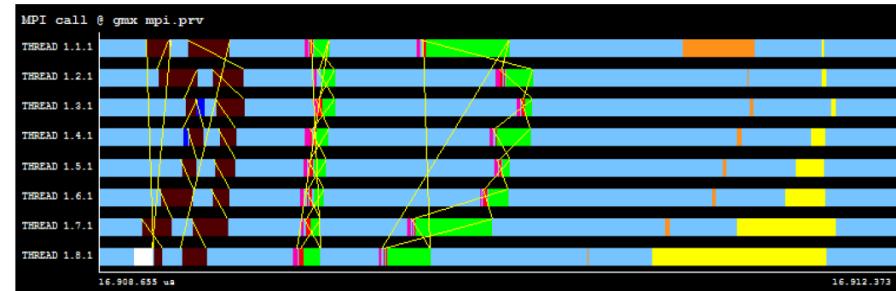
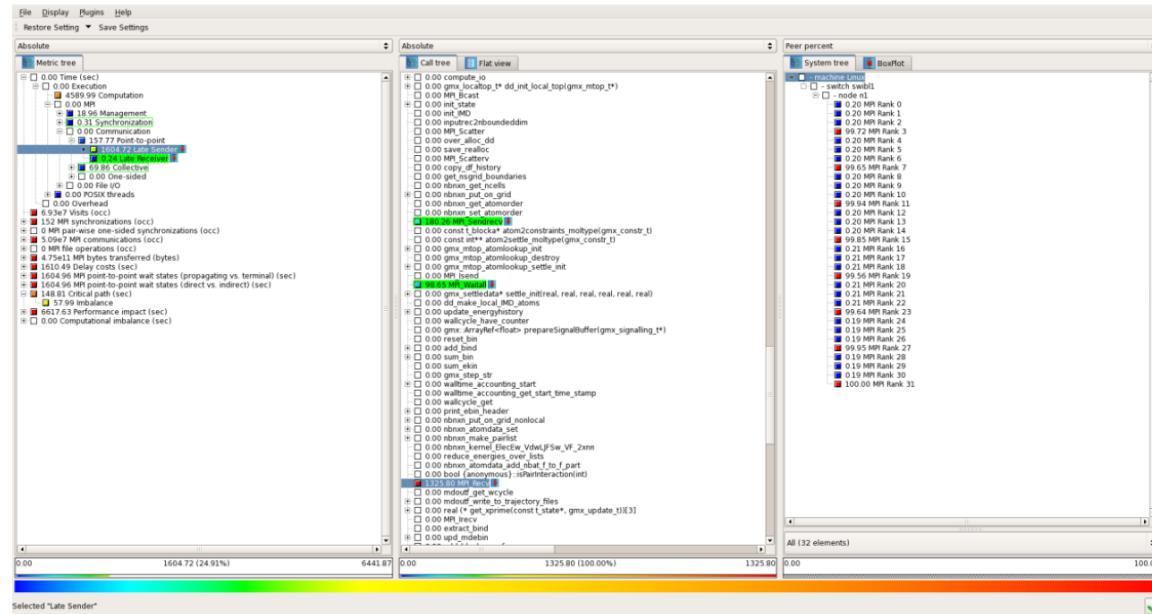
N-body MD





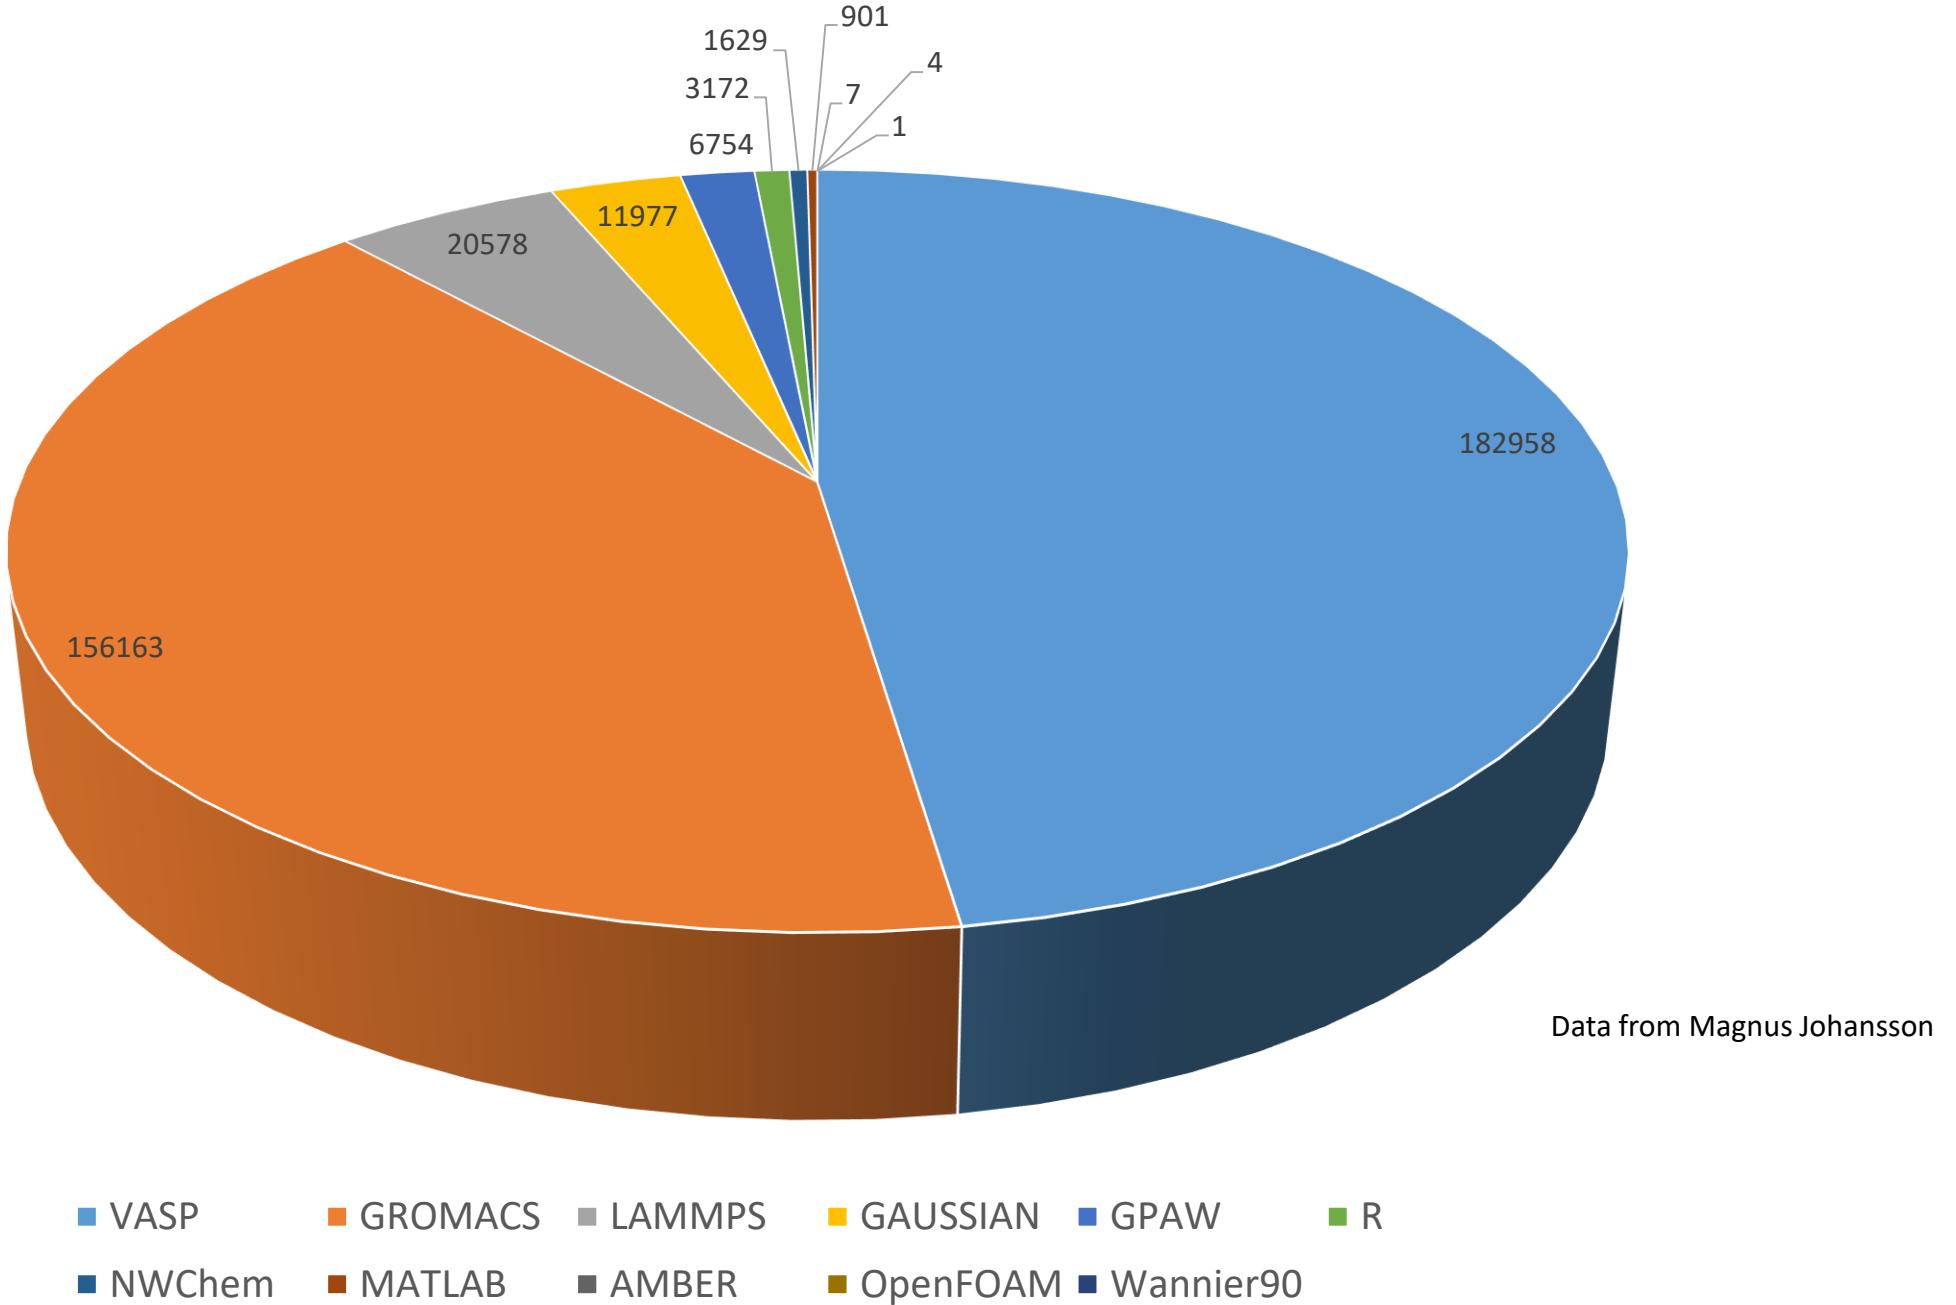
Profiling Tools at HPC2N

- Intel Vtune, Intel Inspector
- SCALASCA, Score-P, Cube
- Extrae/Paraver
- Allinea DDT
- Application Expert support on demand

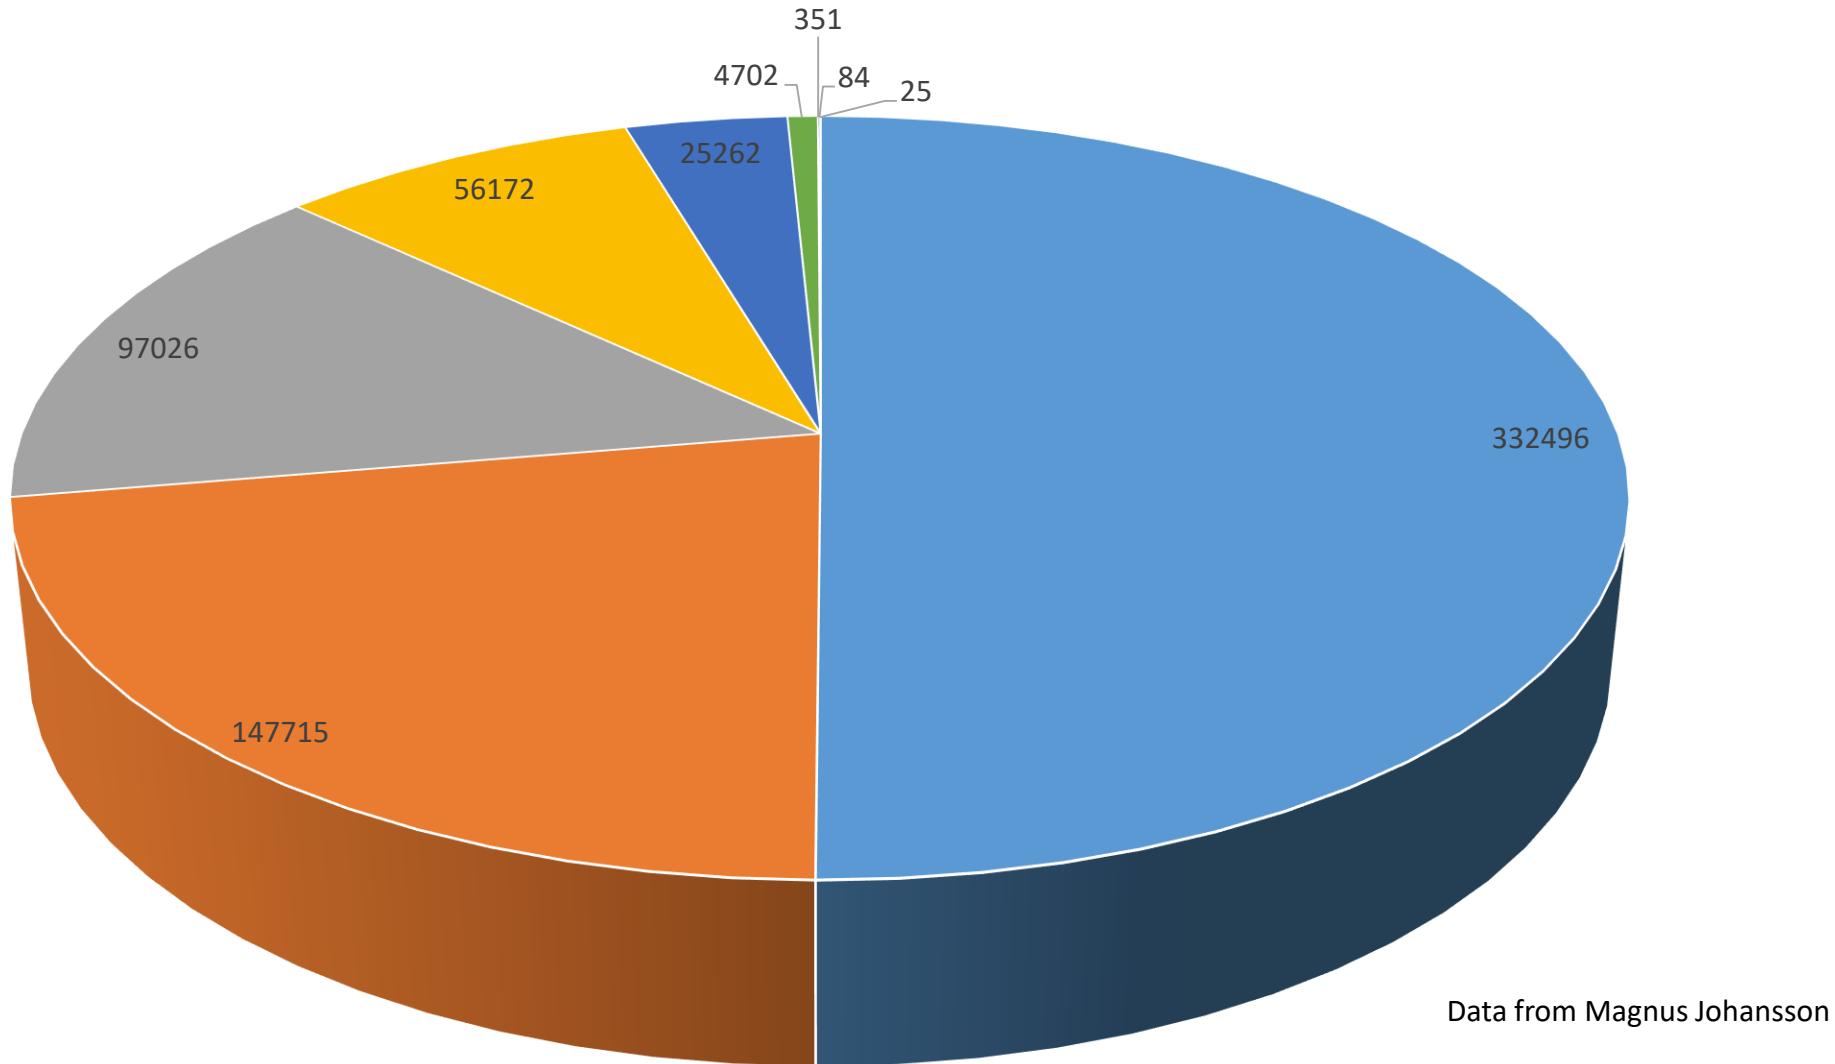


Applications

Software on Kebnekaise



Software on Abisko



■ Singularity ■ GROMACS ■ VASP ■ GAUSSIAN ■ SIESTA ■ AMBER ■ WRF ■ MATLAB ■ NWChem

R/Rstudio

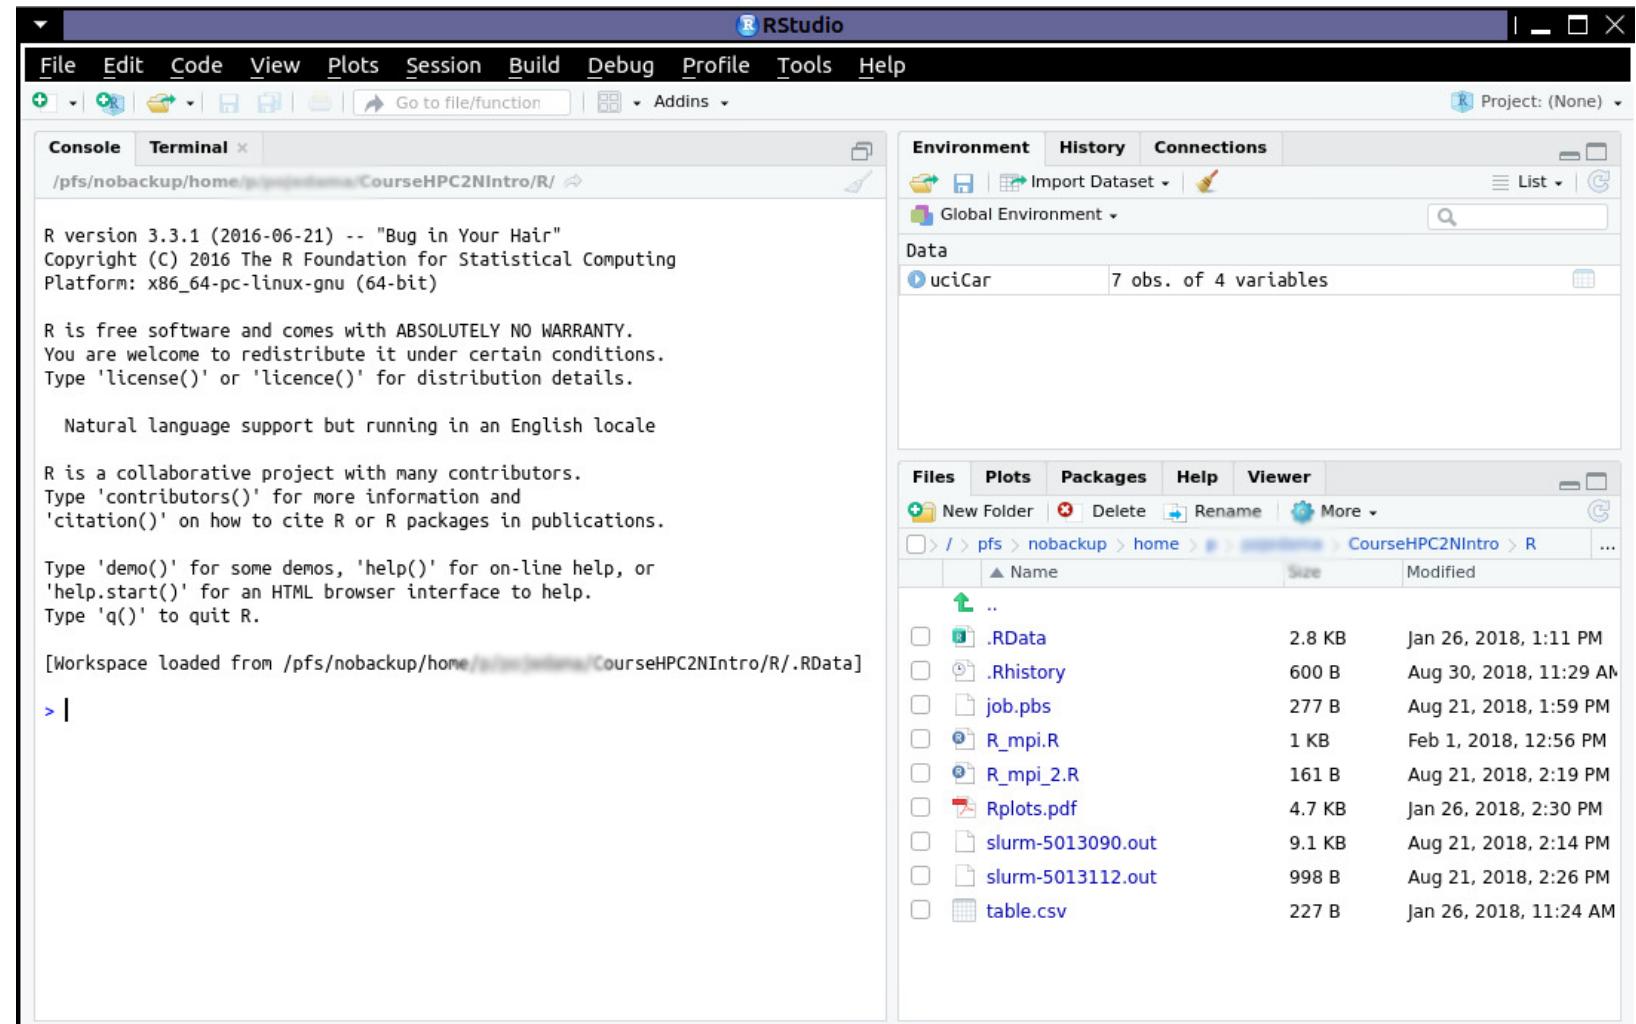
\$ml icc/2017.1.132-GCC-6.3.0-2.27
impi/2017.1.132 ifort/2017.1.132-GCC-
6.3.0-2.27

\$ml R/3.3.1
\$rstudio

Rstudio will use the loaded R version

Please use the GUI just for
lightweight tasks, otherwise
use the batch queues.

Heavy jobs will be cancelled!



R/Rstudio

Installing packages

In case you need to install some R package that is not included in the installed R version, you need to follow this guideline:

https://www.hpc2n.umu.se/resources/software/user_installed/r

Otherwise, the batch system won't have access to the manually installed packages.

R serial job

```
#!/bin/bash
### SNAC project number, enter your own
#SBATCH -A SNICXXXX-YY-ZZ
# Asking for one core
#SBATCH -n 1
#SBATCH --time=00:10:00

ml icc/2017.1.132-GCC-6.3.0-2.27
ml ifort/2017.1.132-GCC-6.3.0-2.27
ml impi/2017.1.132
ml R/3.3.1

R --no-save --quiet < input.R > Rexample.out
```

Further information:
<https://www.hpc2n.umu.se/resources/software/r>

R parallel job (Rmpi)

File: Rmpi.R

```
library("Rmpi")
mpi.spawn.Rslaves(nslaves=5)

x <- c(10,20,30,40,50)
mpi.apply(x,runif)

# Close down the MPI processes and quit R
mpi.close.Rslaves()
mpi.finalize()
```

Further information:

<https://www.hpc2n.umu.se/resources/software/r>

```
#!/bin/bash
#SBATCH -A project_ID
#Asking for 10 min.
#SBATCH -t 00:10:00
#SBATCH -N 1
#SBATCH -n 6

export OMPI_MCA_mpi_warn_on_fork=0

ml GCC/6.4.0-2.28 OpenMPI/2.1.2
ml R/3.4.4-X11-20180131

R -q --slave -f Rmpi.R
```

R parallel job (doParallel)

```
#!/bin/bash
#SBATCH -A project_ID
#Asking for 10 min.
#SBATCH -t 00:10:00
##SBATCH -N 1
#SBATCH -n 8
##SBATCH --exclusive

ml GCC/6.4.0-2.28 OpenMPI/2.1.2
ml R/3.4.4-X11-20180131
```

```
R -q --slave -f doParallel_ML.R
```

doParallel_ML.R

```
library(doParallel)
registerDoParallel(cores=8)
getDoParWorkers()
```

```
library(caret)
library(MASS)
library(klaR)
library(nnet)
library(e1071)
library(rpart)
```

Further information:
<https://www.hpc2n.umu.se/resources/software/r>

R parallel job (doParallel)

doParallel_ML.R

```
#!/bin/bash
#SBATCH -A project_ID
#Asking for 10 min.
#SBATCH -t 00:10:00
##SBATCH -N 1
#SBATCH -n 8
##SBATCH --exclusive

ml GCC/6.4.0-2.28 OpenMPI/2.1.2
ml R/3.4.4-X11-20180131

R -q --slave -f doParallel_ML.R
```

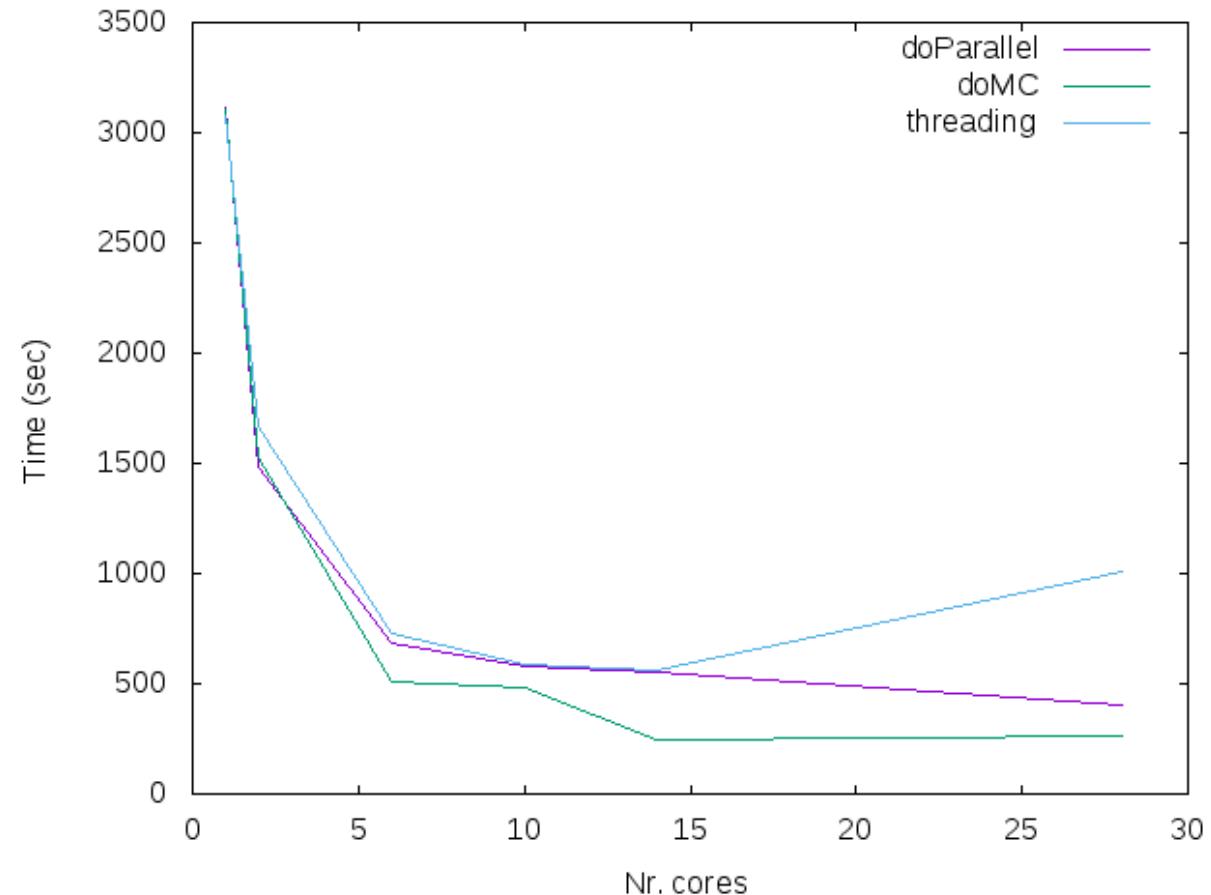
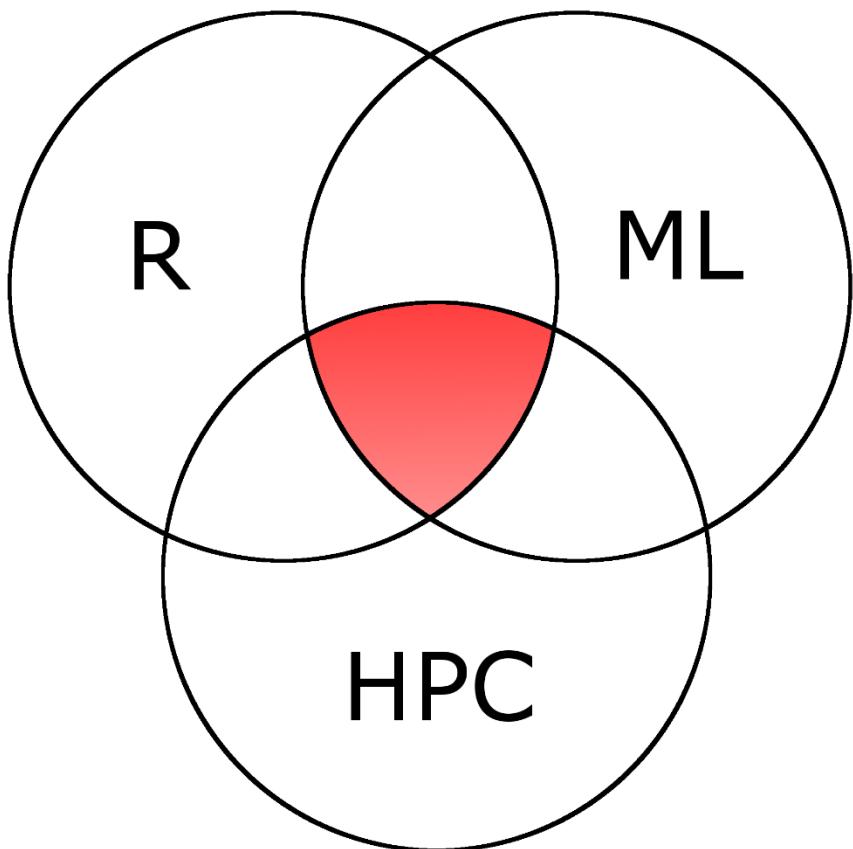
```
#Serial mode
sr <- foreach(1:runs, .combine = rbind) %do% evaluation()

#Parallel mode
pr <- foreach(1:runs, .combine = rbind) %dopar% evaluation()
```

Further information:
<https://www.hpc2n.umu.se/resources/software/r>

R machine learning

Example: Boosted classification model using “xgboost” (7 tuning parameters). Using R package on Kebnekaise (details: <http://appliedpredictivemodeling.com/blog/2018/1/17/parallel-processing>).

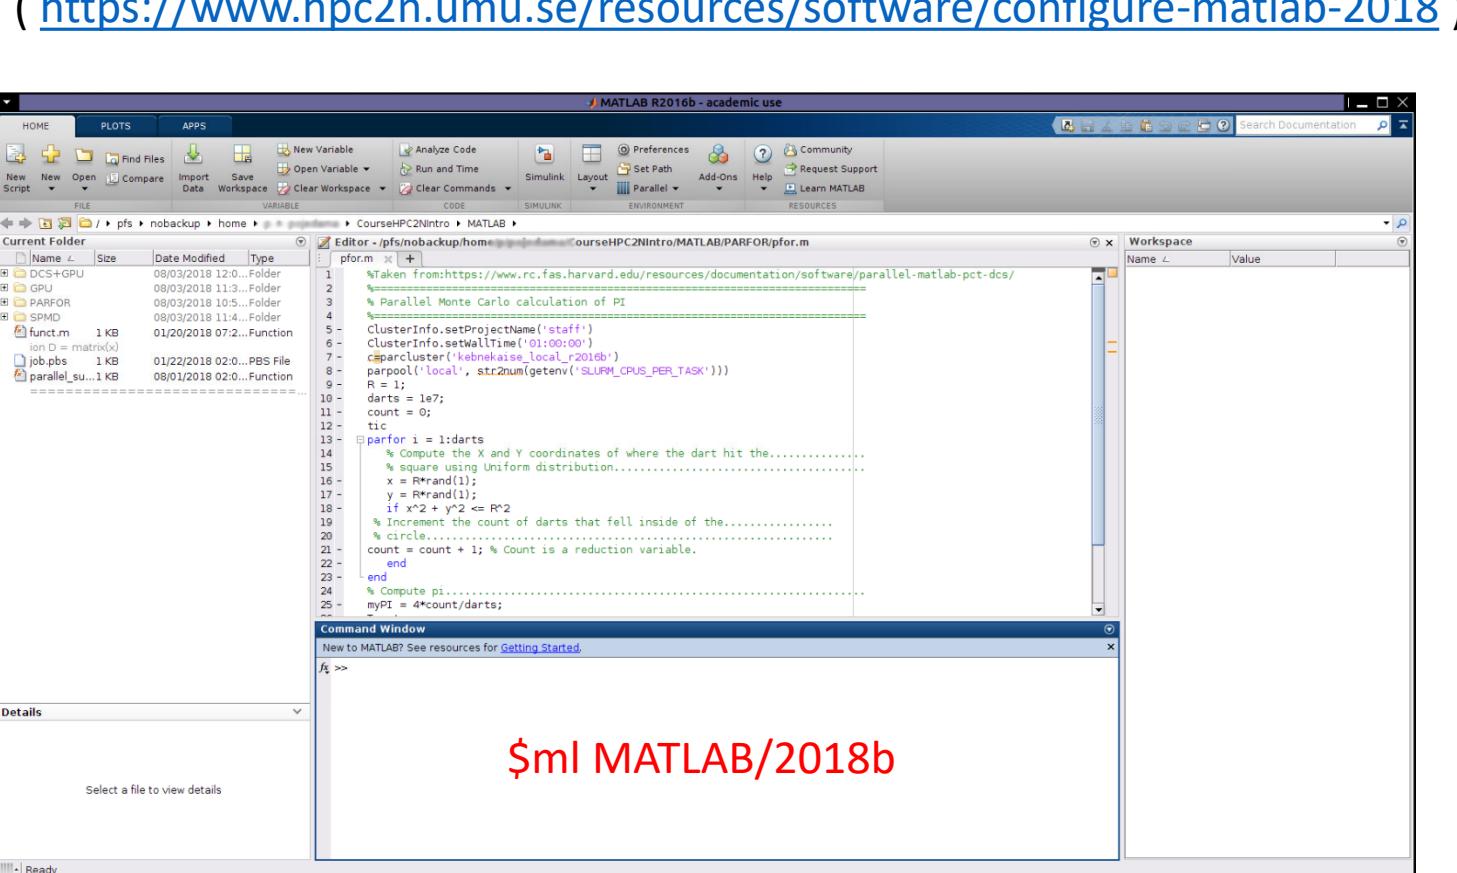


MATLAB

One needs to configure MATLAB the first time one uses it:

*MATLAB 2018b and later

(<https://www.hpc2n.umu.se/resources/software/matlab>)



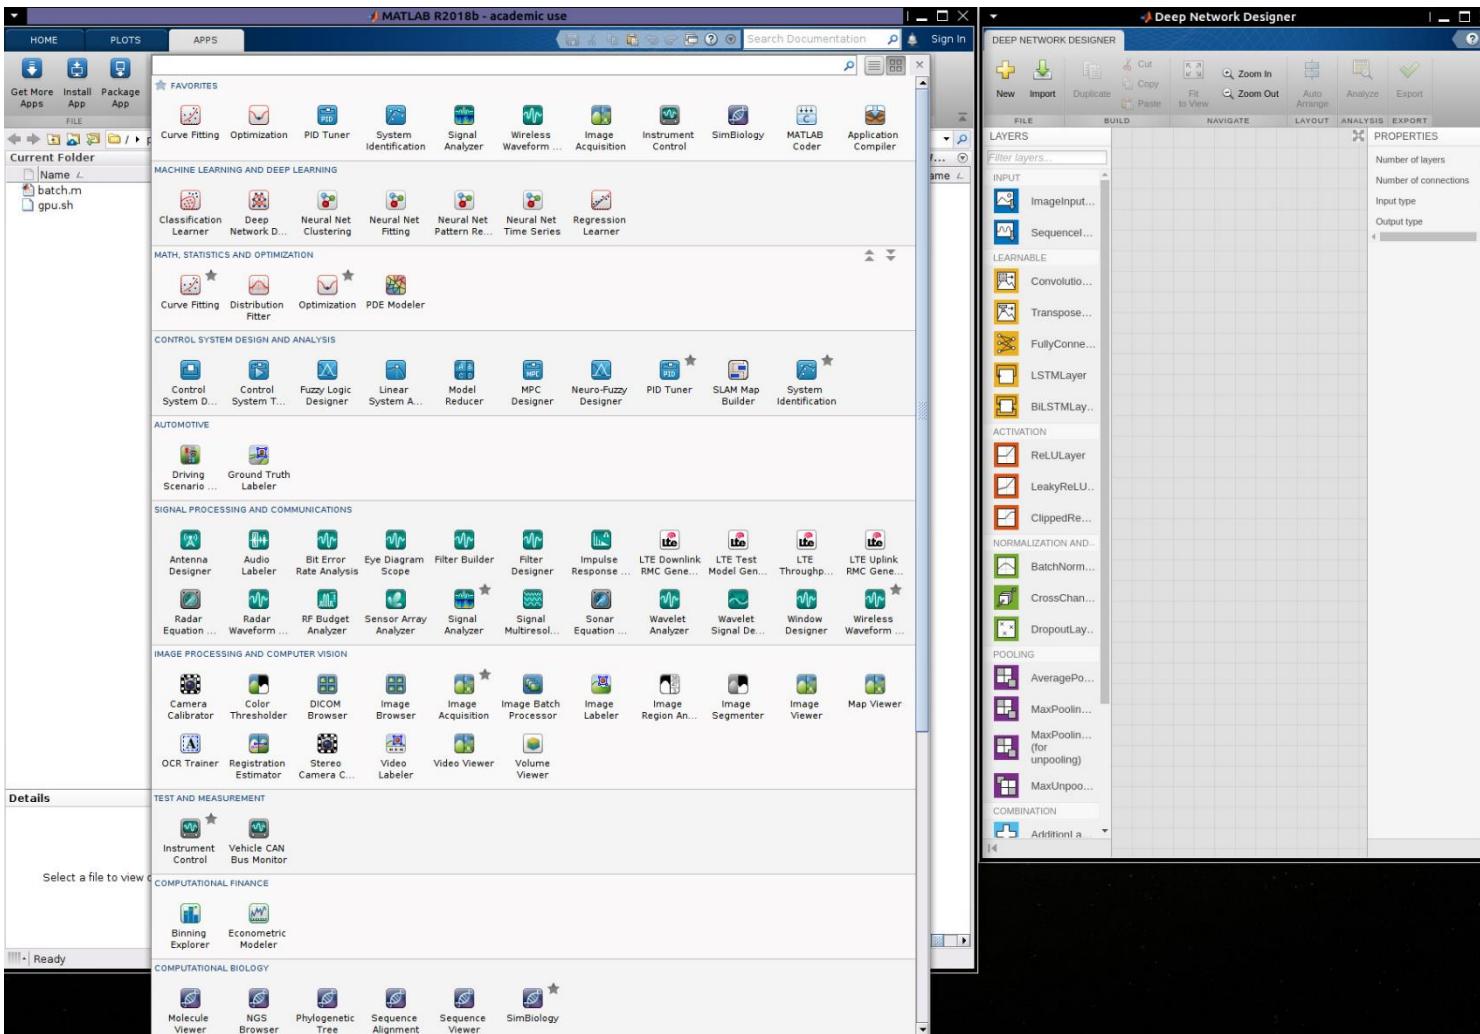
Further information: <https://www.hpc2n.umu.se/resources/software/matlab>

Please use the GUI just for lightweight tasks, otherwise use the batch queues.

Heavy jobs will be cancelled!

Start Matlab with:
matlab -singleCompThread

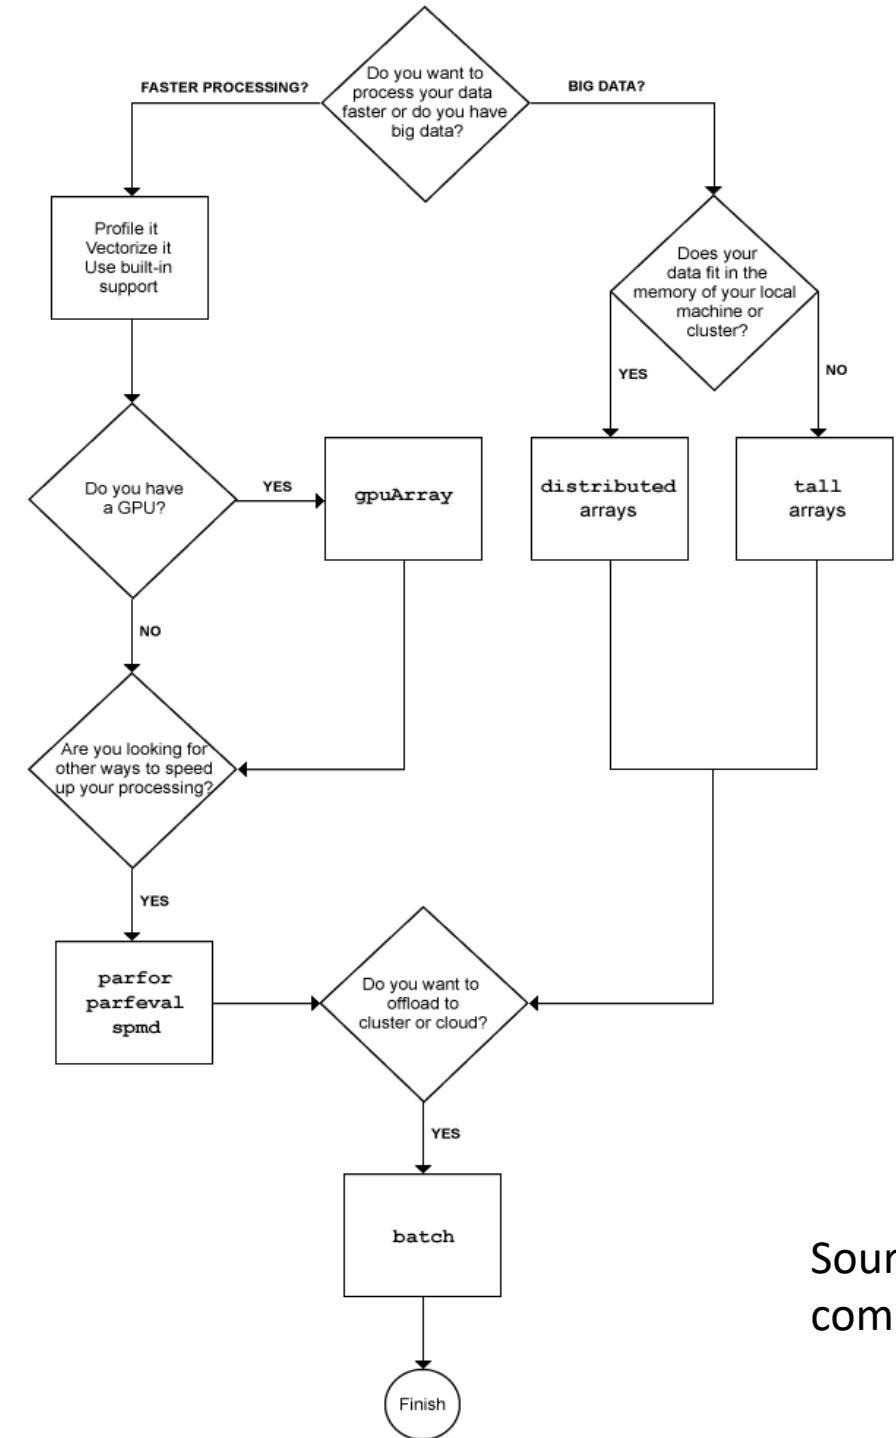
MATLAB



\$ml MATLAB/2018b

Includes improved tools for Machine Learning and Deep Learning

Further information: <https://www.hpc2n.umu.se/resources/software/matlab>



Parallel computing tools in MATLAB

Source: <https://se.mathworks.com/help/parallel-computing/choosing-a-parallel-computing-solution.html>

MATLAB

- Parallel Computing Toolbox (PCT)
 - Parfor loops
 - Single Program Multiple Data, Labs are used to solve a task
- Distributed Computing Server (DCS)

Serial job

File name: funct.m

```
function D  
  
format long  
A1 = rand(10000,10000);  
tic;  
B1 = fft(A1);  
time1 = toc;
```

```
filename = 'log.out';  
mid=fopen(filename,'w');  
fprintf(mid,'Time = %16.12f\n',time1);  
fclose(mid);
```

```
#!/bin/bash  
#SBATCH -A project_ID  
#Asking for 10 min.  
#SBATCH -t 00:10:00  
#SBATCH -n 1  
#SBATCH -p batch  
#Load modules necessary for running MATLAB  
ml MATLAB/2019b.Update2
```

srun matlab -nodisplay -singleCompThread -r
"funct;exit" > out.log

exit 0

Parfor

File name: pfor.m

```
parpool('local', str2num(getenv('SLURM_CPUS_PER_TASK')))
```

```
parfor i = 1:darts
```

```
x = rand(1);
```

```
y = rand(1);
```

```
if x^2 + y^2 <= R^2
```

```
count = count + 1; % reduction variable.
```

```
end
```

```
end
```

Batch script:

```
#SBATCH -c 4  
#SBATCH -p batch
```

```
#Load modules necessary for running MATLAB  
ml MATLAB/2019b.Update2
```

```
srun matlab –nodisplay -r "pfor"
```

Timing	
Nr. Workers	Time(sec)
1	26.43
2	17.17
3	11.82
4	7.71
28	1.82

SPMD

File: spmdex.m

```
parpool('local', str2num(getenv('SLURM_CPUS_PER_TASK')))
```

```
% The expression that we wish to integrate:
```

```
f = @(x) 4./(1 + x.^2);
```

```
spmd
```

```
% Choose a range over which to integrate based on my unique index:
```

```
range_start = (labindex - 1) / numlabs;
```

```
range_end = labindex / numlabs;
```

```
% Calculate my portion of the overall integral
```

```
my_integral = quadl( f, range_start, range_end );
```

```
% Aggregate the result by adding together each value of ?my_integral?
```

```
total_integral = gplus( my_integral );
```

```
end
```

```
total_int = total_integral{1};
```

```
fprintf('The computed value of pi is %8.7f.\n',total_int);
```

```
delete(gcp);
```

```
#!/bin/bash
```

```
#SBATCH -A staff
```

```
#Asking for 10 min.
```

```
#SBATCH -t 00:10:00
```

```
#SBATCH -c 4
```

```
#SBATCH -p batch
```

```
#Load modules necessary for MATLAB  
ml MATLAB/2019b.Update2
```

```
srun matlab -nodisplay -r "spmdex" > out.log
```

MATLAB+GPU

On the MATLAB GUI:

```
c = parcluster;
ClusterInfo.setName('staff');
ClusterInfo.setWallTime('00:10:00');
% Tell the scheduler that you want to use the GPUs
ClusterInfo.setUseGpu(true)
% number of GPUs per node to use
ClusterInfo.setGpusPerNode(2)
ClusterInfo.setUserDefinedOptions('--gres=gpu:k80:2,mps')
%gpuDevice

j = c.batch(@parallelex, 1, {1}, 'pool', 2);
```

File: parallelex.m

```
function speedUp = parallel(x)
format long
A1 = rand(10000,10000);
tic;
B1 = fft(A1);
time1 = toc;

A2 = gpuArray(A1);
tic;
B2 = fft(A2);
time2 = toc;

speedUp = time1/time2;
end
```

MATLAB+GPU

Using a batch file: gpu.sh

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 00:10:00
#SBATCH -N 1
#SBATCH -n 1
#SBATCH --gres=gpu:k80:2
#SBATCH -p batch
#SBATCH --exclusive
#Load modules necessary for running MATLAB
ml MATLAB/2019b.Update2
srun matlab -nodisplay -singleCompThread -r "parallelex" > test_parallel.log
```

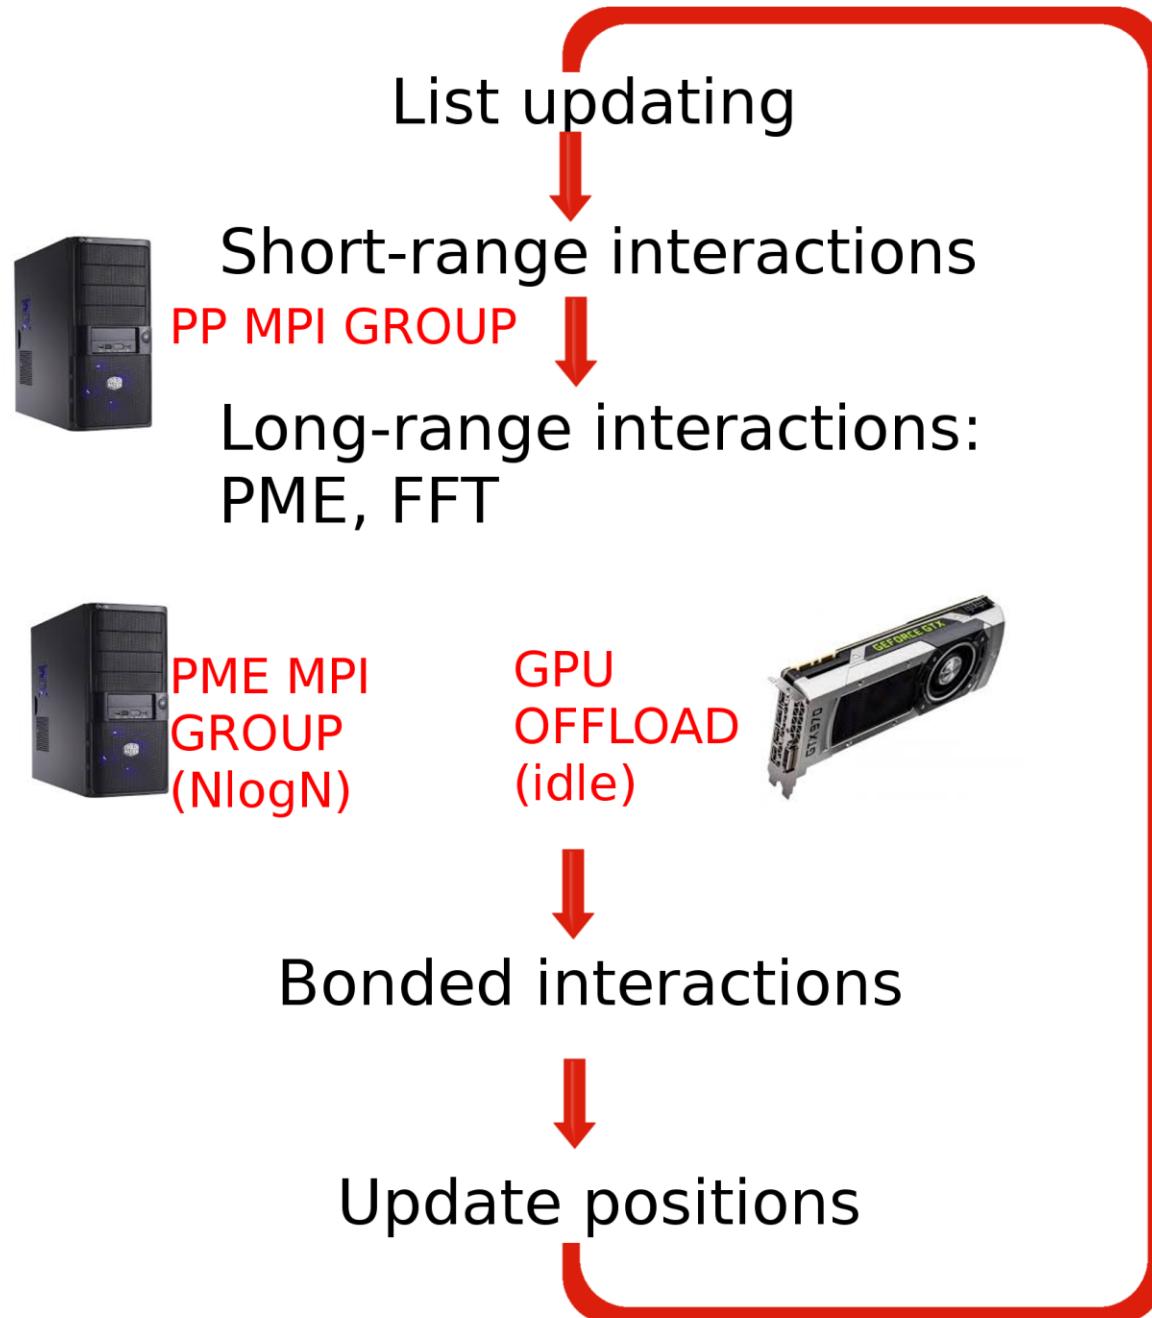
File: parallelex.m

```
function speedUp = parallel(x)
format long
A1 = rand(10000,10000);
tic;
B1 = fft(A1);
time1 = toc;

A2 = gpuArray(A1);
tic;
B2 = fft(A2);
time2 = toc;

speedUp = time1/time2;
end
```

MD Codes



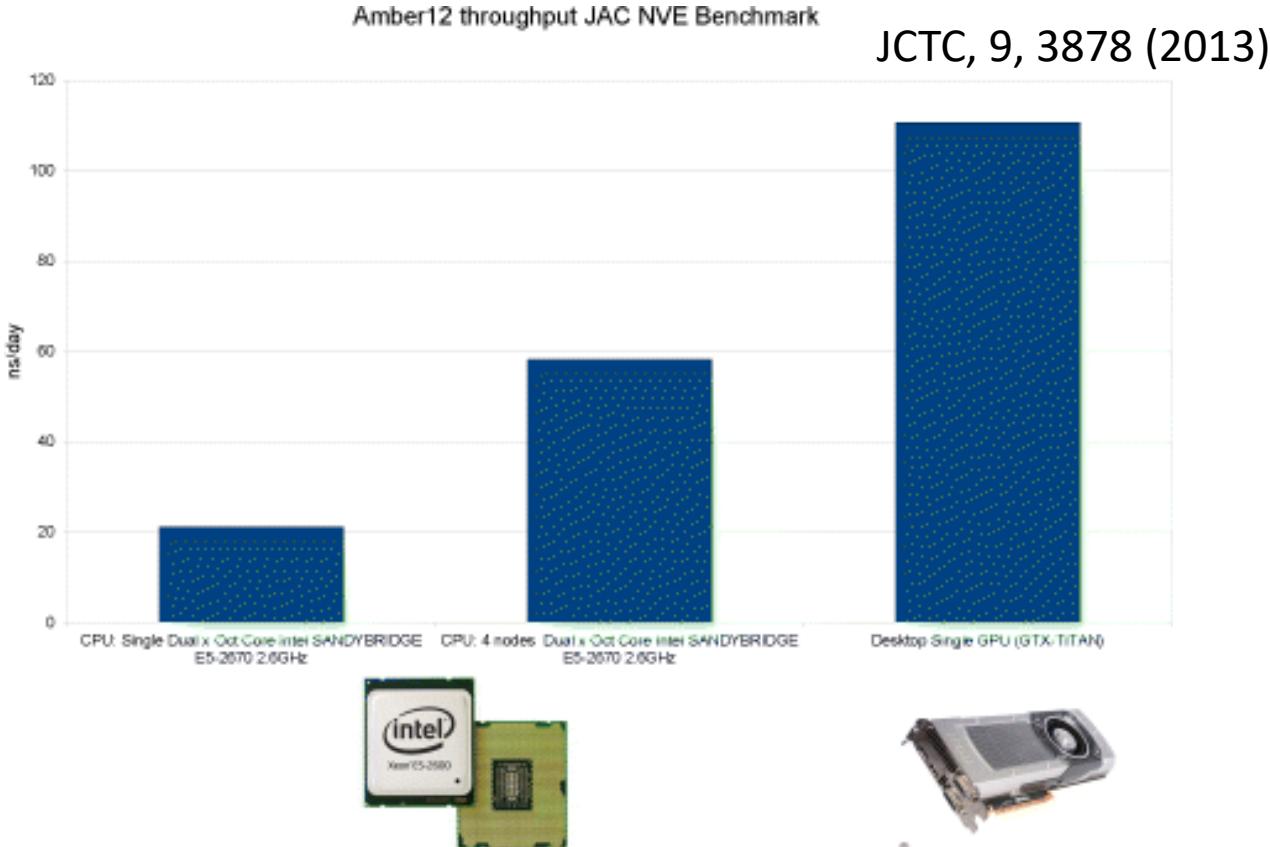
Benchmark

- Solvated protein
- 158945 atoms
- 1.2 nm cutoff radius
- 1 fs time step
- PME for electrostatics

<http://www.charmm-gui.org/>

Amber

- Collection of independent routines
- It uses Sander/PMEMD for solving the Newton's equations
- It offers a robust set of analysis tools



```
module load gcc/2017.1.132-GCC-5.4.0-2.26 CUDA/8.0.44
module load impi/2017.1.132
module load ifort/2017.1.132-GCC-5.4.0-2.26 CUDA/8.0.44
ml Amber/16-AmberTools-16-patchlevel-20-7-hpc2n
srun pmemd.MPI -O -i input.mdin
srun pmemd.cuda.MPI -O -I input.mdin
```

AMBER Tools

- For setting up a simulation (initial structure, solvation, ions, ...):

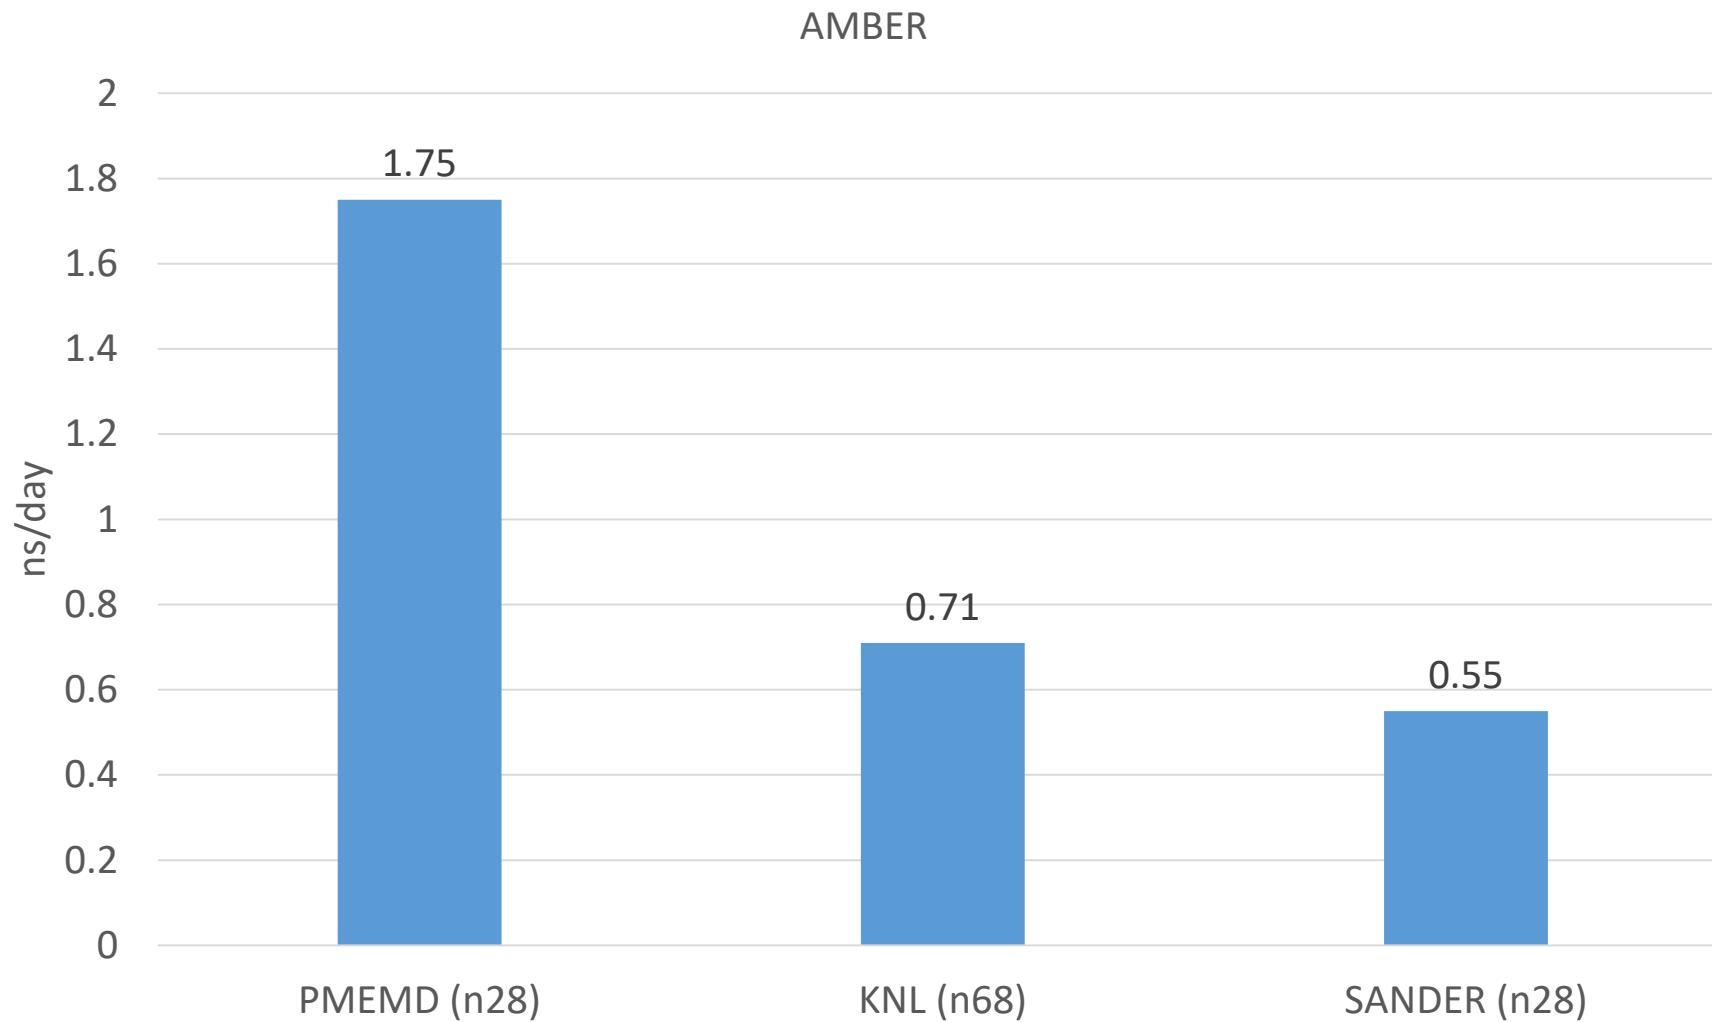
Load the modules:

```
$ml icc/2017.4.196-GCC-6.4.0-2.28 ifort/2017.4.196-GCC-6.4.0-2.28 impi/2017.3.196  
$ml Amber/16-AmberTools-17-patchlevel-8-12
```

On the command line:

```
$tleap, antechamber, cpptraj, ...
```

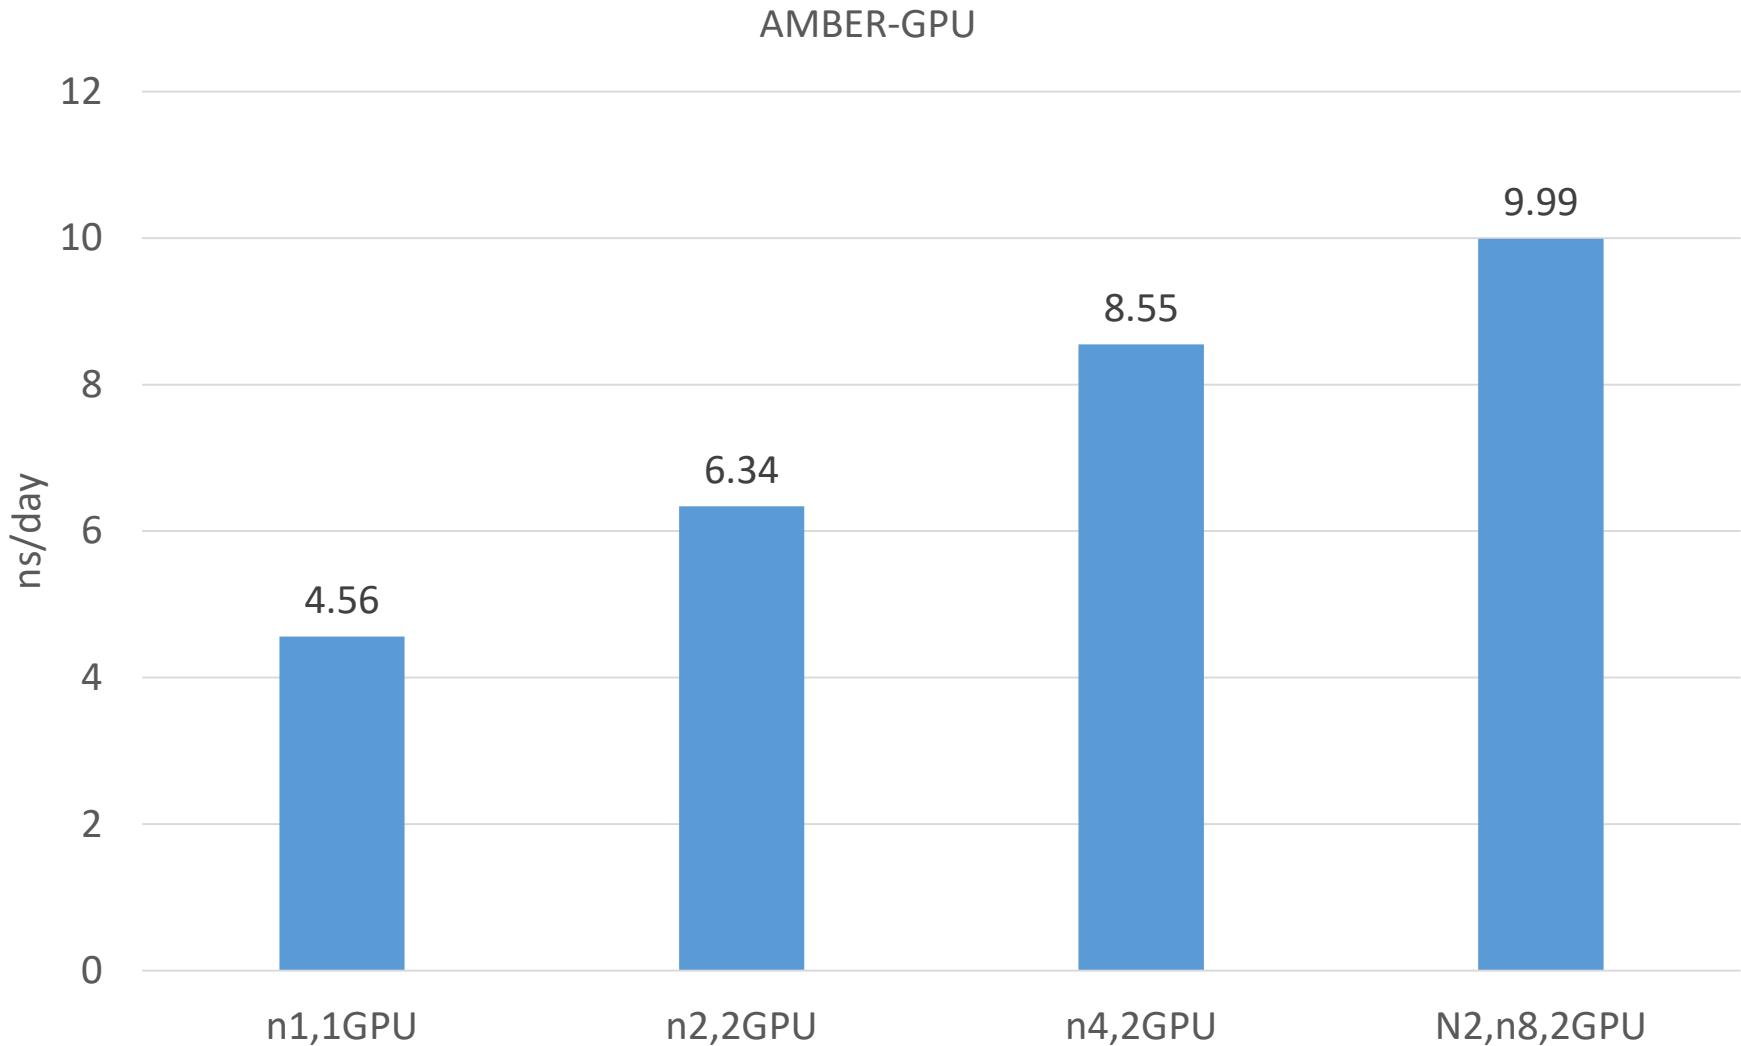
AMBER



AMBER

On Kebnekaise,
best performance
is achieved with 4
MPIs/Node and
using 2 GPU cards

Notice that, the
remaining CPUs
are not used.



Comment on the CUDA version

- AMBER-CUDA does all computations on the GPUs. Thus, the remaining cores remain idle. One can use those cores by launching an additional MPI simulation:

```
#!/bin/bash
#SBATCH -N 1
#SBATCH -n 28
#SBATCH --gres=gpu:k80:2
#Load the AMBER-CUDA version

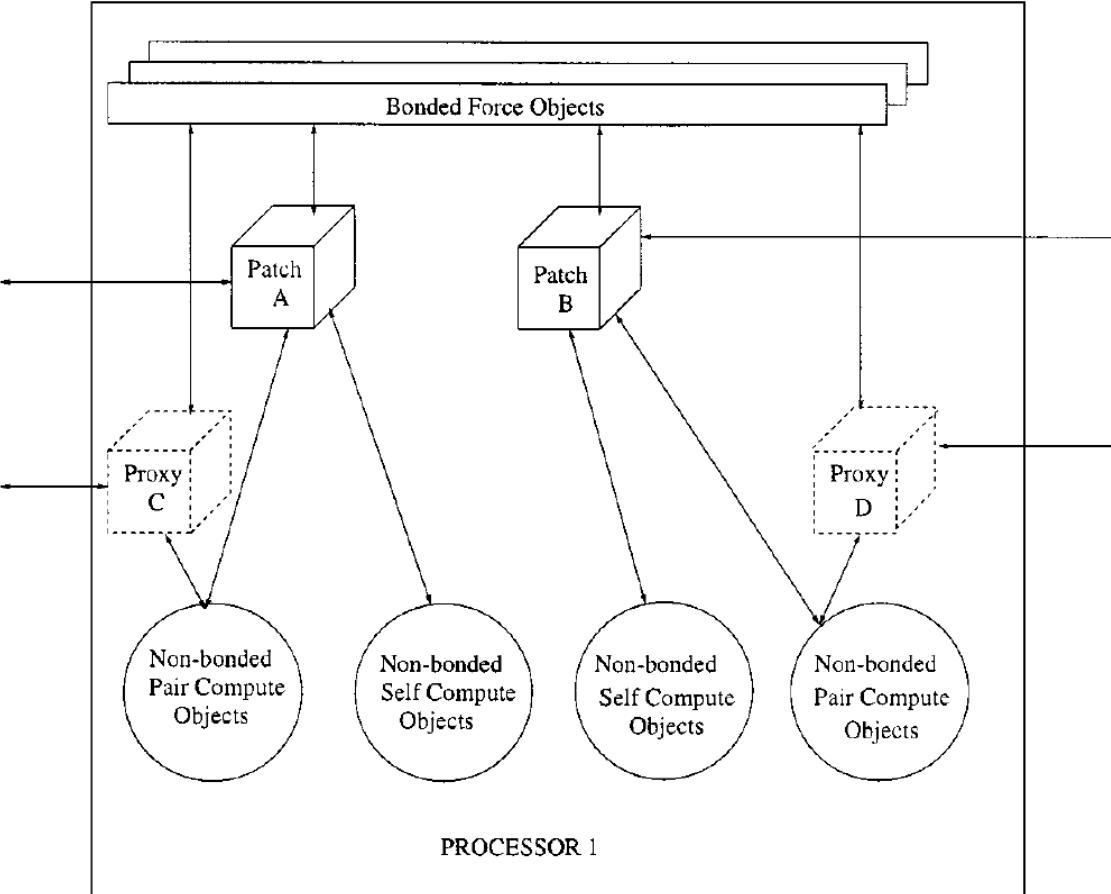
export init="step3_charmm2amber"
export pstep="step4.0_minimization"
export istep="step4.1_equilibration"
mpirun -np 4 pmemd.cuda.MPI -O -i ${istep}.mdin -p ${init}.parm7 -c ${pstep}.rst7 -o ${istep}.mdout -r ${istep}.rst7 -inf ${istep}.mdinfo -ref ${init}.rst7 -x ${istep}.nc &

export istep2="step4.1_equilibration_parallel"
mpirun -np 24 pmemd.MPI -O -i ${istep2}.mdin -p ${init}.parm7 -c ${pstep}.rst7 -o ${istep2}.mdout -r ${istep2}.rst7 -inf ${istep2}.mdinfo -ref ${init}.rst7 -x ${istep2}.nc &

wait
```

NAMD

- Based on charm++ communication protocol
- It is object-oriented
- Versions: single node, multi-node, GPU, and KNL
- Highly scalable



NAMD (SMP)

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 00:50:00
#Number of nodes
#SBATCH -N 1
#Ask for 28 processes
#SBATCH -n 28
#SBATCH --exclusive
#Load modules necessary for running NAMD
module add intel/2017.1.132-GCC-6.3.0-2.27 impi/2017.1.132
module add NAMD/2.12-nompi
#Execute NAMD
namd2 +p 28 +setcpuaffinity step4_equilibration.inp > output_smp.dat
```

NAMD (GPU) (single node)

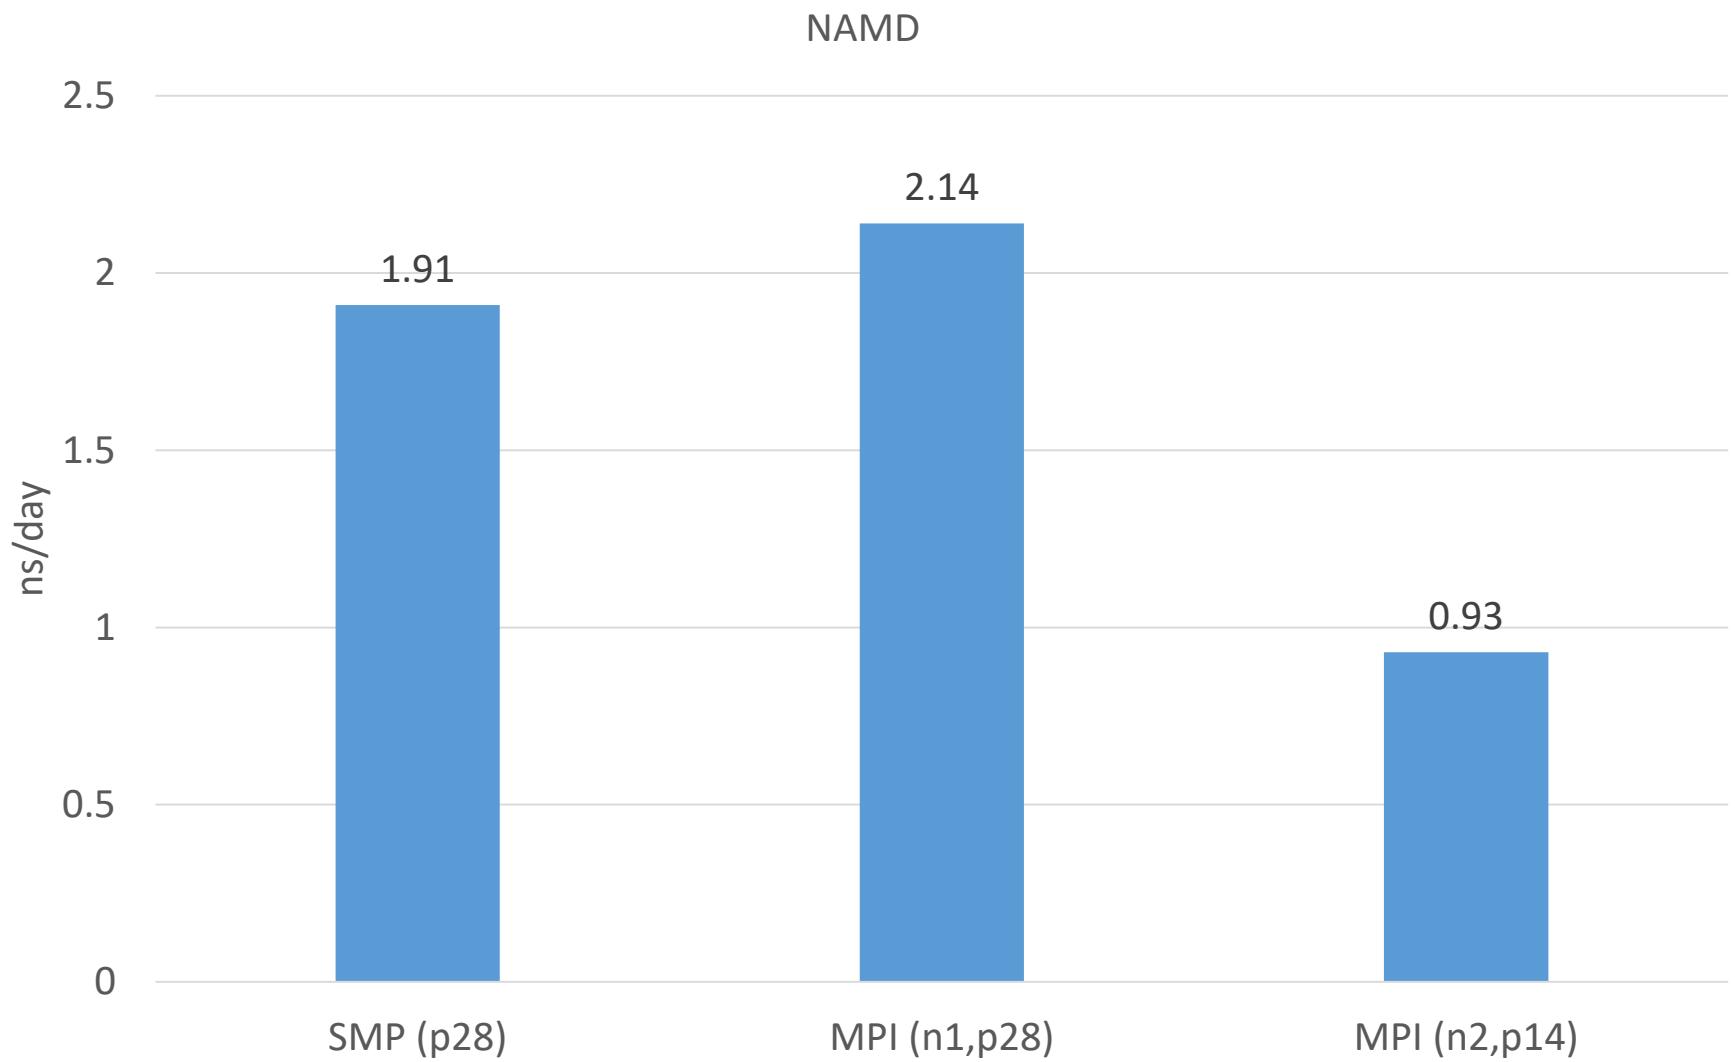
```
#!/bin/bash
#SBATCH -A staff
#SBATCH -t 00:50:00
#SBATCH -N 1
#SBATCH -n 28
#SBATCH --exclusive
#Ask for 2 GPU cards
#SBATCH --gres=gpu:k80:2
#Load modules necessary for running NAMD
module add GCC/5.4.0-2.26 CUDA/8.0.61_375.26 OpenMPI/2.0.2
module add NAMD/2.12-nompi
#Execute NAMD
namd2 +p 28 +setcpuaffinity step4_equilibration.inp > output_smp.dat
```

NAMD

- Colvars module for free energy calculations can be run on GPUs.

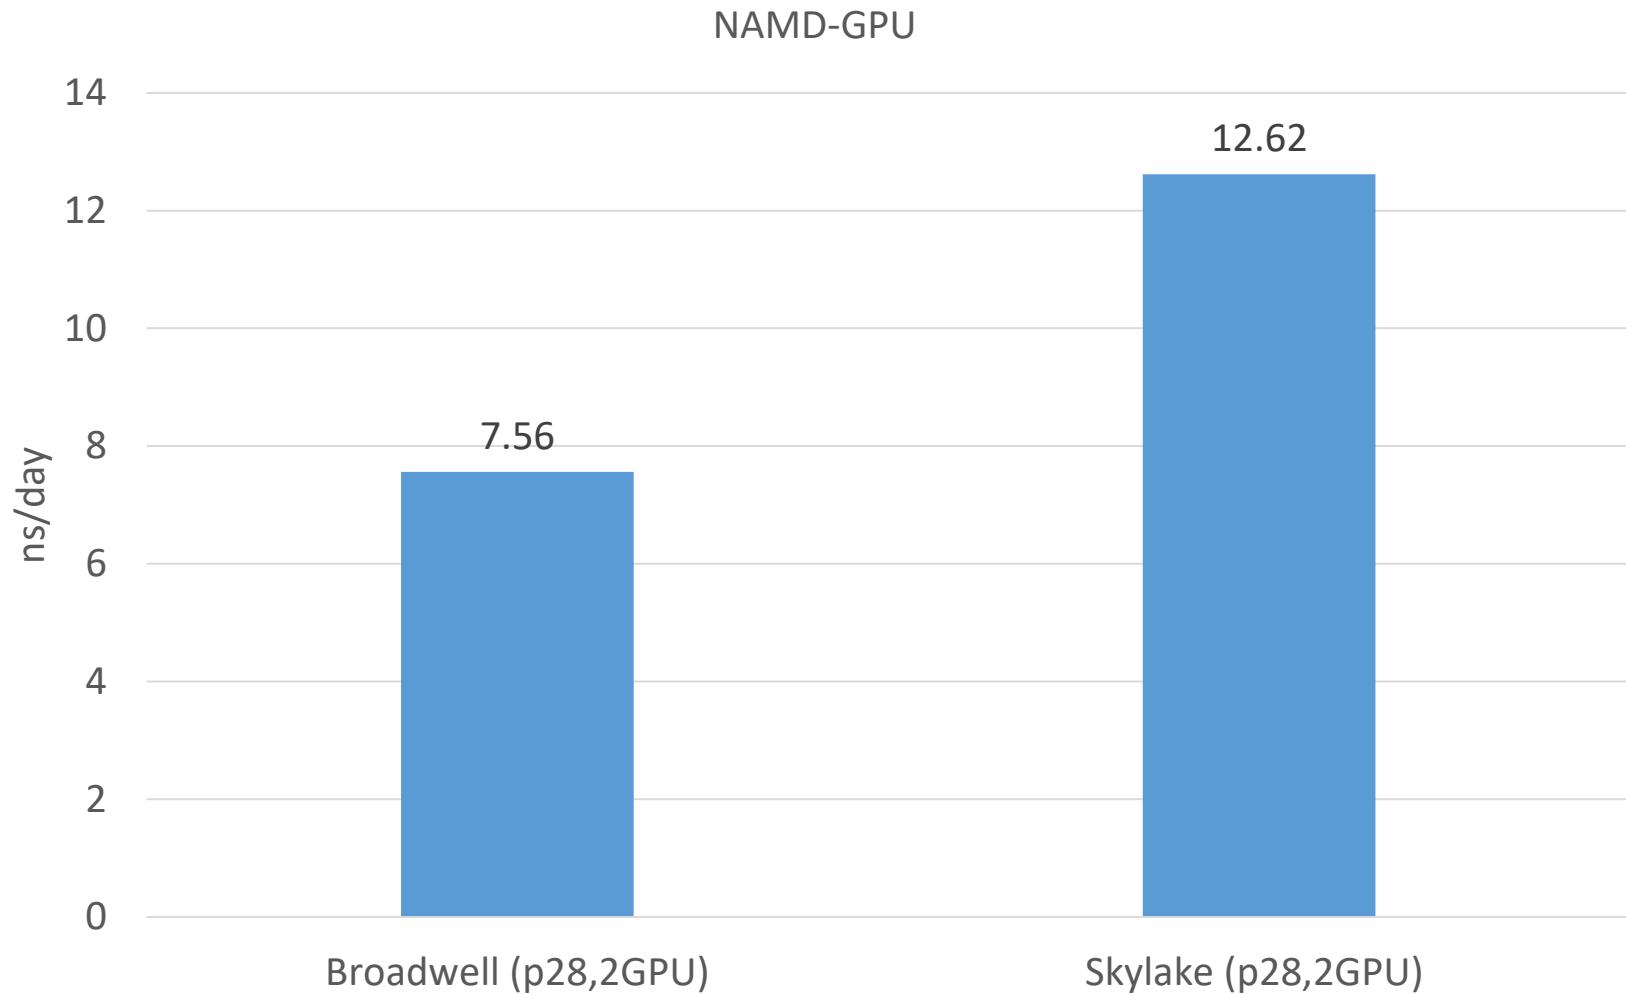
NAMD

Use the `+setcpuaffinity`
flag for a 10% speedup



NAMD

With the new Volta GPU cards one can speed up the simulation by **1.66x**



GROMACS

```
#SBATCH -n 4
#SBATCH -c 7
# Asking for 2 GPUs
#SBATCH --gres=gpu:k80:2
#SBATCH -p batch
```

```
ml GCC/5.4.0-2.26
ml CUDA/8.0.44
ml OpenMPI/2.0.1
ml GROMACS/2016-hybrid
```

```
if [ -n "$SLURM_CPUS_PER_TASK" ]; then
    mdargs="-ntomp $SLURM_CPUS_PER_TASK"
else
    mdargs="-ntomp 1"
fi
```

```
srun -n $SLURM_NTASKS gmx_mpi mdrun $mdargs -npme 0 -dlb yes -v -deffnm step4.1_equilibration
```

GROMACS recognizes the number of available GPU cards

gmx tune_pme

Individual timings for input file 0 (npt_bench00.tpr):

PME ranks	Gcycles	ns/day	PME/f	Remark
0	4355.019	57.776	-	OK.
0	4547.105	55.335	-	OK.
0	4289.420	58.659	-	OK.
-1(0)	4455.791	56.469	-	OK.
-1(0)	4440.157	56.668	-	OK.
-1(0)	4275.551	58.850	-	OK.

Tuning took 7.7 minutes.

Summary of successful runs:

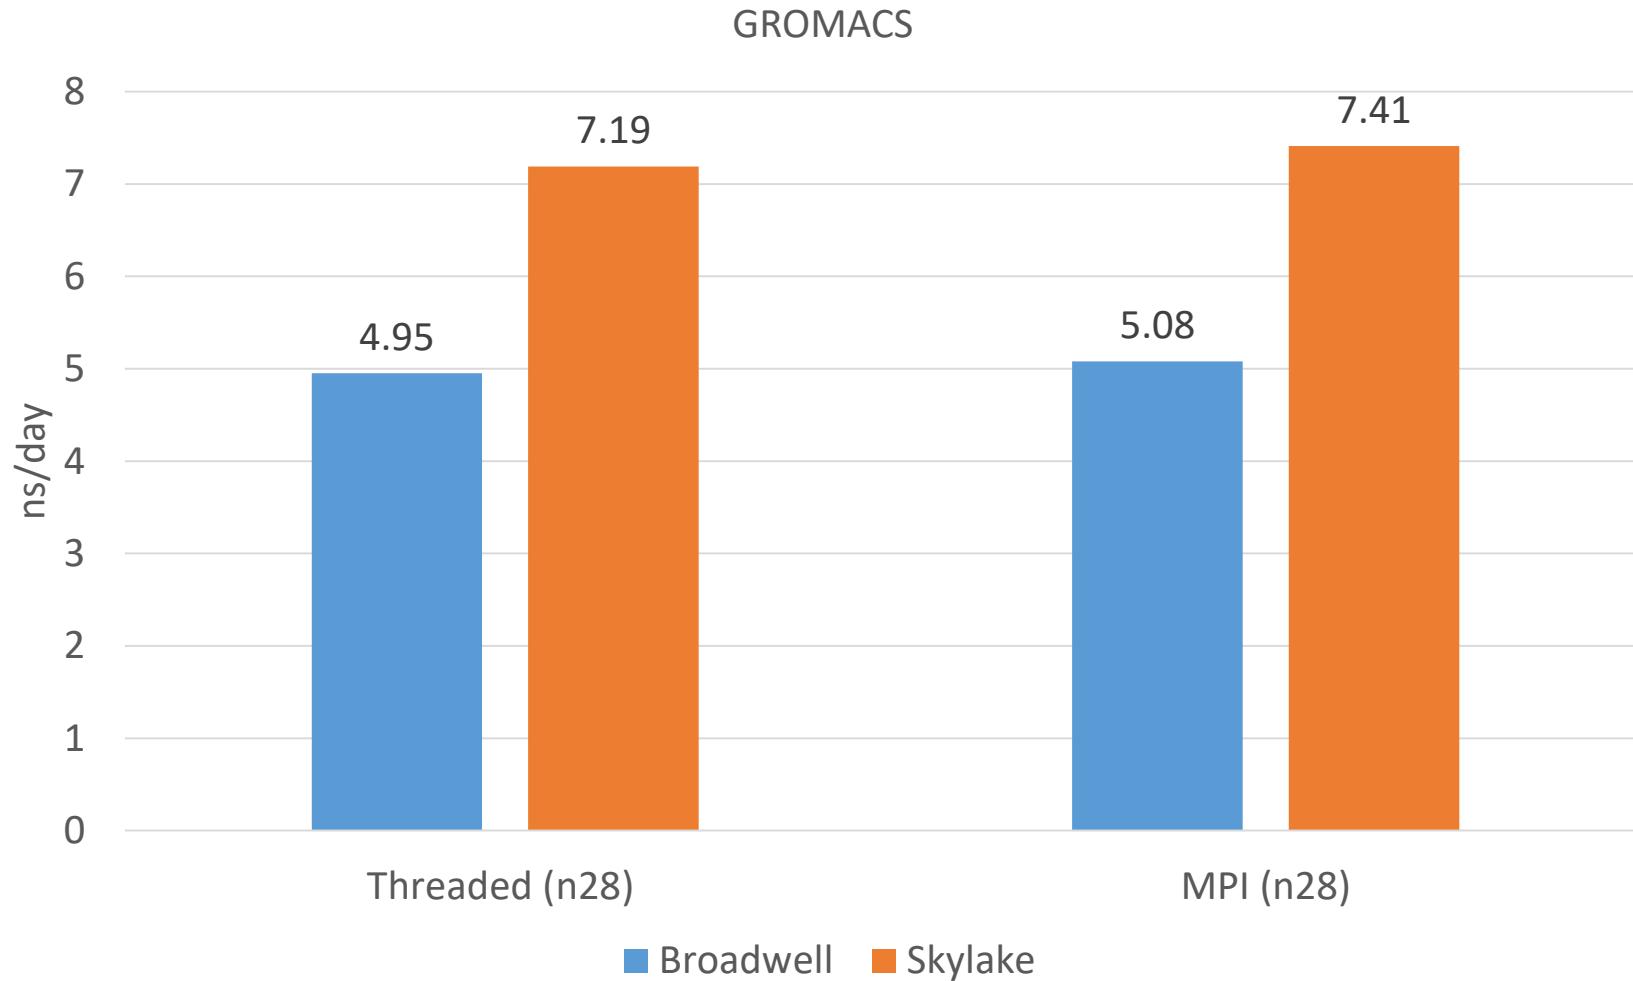
Line	tpr	PME ranks	Gcycles	Av.	Std.dev.	ns/day	PME/f	DD grid
0	0	0	4397.181	133.917	57.257	-	4 1 1	
1	0	-1(0)	4390.500	99.855	57.329	-	4 1 1	

Best performance was achieved with the automatic number of PME ranks (see line 1)

Please use this command line to launch the simulation:

mpirun -np 4 gmx_mpi mdrun -npme -1 -s npt.tpr -ntomp 7 -dlb yes

GROMACS



GROMACS KNL

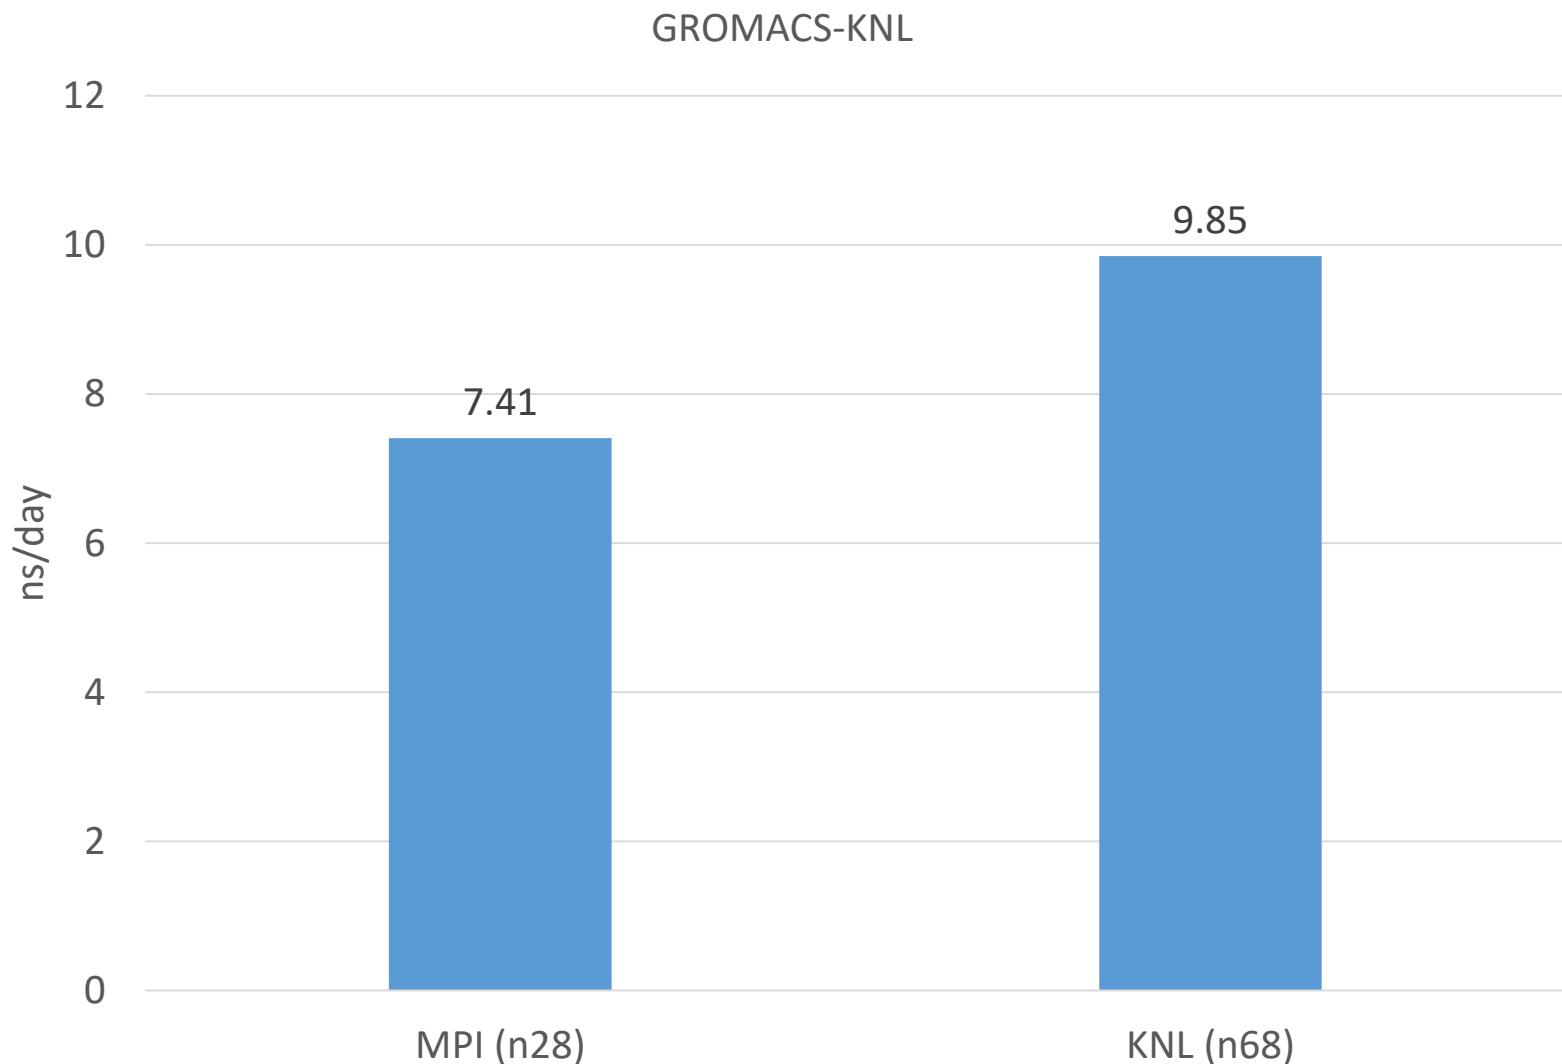
```
#!/bin/bash
#SBATCH -A project_ID
#SBATCH -t 00:50:00
#SBATCH -N 1
#SBATCH -n 68
#SBATCH -p knl
#SBATCH --constraint=cache,quad
#SBATCH --exclusive

ml GCC/7.3.0-2.30 OpenMPI/3.1.1
ml GROMACS/2018.3

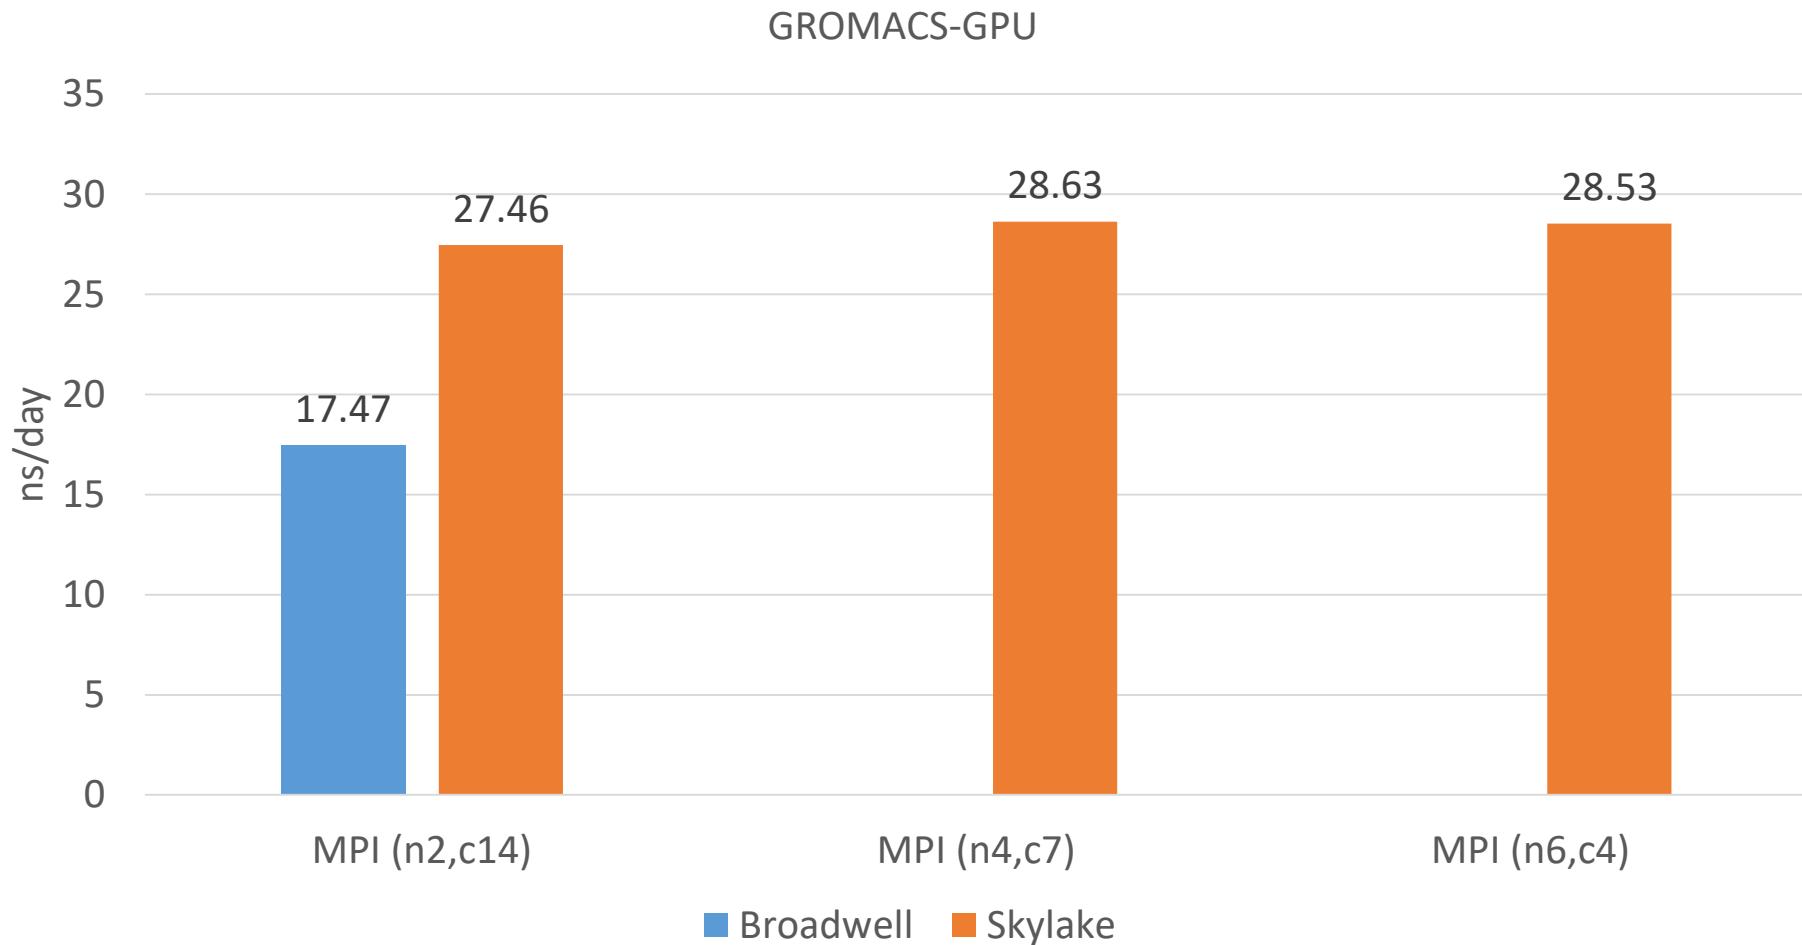
export OMP_NUM_THREADS=2
gmx mdrun -ntmpi 68 -npme 18 -ntomp 2 -pin on -pinoffset 0 -pinstride 2 -dlb auto -v -deffnm step4.1_equilibration
```

GROMACS KNL

- Login to `kebnekaise-knl.hpc2n.umu.se` and submit your script from there
- KNL queue is most of the time available compared to GPU one



GROMACS



GROMACS

- PLUMED, threaded MPI version is not supported.

Instead of:

`gmx mdrun –ntmpi X`

Use:

`mpirun gmx_mpi ...`

- Avoid writing energies
- Offloading features starting from 2018 version

<https://www.hpc2n.umu.se/events/courses/md-course-spring-2019>

LAMMPS (MPI)

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 02:10:00
#Number of nodes
#SBATCH -N 1
#Ask for 28 processes
#SBATCH -n 28

#Load modules necessary for running LAMMPS
module load intel/2017.1.132-GCC-6.3.0-2.27 ifort/2017.1.132-GCC-6.3.0-2.27
impi/2017.1.132
module load LAMMPS/31Mar17

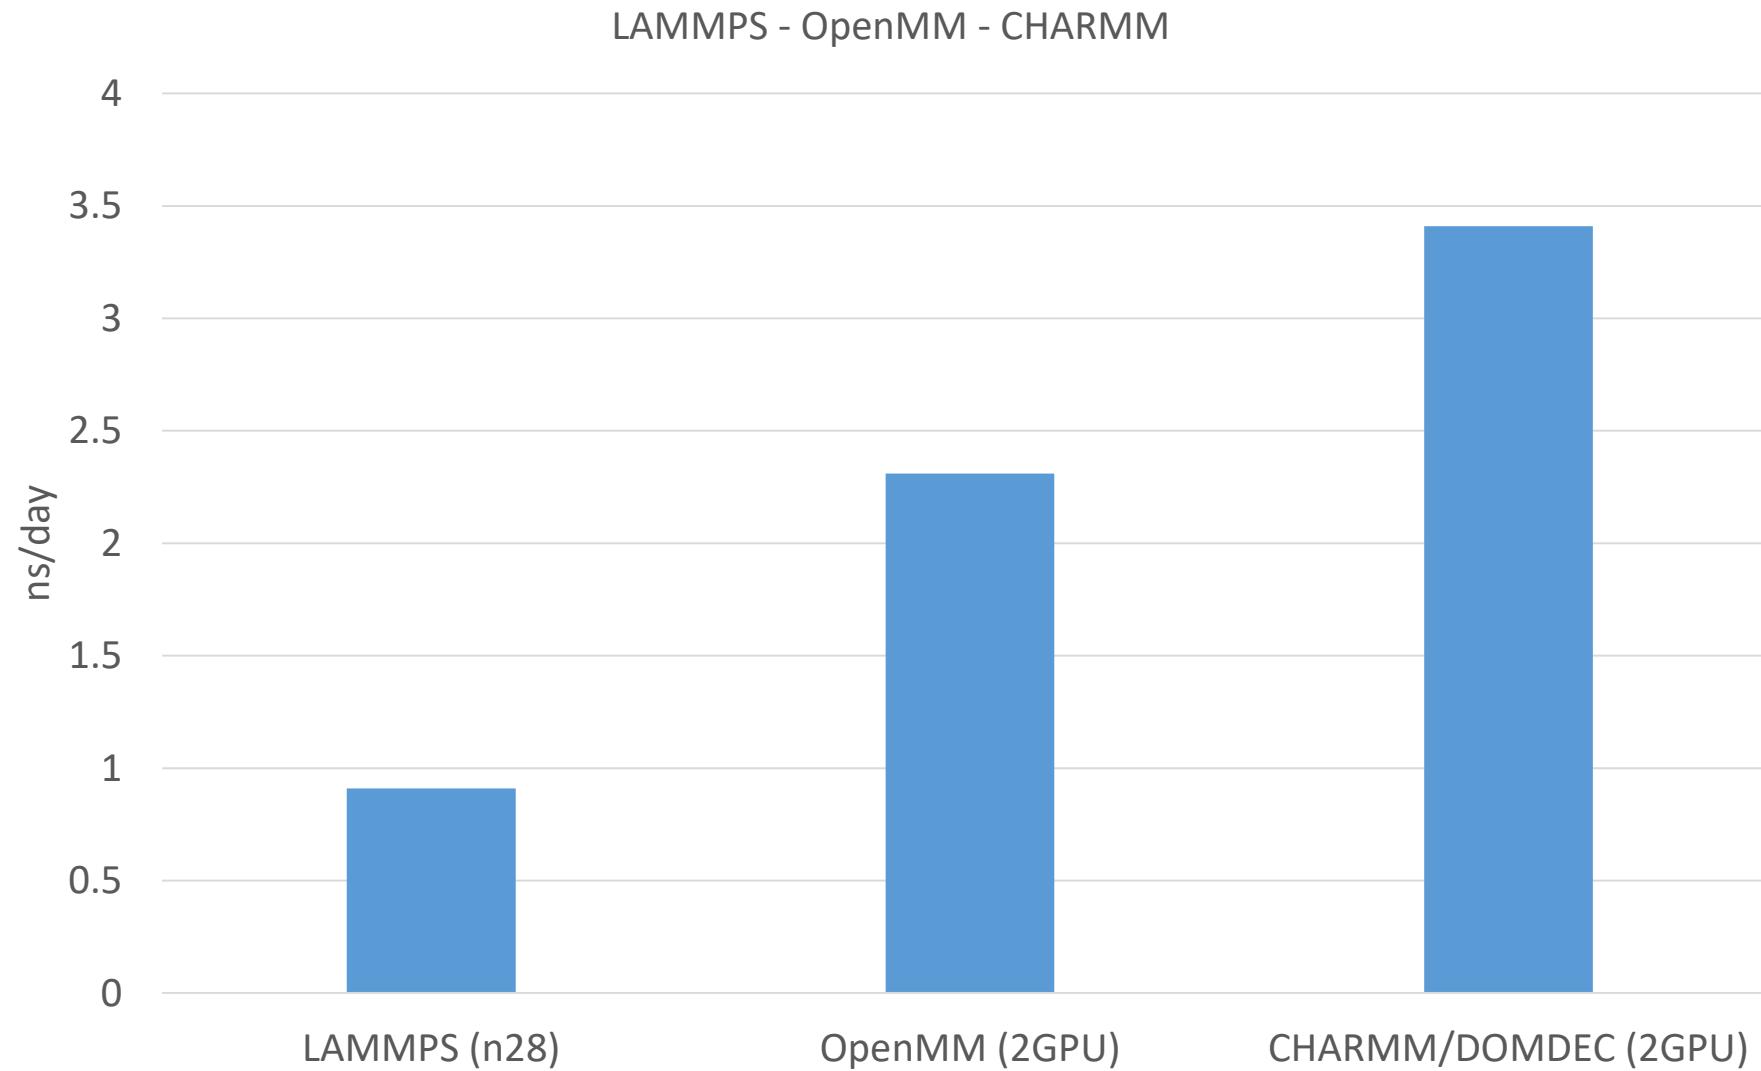
#Execute LAMMPS
srun lmp_intel_cpu -in step4.1_equilibration.inp
```

LAMMPS (OpenMP)

```
#!/bin/bash
#SBATCH -A staff
#Asking for 10 min.
#SBATCH -t 02:10:00
#Number of nodes
#SBATCH -N 1
#Ask for 28 processes
#SBATCH -n 14
#SBATCH -c 2

#Load modules necessary for running LAMMPS
module load intel/2017.1.132-GCC-6.3.0-2.27 ifort/2017.1.132-GCC-6.3.0-2.27
module load impi/2017.1.132
module load LAMMPS/31Mar17
export OMP_NUM_THREADS=2
#Execute LAMMPS
srun lmp_intel_cpu -in step4.1_equilibration.inp
```

LAMMPS - OpenMM - CHARMM



Scratch directory

- If the program writes frequently data, you can get an additional speed up by using the local scratch directory:

```
rm -rf /scratch/test  
export parent=/pfs/nobackup/home/u/username/benchmarks/  
mkdir /scratch/test  
rsync -avzh $parent/ /scratch/test/
```

```
cd /scratch/test
```

~10-15% speedup

```
namd2 +p 28 +setcpuaffinity input.inp > output.dat
```

```
cd $parent
```

```
rsync -avzh /scratch/test/ $parent/
```

```
rm -rf /scratch/test
```

Abisko (all nodes): 352 GB
Kebnekaise, standard nodes: 171 GB
Kebnekaise, GPU nodes: 171 GB
Kebnekaise, Largemem nodes: 352 GB
(a few have 391 GB)

Suggestions

- Check if the software version you are trying is MPI (-n), OpenMP (-c), or hybrid (-n -c). Also, if you are using a GPU version.
- Is your simulation scaling properly?
- Is the cutoff radius for long-range interactions appropriate? PME (gmx tune_pme for GROMACS)?

Tensorflow

```
#!/usr/share/python
```

```
import tensorflow as tf
hello = tf.constant('Hello, TensorFlow!')
sess = tf.Session()
print(sess.run(hello))
```

```
#!/bin/bash
#SBATCH -A staff
#SBATCH -t 00:50:00
#SBATCH -N 1
#SBATCH --exclusive
#SBATCH --gres=gpu:k80:2
```

```
ml icc/2017.1.132-GCC-5.4.0-2.26
ml ifort/2017.1.132-GCC-5.4.0-2.26
ml CUDA/8.0.44 impi/2017.1.132
ml Python/3.6.1
ml Tensorflow/1.3.0-Python-3.6.1
python tensor.py
```

```
name: Tesla K80
major: 3 minor: 7 memoryClockRate (GHz) 0.8235
pciBusID 0000:0f:00.0
Total memory: 11.17GiB
Free memory: 11.10GiB
2017-12-05 17:13:03.555397: W
Found device 1 with properties:
name: Tesla K80
>>> print(sess.run(hello))
b'Hello, TensorFlow!'
```

Tensorflow

```
#!/usr/share/python
import tensorflow as tf

#Parameters
W = tf.Variable([.3],tf.float32)
b = tf.Variable([-3],tf.float32)
#Input and output
x = tf.placeholder(tf.float32)
linear_model = W*x+b
y = tf.placeholder(tf.float32)

#Loss
square_delta = tf.square(linear_model-y)
loss = tf.reduce_sum(square_delta)
#Optimize
optimizer = tf.train.GradientDescentOptimizer(0.01)
train = optimizer.minimize(loss)

init = tf.global_variables_initializer()
sess = tf.Session()
sess.run(init)

for i in range (1000):
    sess.run(train,{x:[1,2,3,4],y:[0,-1,-2,-3]})

print(sess.run([W,b]))
```

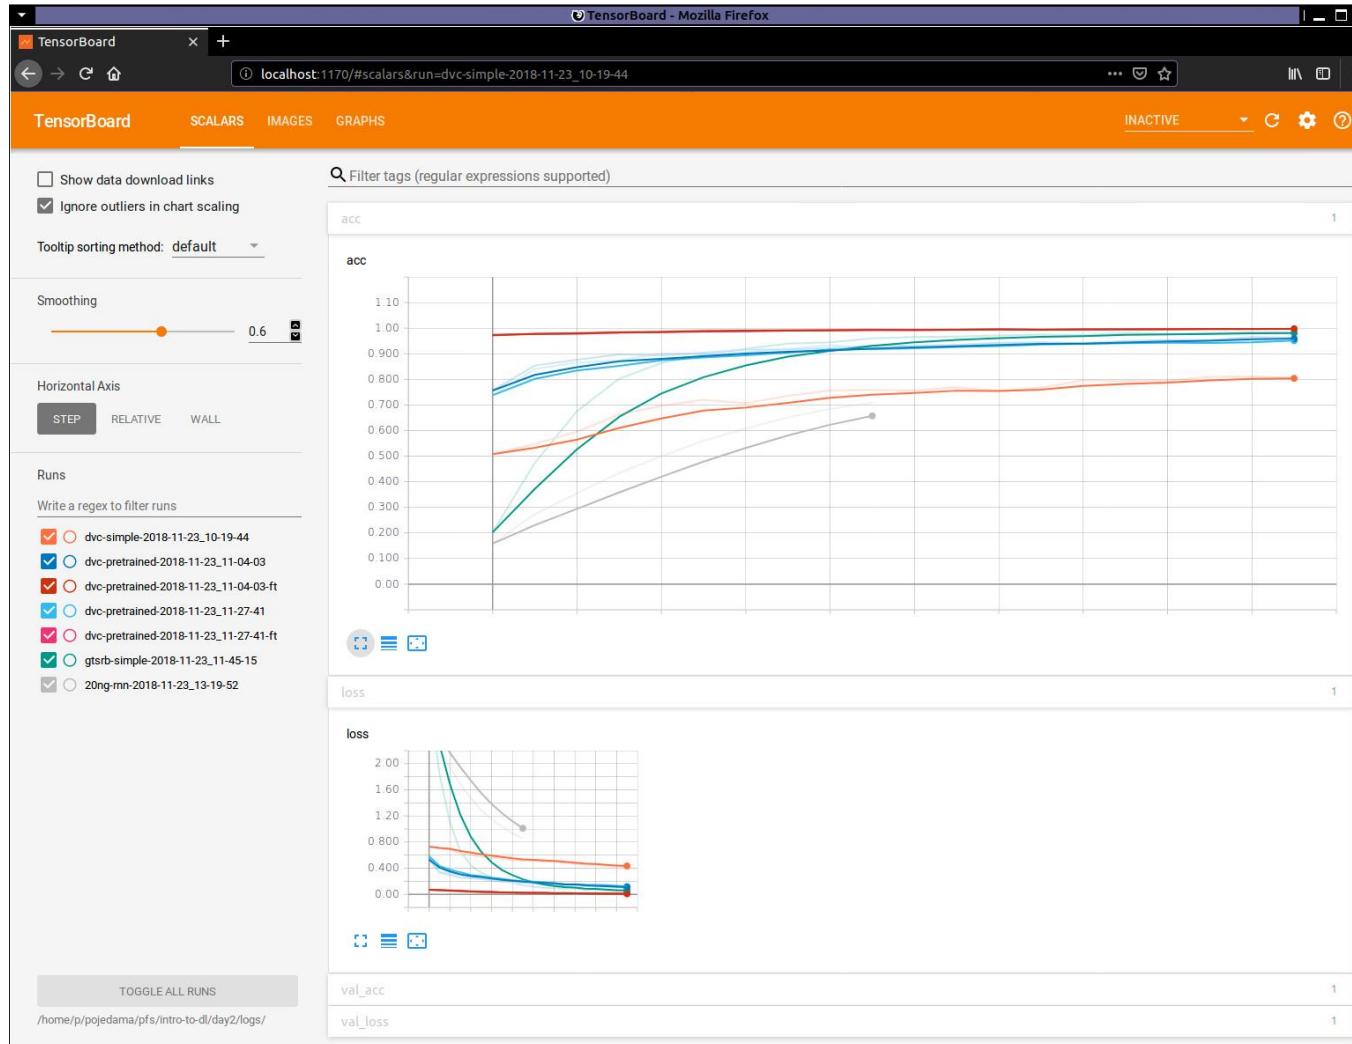
Computing the loss

name: Quadro K6000
major: 3 minor: 5 memoryClockRate (GHz) 0.9015
pciBusID 0000:06:00.0
Total memory: 11.92GiB
Free memory: 11.82GiB
Creating TensorFlow device (/gpu:0) -> (device: 0, name: Quadro K6000, pci bus id: 0000:06:00.0)

[array([-0.9999969], dtype=float32), array([0.99999082], dtype=float32)]

Tensorflow+Tensorboard

```
module load GCC/7.3.0-2.30 CUDA/9.2.88 OpenMPI/3.1.1 #On the command line execute these instructions  
module load TensorFlow/1.10.1-Python-2.7.15 Keras/2.2.2-Python-2.7.15  
tensorboard --logdir=/home/p/<path-to-directory>/logs --port=1170
```



To monitor the behavior of your jobs open a web browser and type:

`localhost:1170`

GAUSSIAN+GPU

Initial input file:

```
$more input_bk.com
%chk=geom_optim.chk
%mem=16GB
#UB3LYP/6-31+G(d) OPT=(ModRedun) SCF=(MaxCycle=256) pop=none NoSymm

45 atoms structure, RESP

+5 15
0      3.744336      -1.126487      6.111505
P      2.893853      -1.251776      4.246949
O      4.150424      -3.154051      4.078061
```

Corrected input file:

```
$more input.com
%cpu=0,1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,25,26,27
%gpucpu=0-3=0,7,14,21
%chk=geom_optim.chk
%mem=16GB
#UB3LYP/6-31+G(d) OPT=(ModRedun) SCF=(MaxCycle=256) pop=none NoSymm
```

```
#!/bin/bash
#SBATCH -A staff
#SBATCH -N 1
#SBATCH -c 28
#SBATCH --exclusive
#SBATCH --gres=gpu:k80:2
#SBATCH --time=00:10:00

module add gaussian/16.A.03-AVX2
# Assume that the job file are located
g16.set-cpu+gpu-list input.com
time g16 input
```

Integral=(FineGrid,Acc2E=10) Constants=2006

VASP

Batch script:

```
#SBATCH -c 10  (number of cores per task)
#SBATCH -n 1    (number of tasks)
```

```
ml ifort/2017.1.132-GCC-6.3.0-2.27 OpenMPI/2.0.2
ml NWChem/6.6.revision27746-2015-10-20-Python-2.7.12
```

```
mpirun -n 10 nwchem input.nw      (this script will fail)
```

VASP

Batch script:

```
#SBATCH -c 1      (number of cores per task)
#SBATCH -n 10     (number of tasks)
```

```
ml ifort/2017.1.132-GCC-6.3.0-2.27 OpenMPI/2.0.2
ml NWChem/6.6.revision27746-2015-10-20-Python-2.7.12
```

```
mpirun -n 10 nwchem input.nw      (this script is correct!)
```

Suggestions

- One common issue with Comp. Chem. Software is related to the available memory. A possible solution is to request a larger number of cpus (with `-n` or `-c` in SLURM) but make use of less than that number.

Final comments

- Check if the software version you are trying is MPI (-n), OpenMP (-c), or hybrid (-n -c). Also, if you are using a GPU version.
- Is the number of cores/nodes you request optimal? Trying a short simulation could help.
- Is your simulation scaling well?
- Could KNL nodes speed up your workflow?
- Use **job-usage** script on the command line

Final comments

- Upon asking a question to the support team: **Make a folder with the case which displays the issue. Including all relevant files would be useful.**
- Avoid using SLURM batch scripts from other centers, use those provided in our web site.
- Don't load modules in your .bashrc file, this can caused troubles with software.
- Don't forget to cite HPC2N and the SNIC project upon publishing.